



## PyG 2.0

# Advanced Representation Learning on Graphs



**Matthias Fey**

[matthias@pyg.org](mailto:matthias@pyg.org)

<https://pyg.org>

[GitHub: pyg-team/pytorch-geometric](https://github.com/pyg-team/pytorch-geometric)

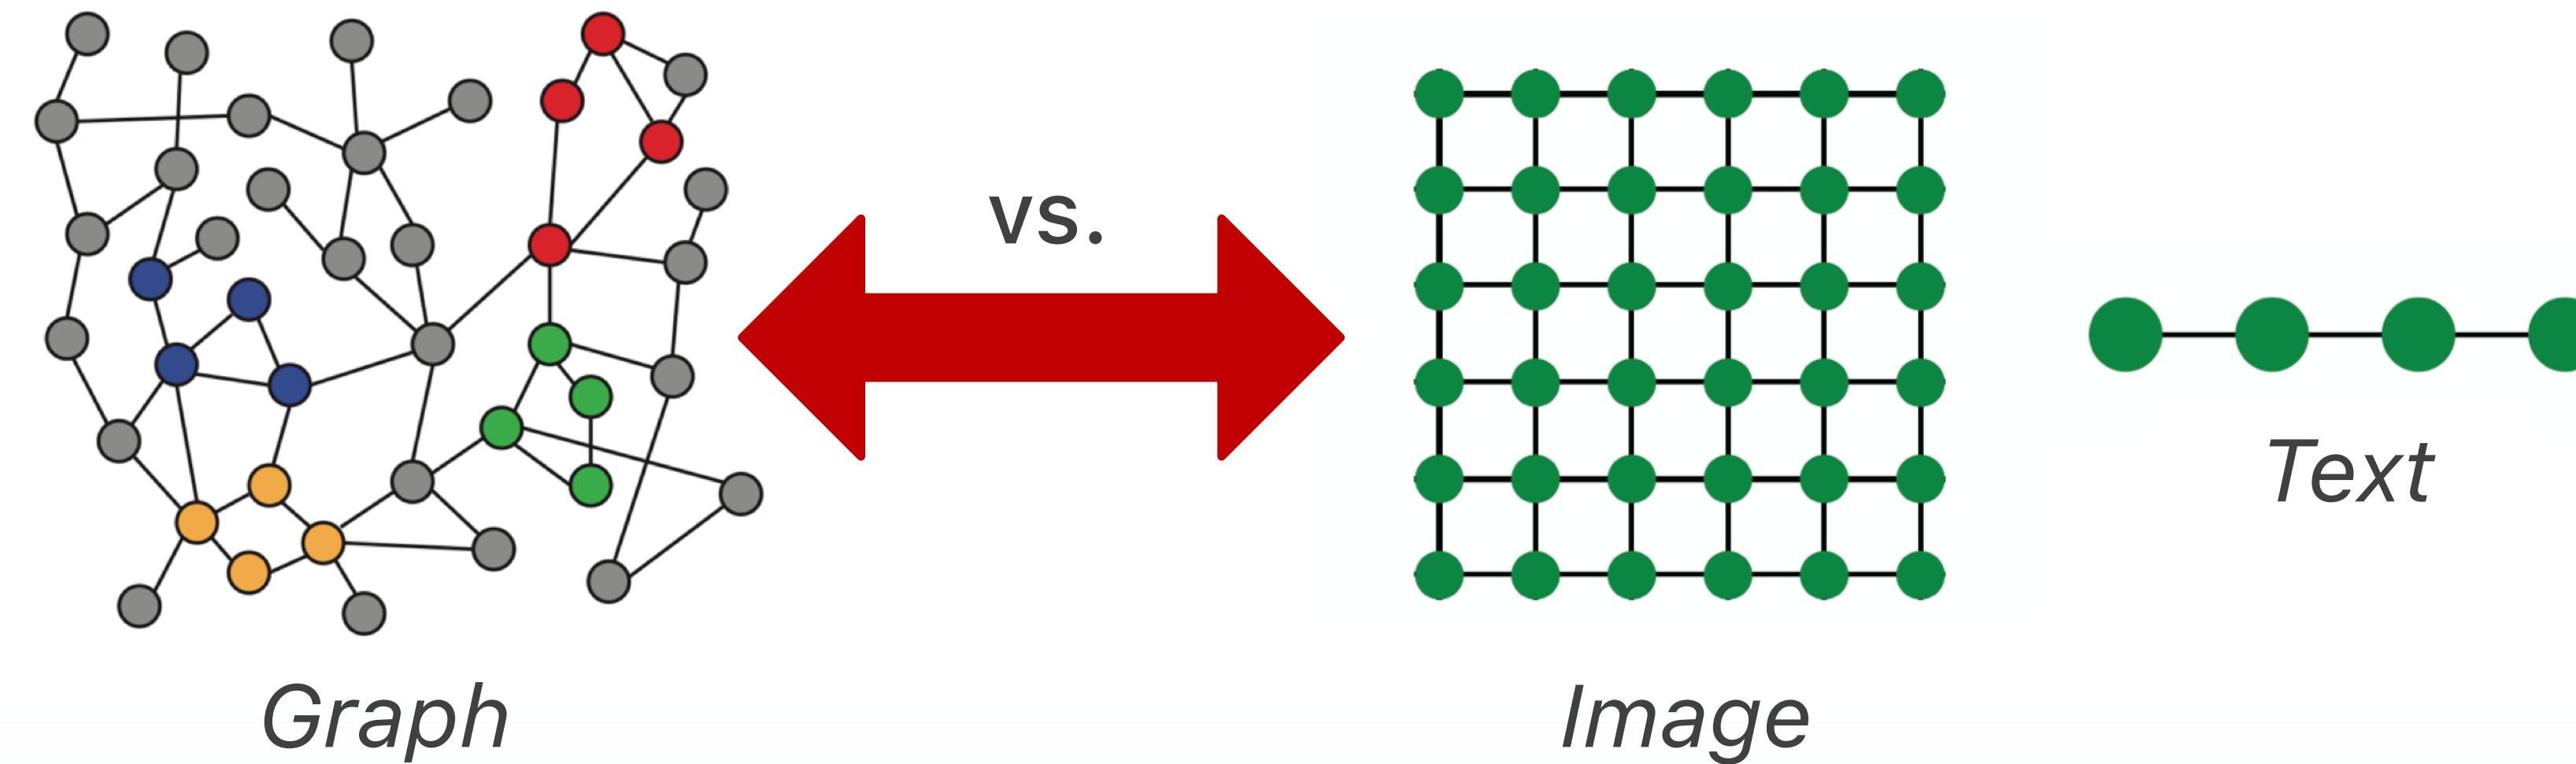


```
conda install pyg -c pyg
```



# Motivation

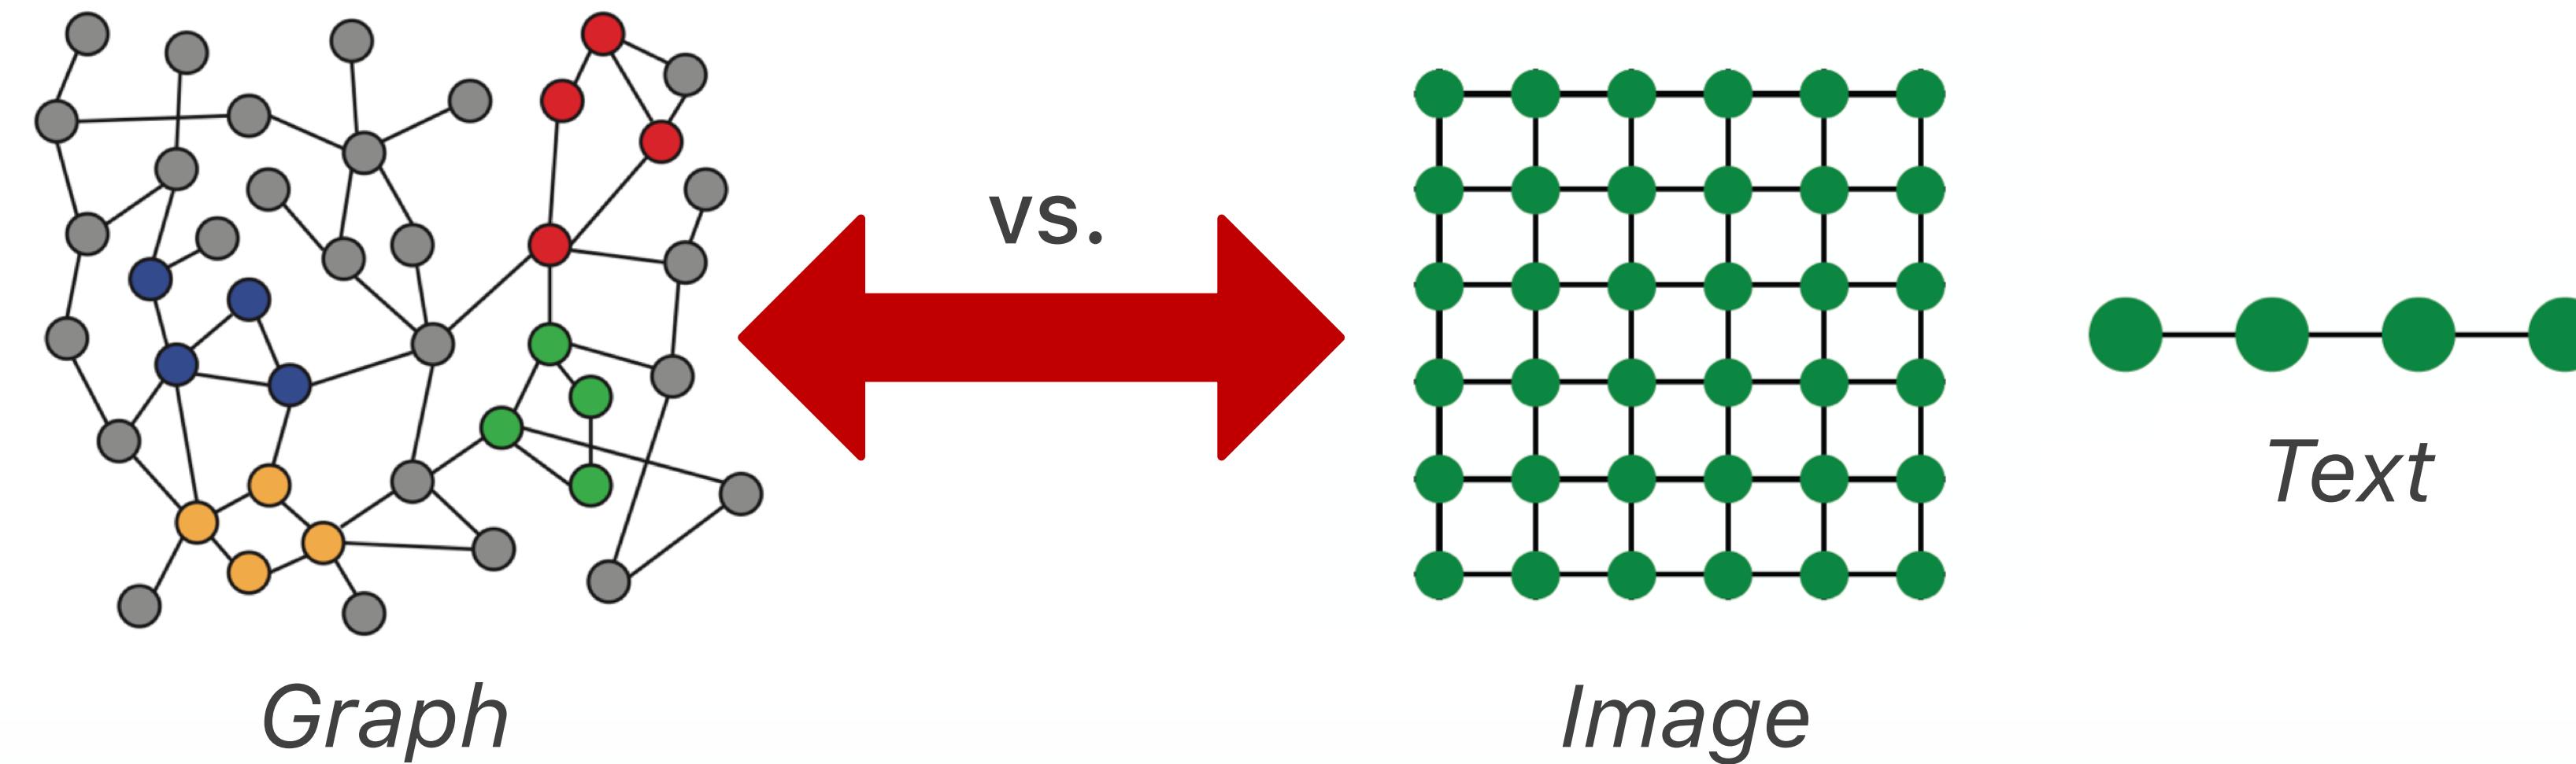
The modern deep learning toolbox is designed for sequences and grids





# Motivation

The modern deep learning toolbox is designed for sequences and grids

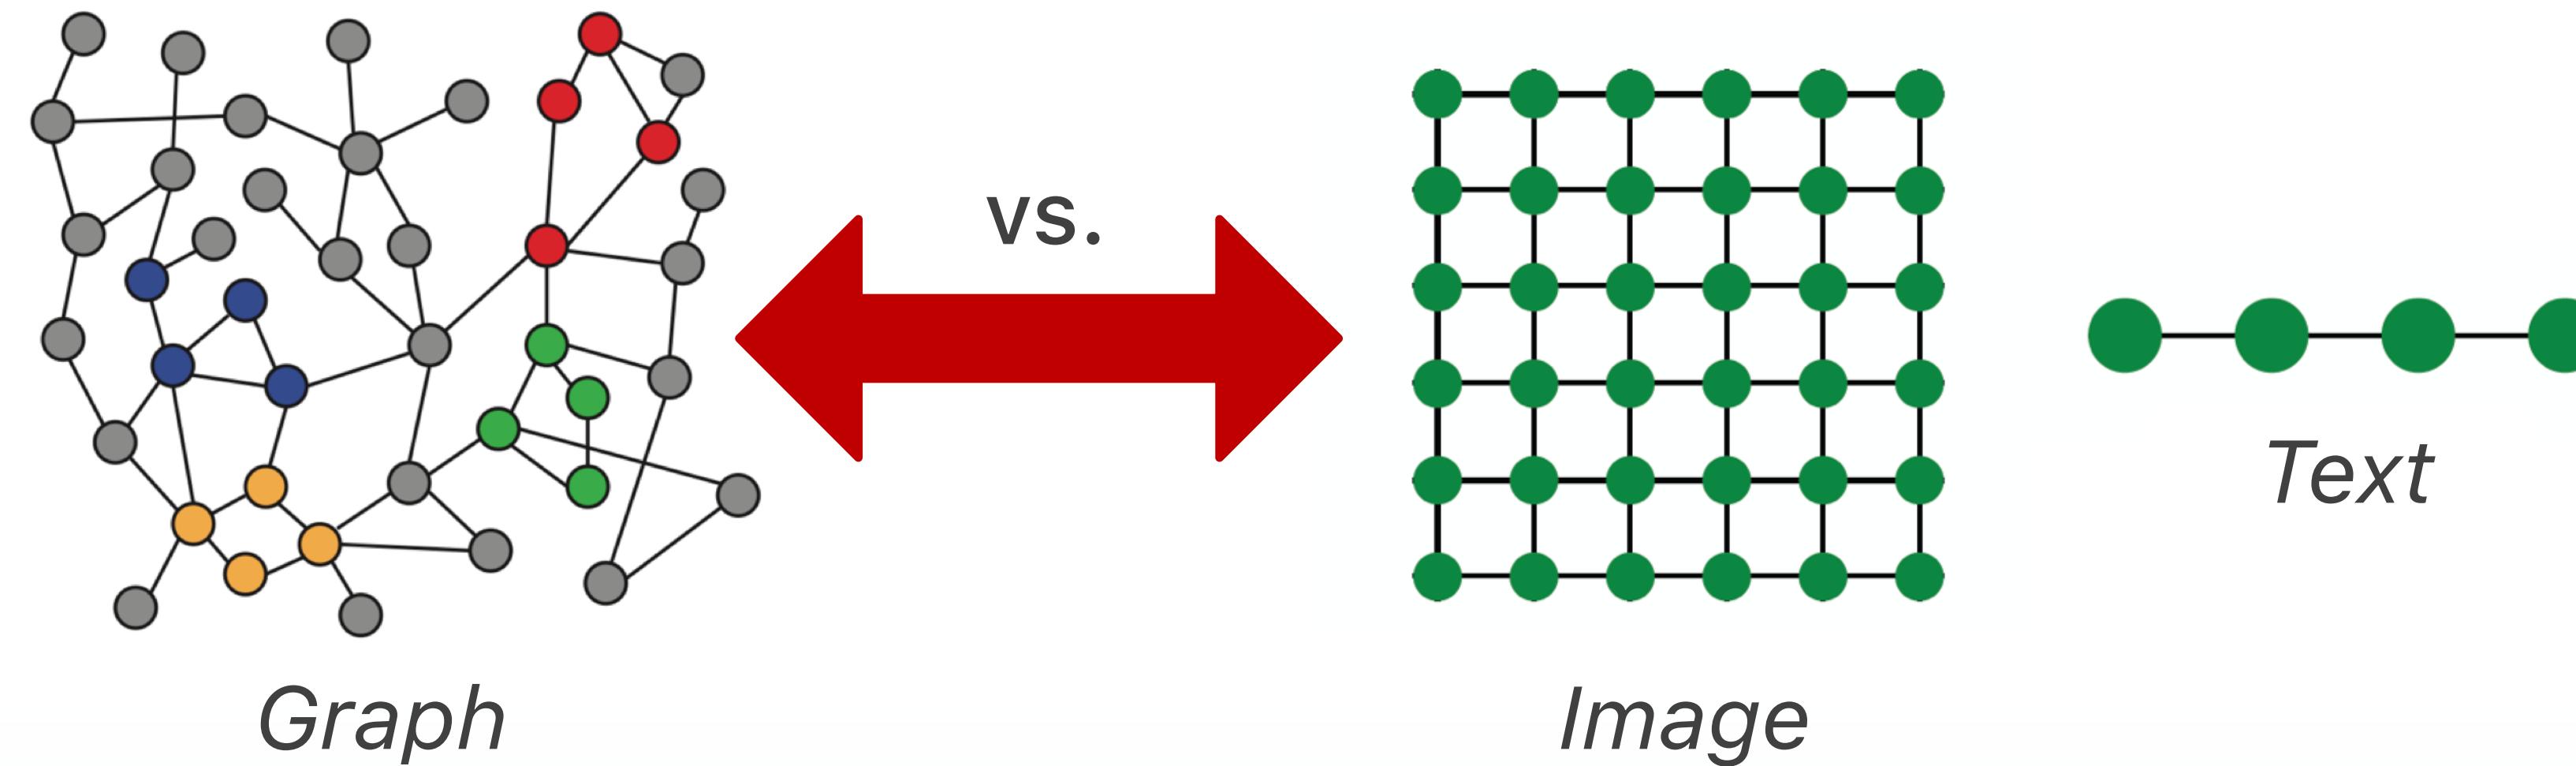


- Arbitrary size and complex topological structure



# Motivation

The modern deep learning toolbox is designed for sequences and grids

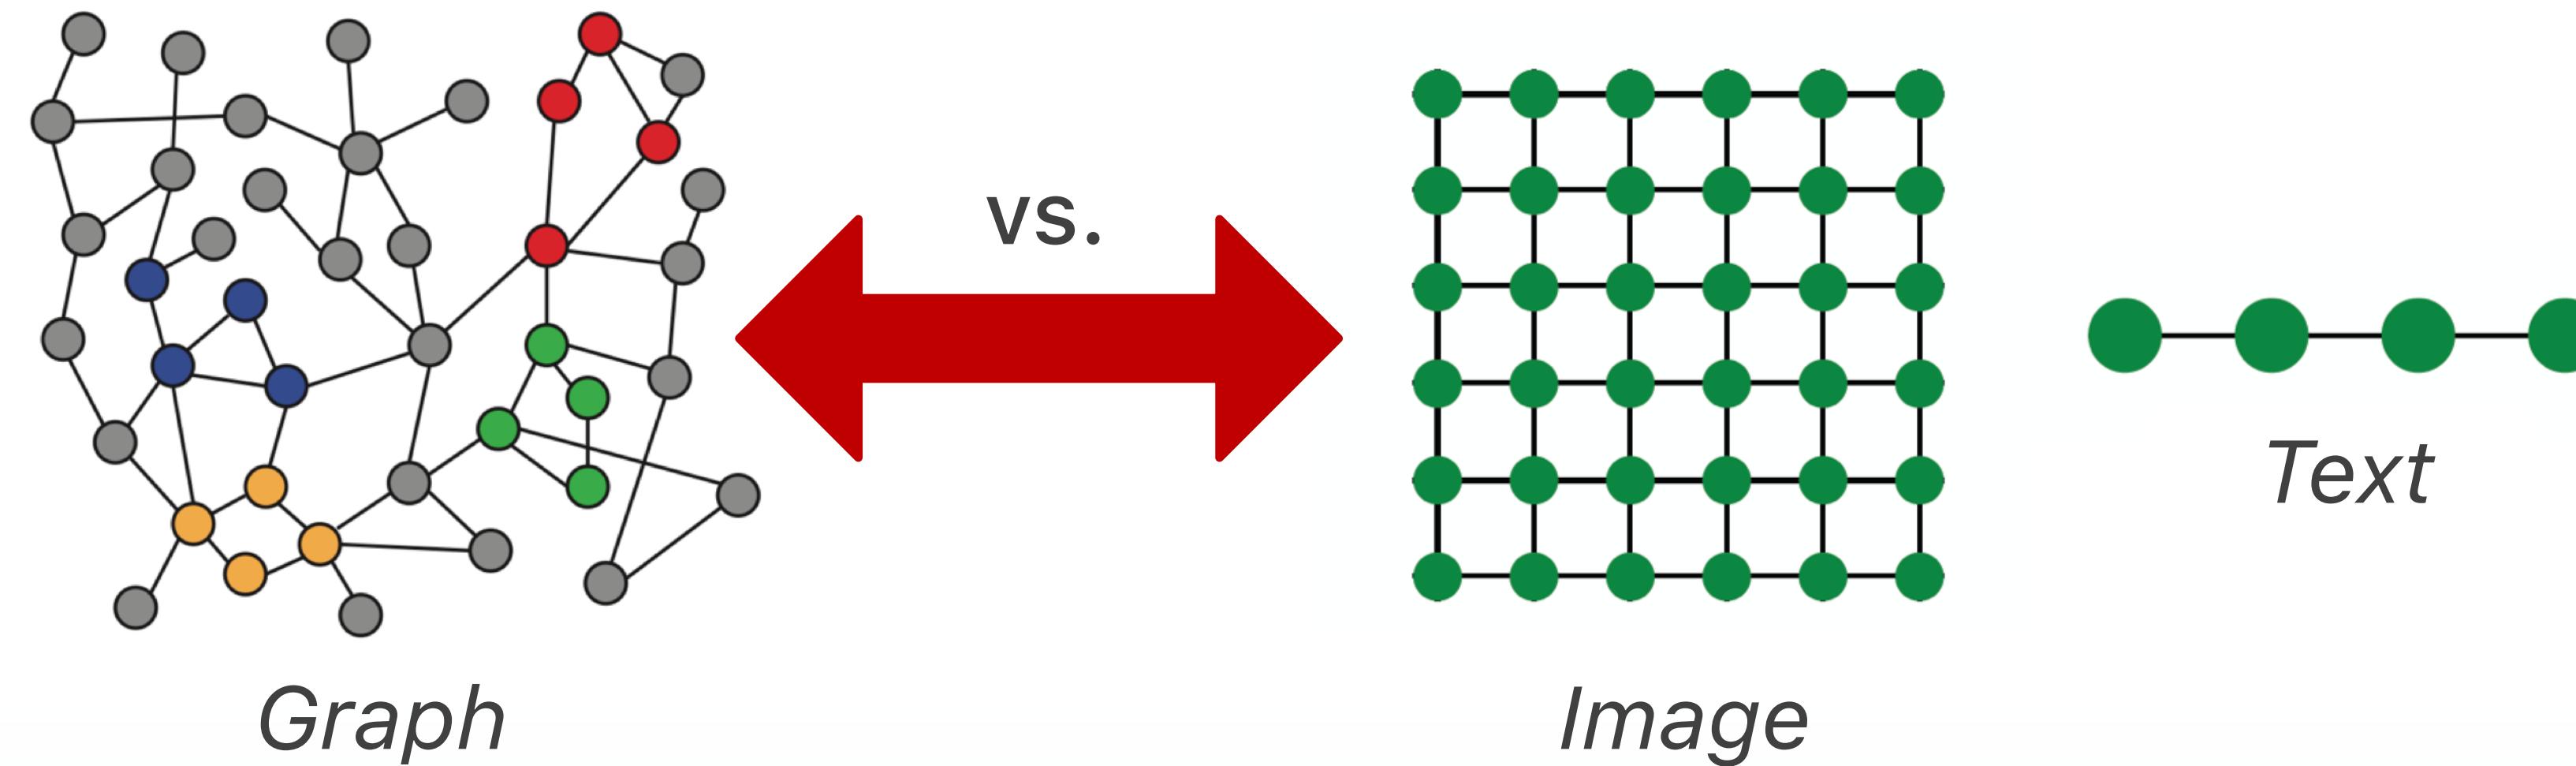


- Arbitrary size and complex topological structure
- No fixed node ordering or reference point



# Motivation

The modern deep learning toolbox is designed for sequences and grids

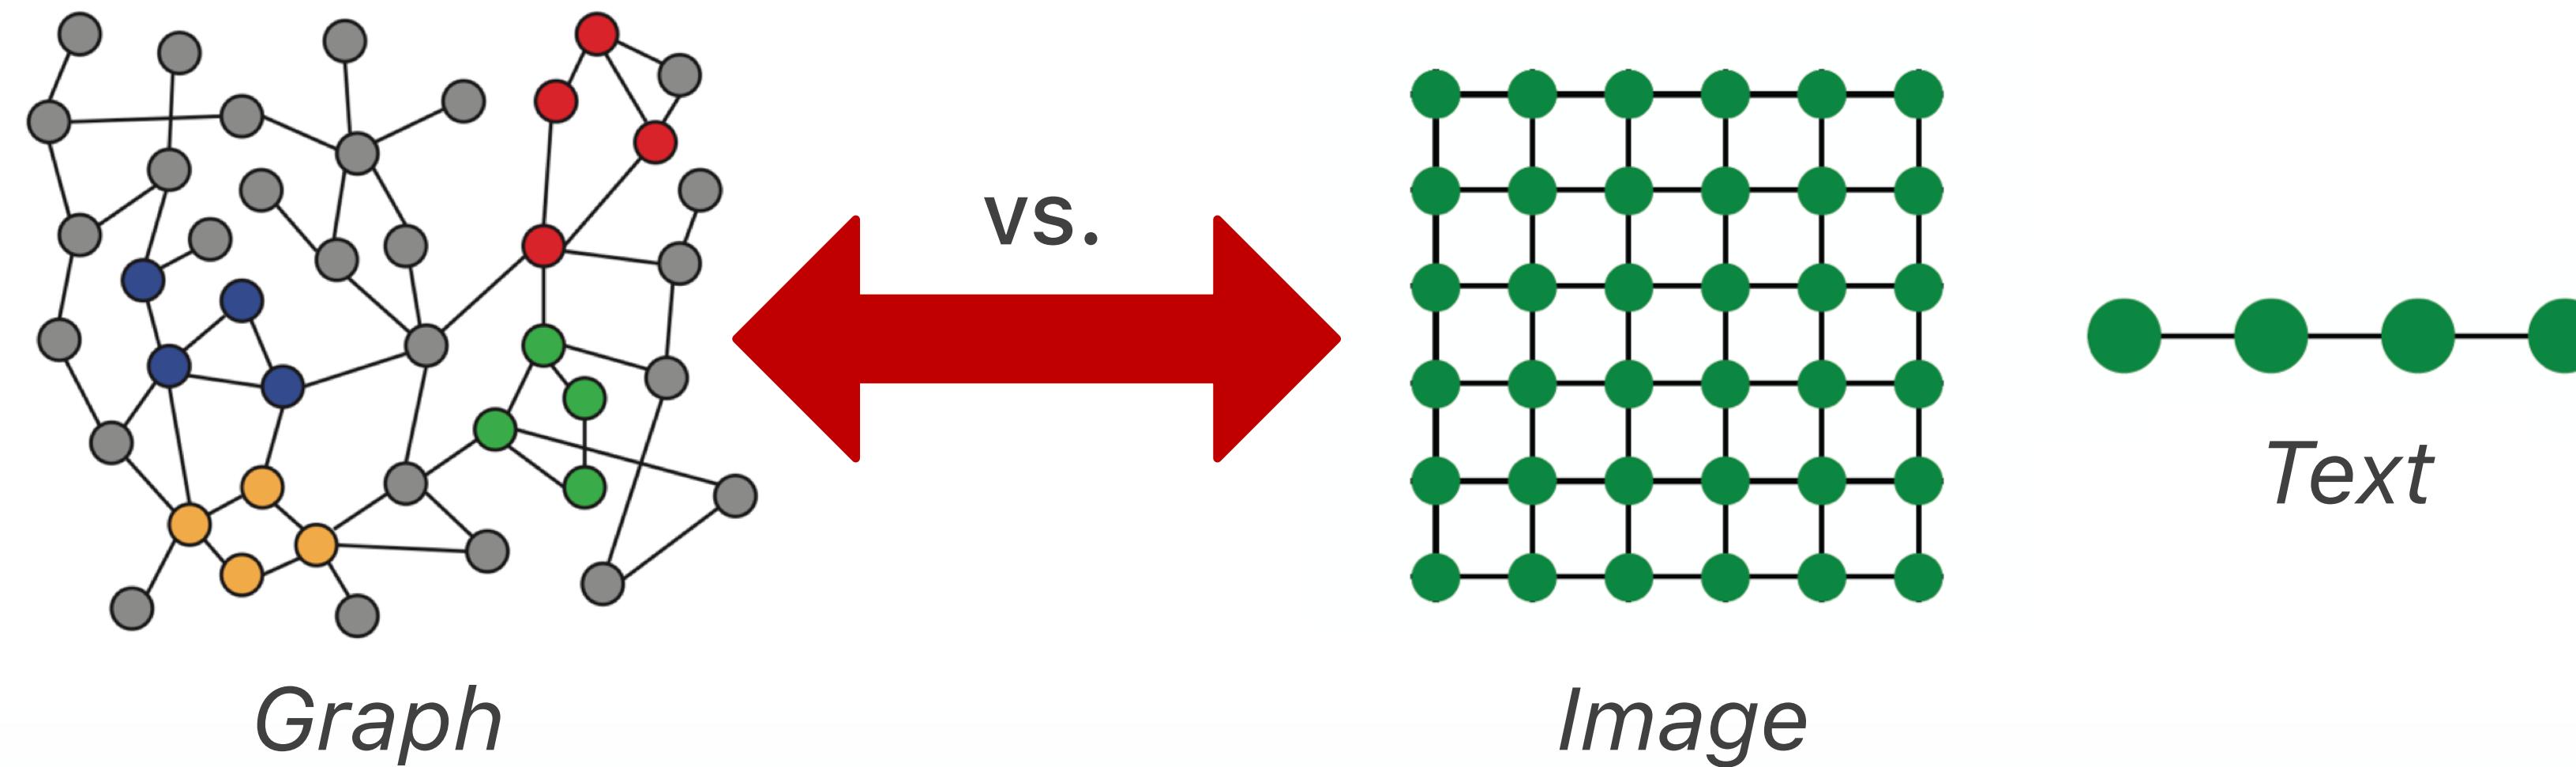


- Arbitrary size and complex topological structure
- No fixed node ordering or reference point
- Often dynamic



# Motivation

The modern deep learning toolbox is designed for sequences and grids

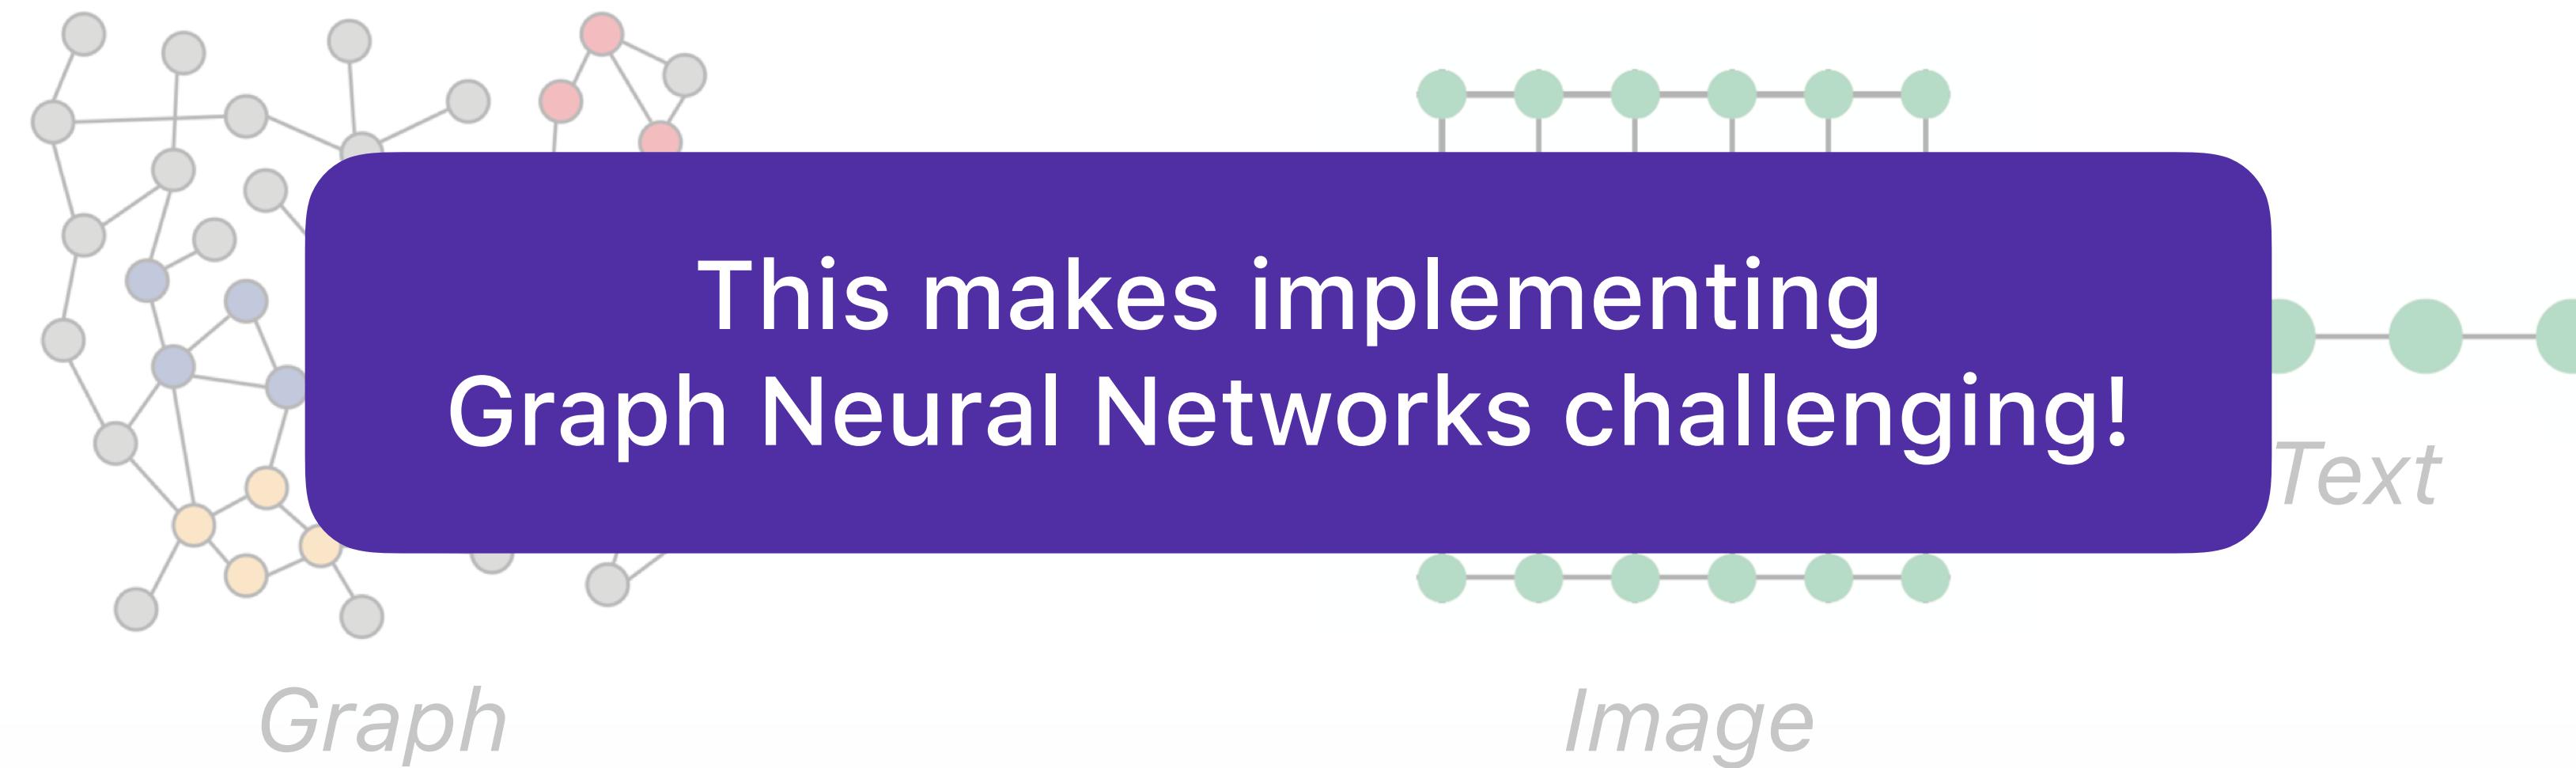


- Arbitrary size and complex topological structure
- No fixed node ordering or reference point
- Often dynamic
- Multimodal node and edge features



# Motivation

The modern deep learning toolbox is designed for sequences and grids



- Arbitrary size and complex topological structure
- No fixed node ordering or reference point
- Often dynamic
- Multimodal node and edge features



# Motivation

 **PyG** (*PyTorch Geometric*) is a  PyTorch library to enable deep learning on graphs, point clouds and manifolds



# Motivation

 **PyG** (*PyTorch Geometric*) is a  PyTorch library to enable deep learning on graphs, point clouds and manifolds

- simplifies implementing and working with Graph Neural Networks (GNNs)



# Motivation

 **PyG** (*PyTorch Geometric*) is a  PyTorch library to enable deep learning on graphs, point clouds and manifolds

- simplifies implementing and working with Graph Neural Networks (GNNs)
- bundles state-of-the-art GNN architectures and training procedures



# Motivation

 **PyG** (*PyTorch Geometric*) is a  PyTorch library to enable deep learning on graphs, point clouds and manifolds

- simplifies implementing and working with Graph Neural Networks (GNNs)
- bundles state-of-the-art GNN architectures and training procedures
- achieves high GPU throughput on highly sparse data of varying size



# Motivation

 **PyG** (*PyTorch Geometric*) is a  PyTorch library to enable deep learning on graphs, point clouds and manifolds

- simplifies implementing and working with Graph Neural Networks (GNNs)
- bundles state-of-the-art GNN architectures and training procedures
- achieves high GPU throughput on highly sparse data of varying size
- suited for both academia and industry



# Design Principles

 **PyG** is  PyTorch-on-the-rocks:



# Design Principles

 **PyG** is  PyTorch-on-the-rocks:

-  **PyG** is framework-specific
  - allows us to make use of recently released features right away:  
`TorchScript` for deployment, `torch.fx` for model transformation*



# Design Principles

 **PyG** is  PyTorch-on-the-rocks:

-  **PyG** is framework-specific
  - allows us to make use of recently released features right away:  
`TorchScript` for deployment, `torch.fx` for model transformation*
-  **PyG** keeps design principles close to vanilla  PyTorch
  - If you are familiar with PyTorch, you already know most about  PyG*



# Design Principles

 **PyG** is  PyTorch-on-the-rocks:

-  **PyG** is framework-specific
  - allows us to make use of recently released features right away:  
`TorchScript` for deployment, `torch.fx` for model transformation*
-  **PyG** keeps design principles close to vanilla  PyTorch
  - If you are familiar with PyTorch, you already know most about  PyG*
-  **PyG** fits nicely into the  PyTorch ecosystem
  - Scaling up models via  PyTorch Lightning*
  - Explaining models via  Captum*



# Design Principles

 PyG is  PyTorch-on-the-rocks:



```
from torch.nn import Conv2d

class CNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = Conv2d(3, 64)
        self.conv2 = Conv2d(64, 64)

    def forward(self, input):
        h = self.conv1(input)
        h = h.relu()
        h = self.conv2(h)
        return h
```



# Design Principles

 PyG is  PyTorch-on-the-rocks:



```
from torch.nn import Conv2d

class CNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = Conv2d(3, 64)
        self.conv2 = Conv2d(64, 64)

    def forward(self, input):
        h = self.conv1(input)
        h = h.relu()
        h = self.conv2(h)
        return h
```



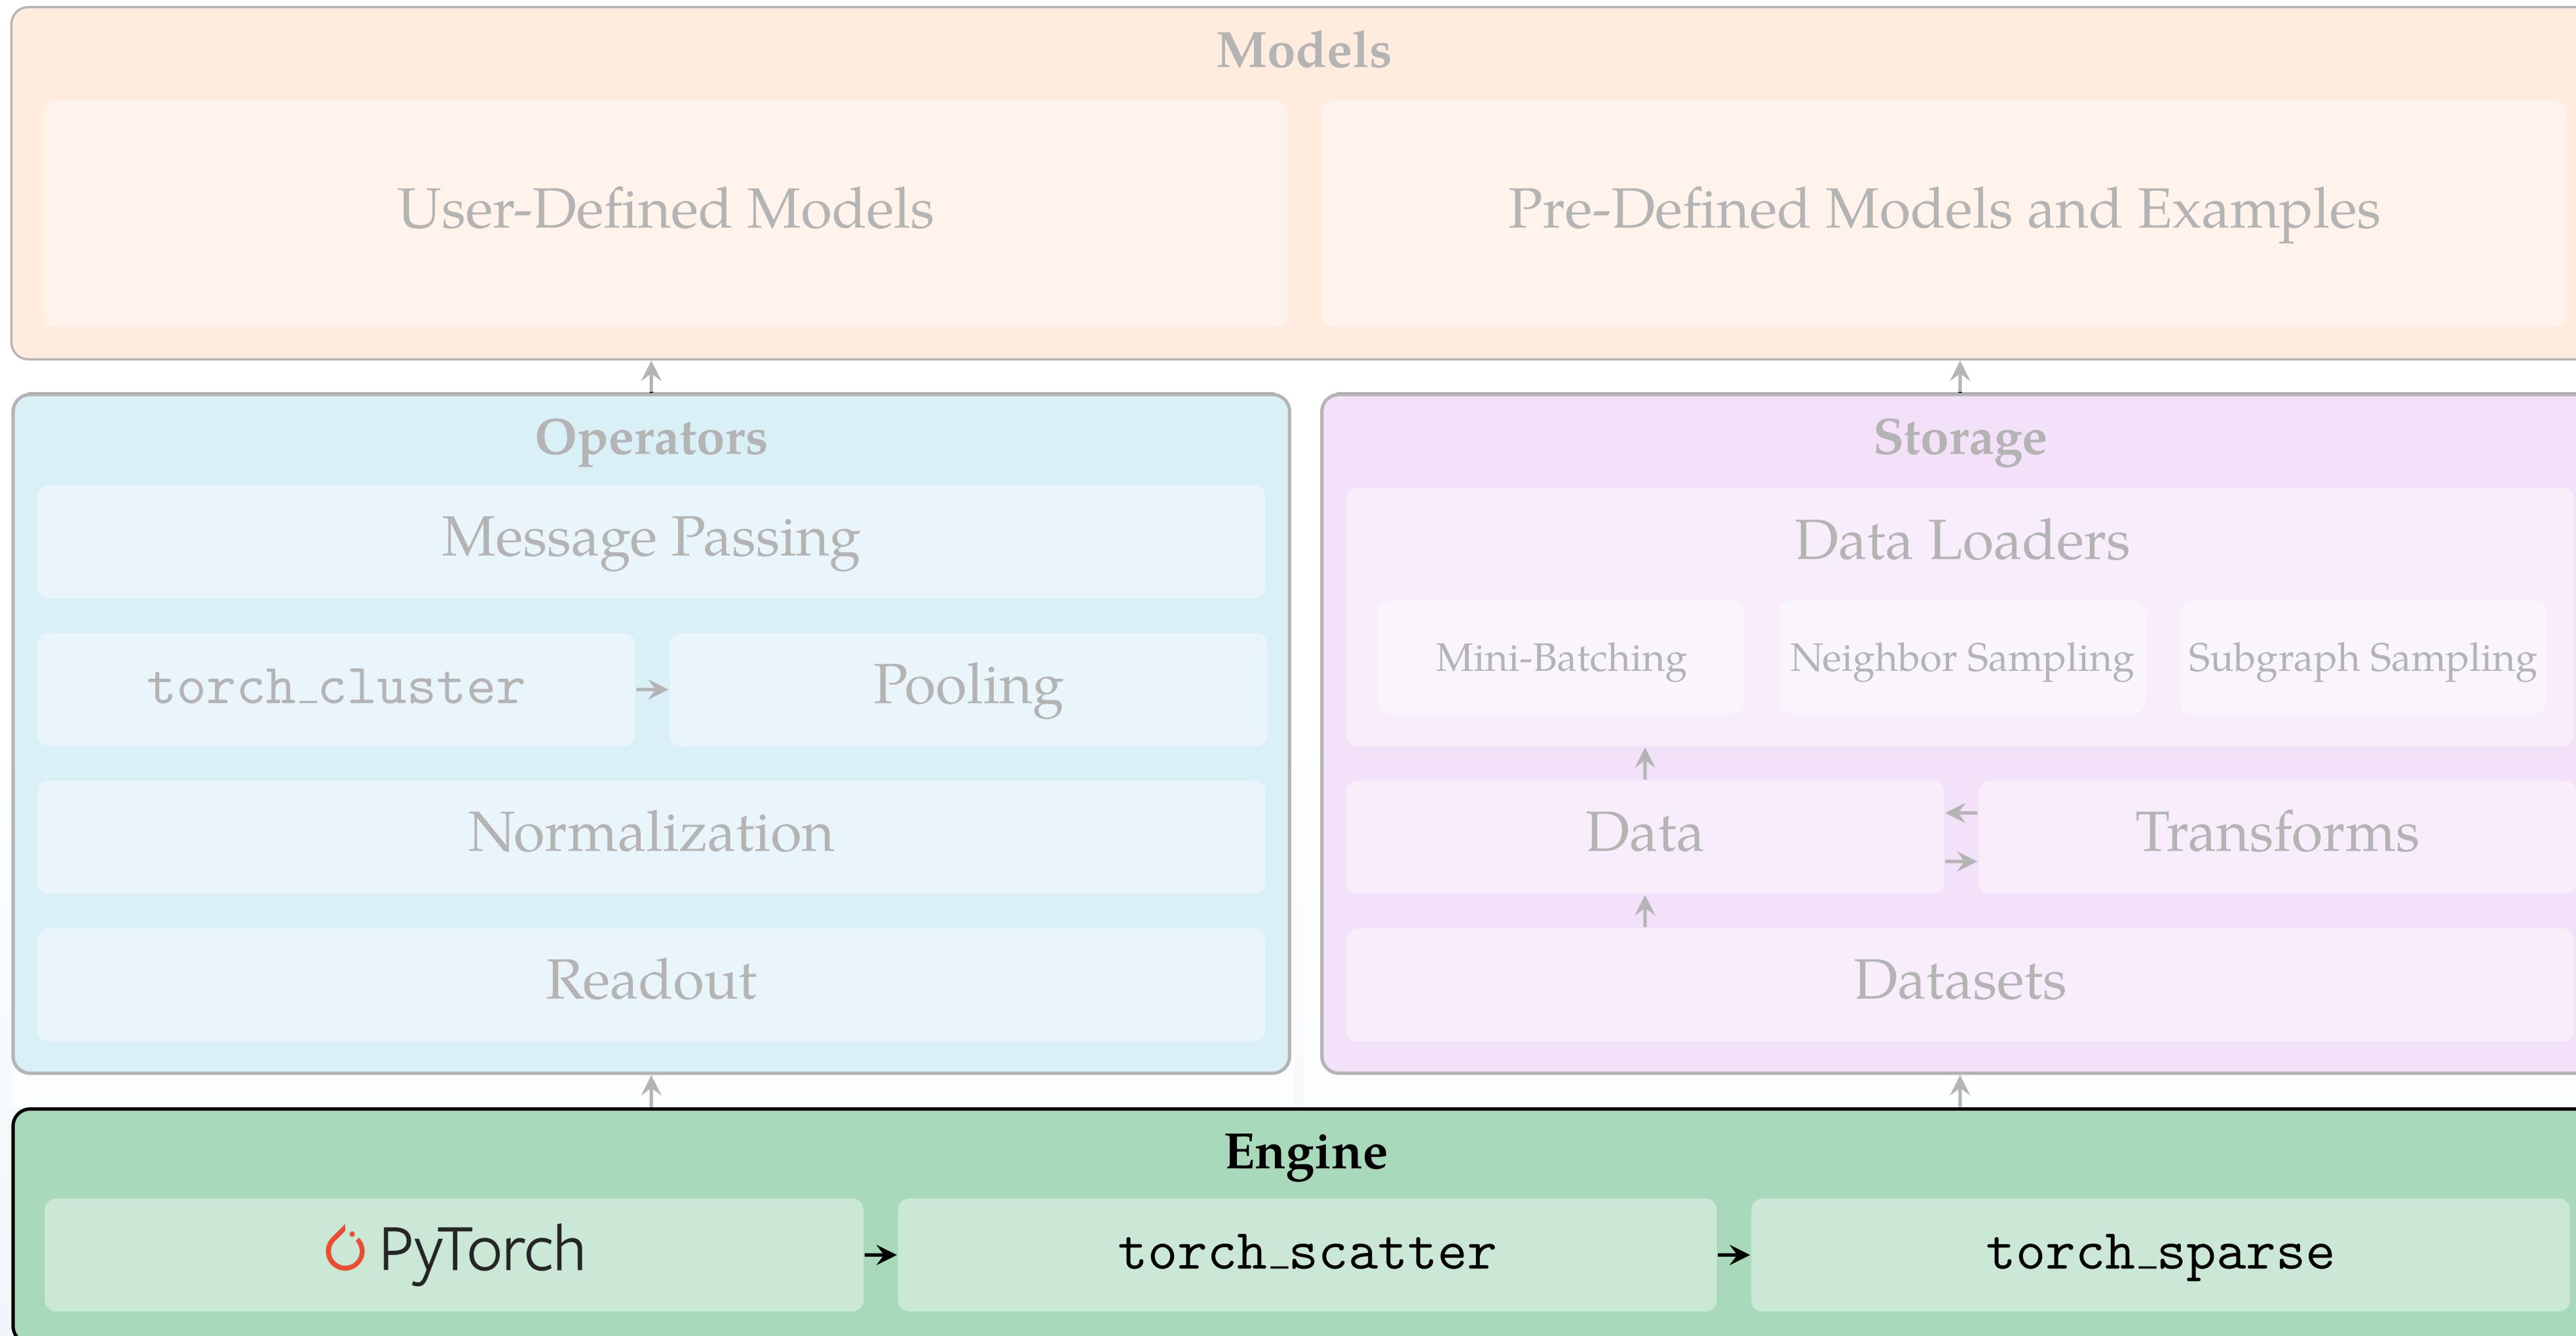
```
from torch_geometric.nn import GCNConv

class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = GCNConv(3, 64)
        self.conv2 = GCNConv(64, 64)

    def forward(self, input, edge_index):
        h = self.conv1(input, edge_index)
        h = h.relu()
        h = self.conv2(h, edge_index)
        return h
```

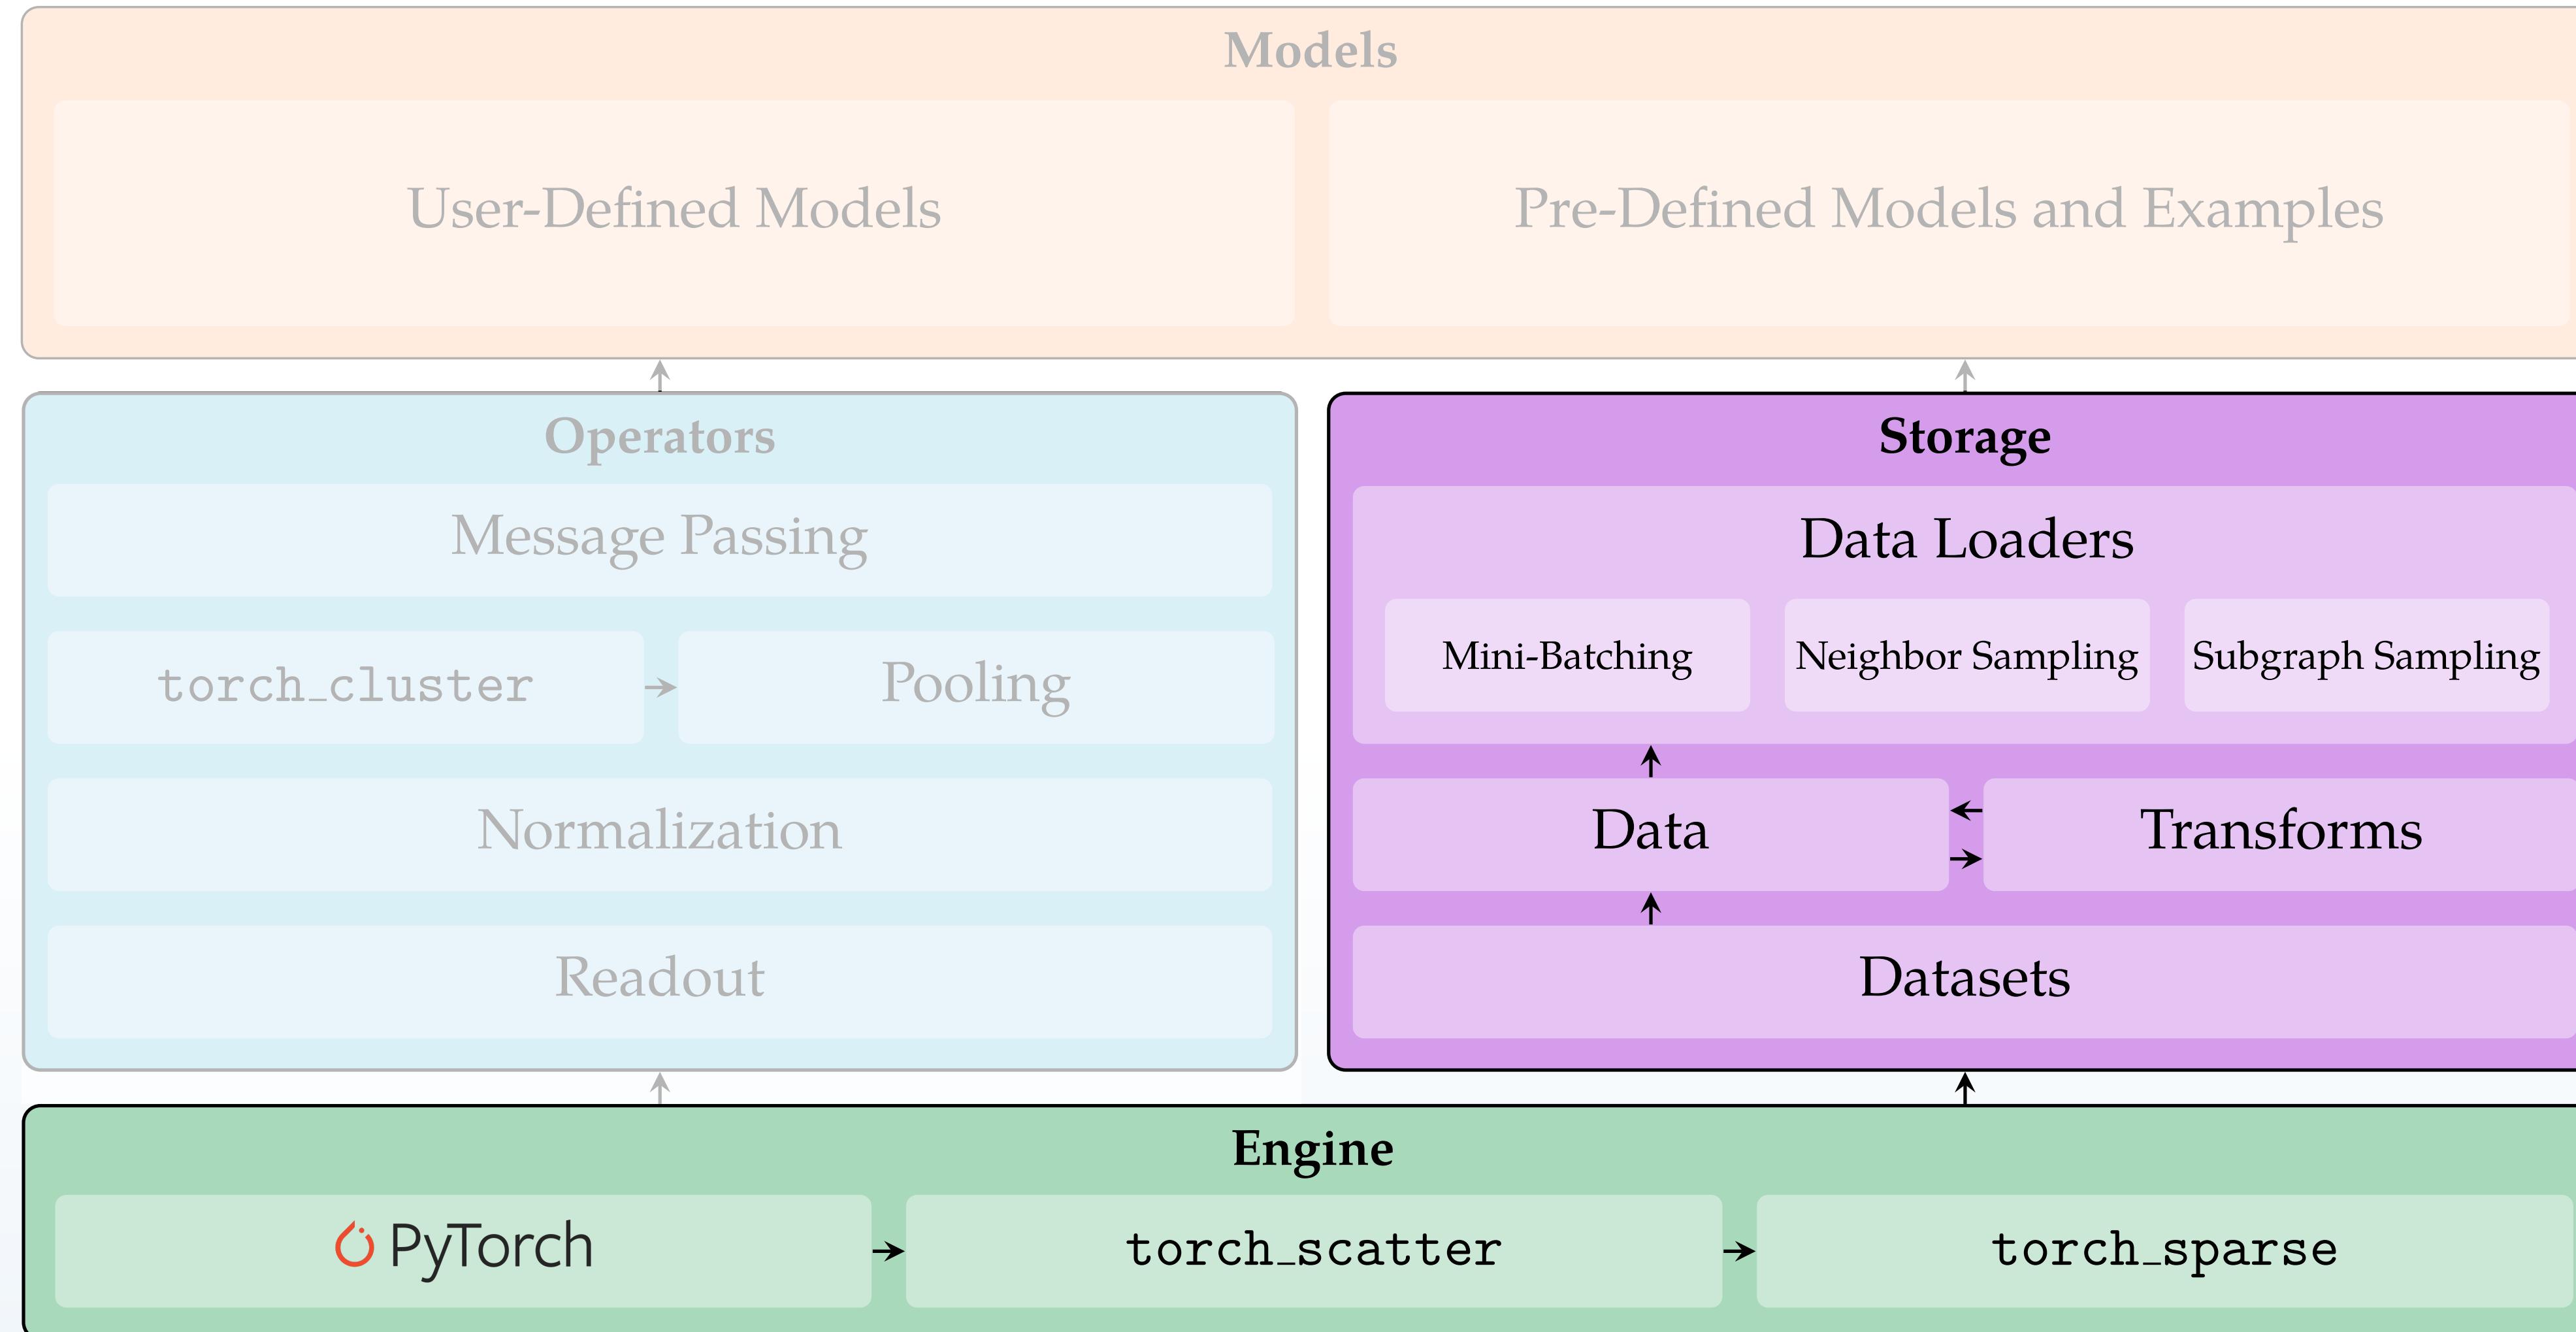


# Design Principles



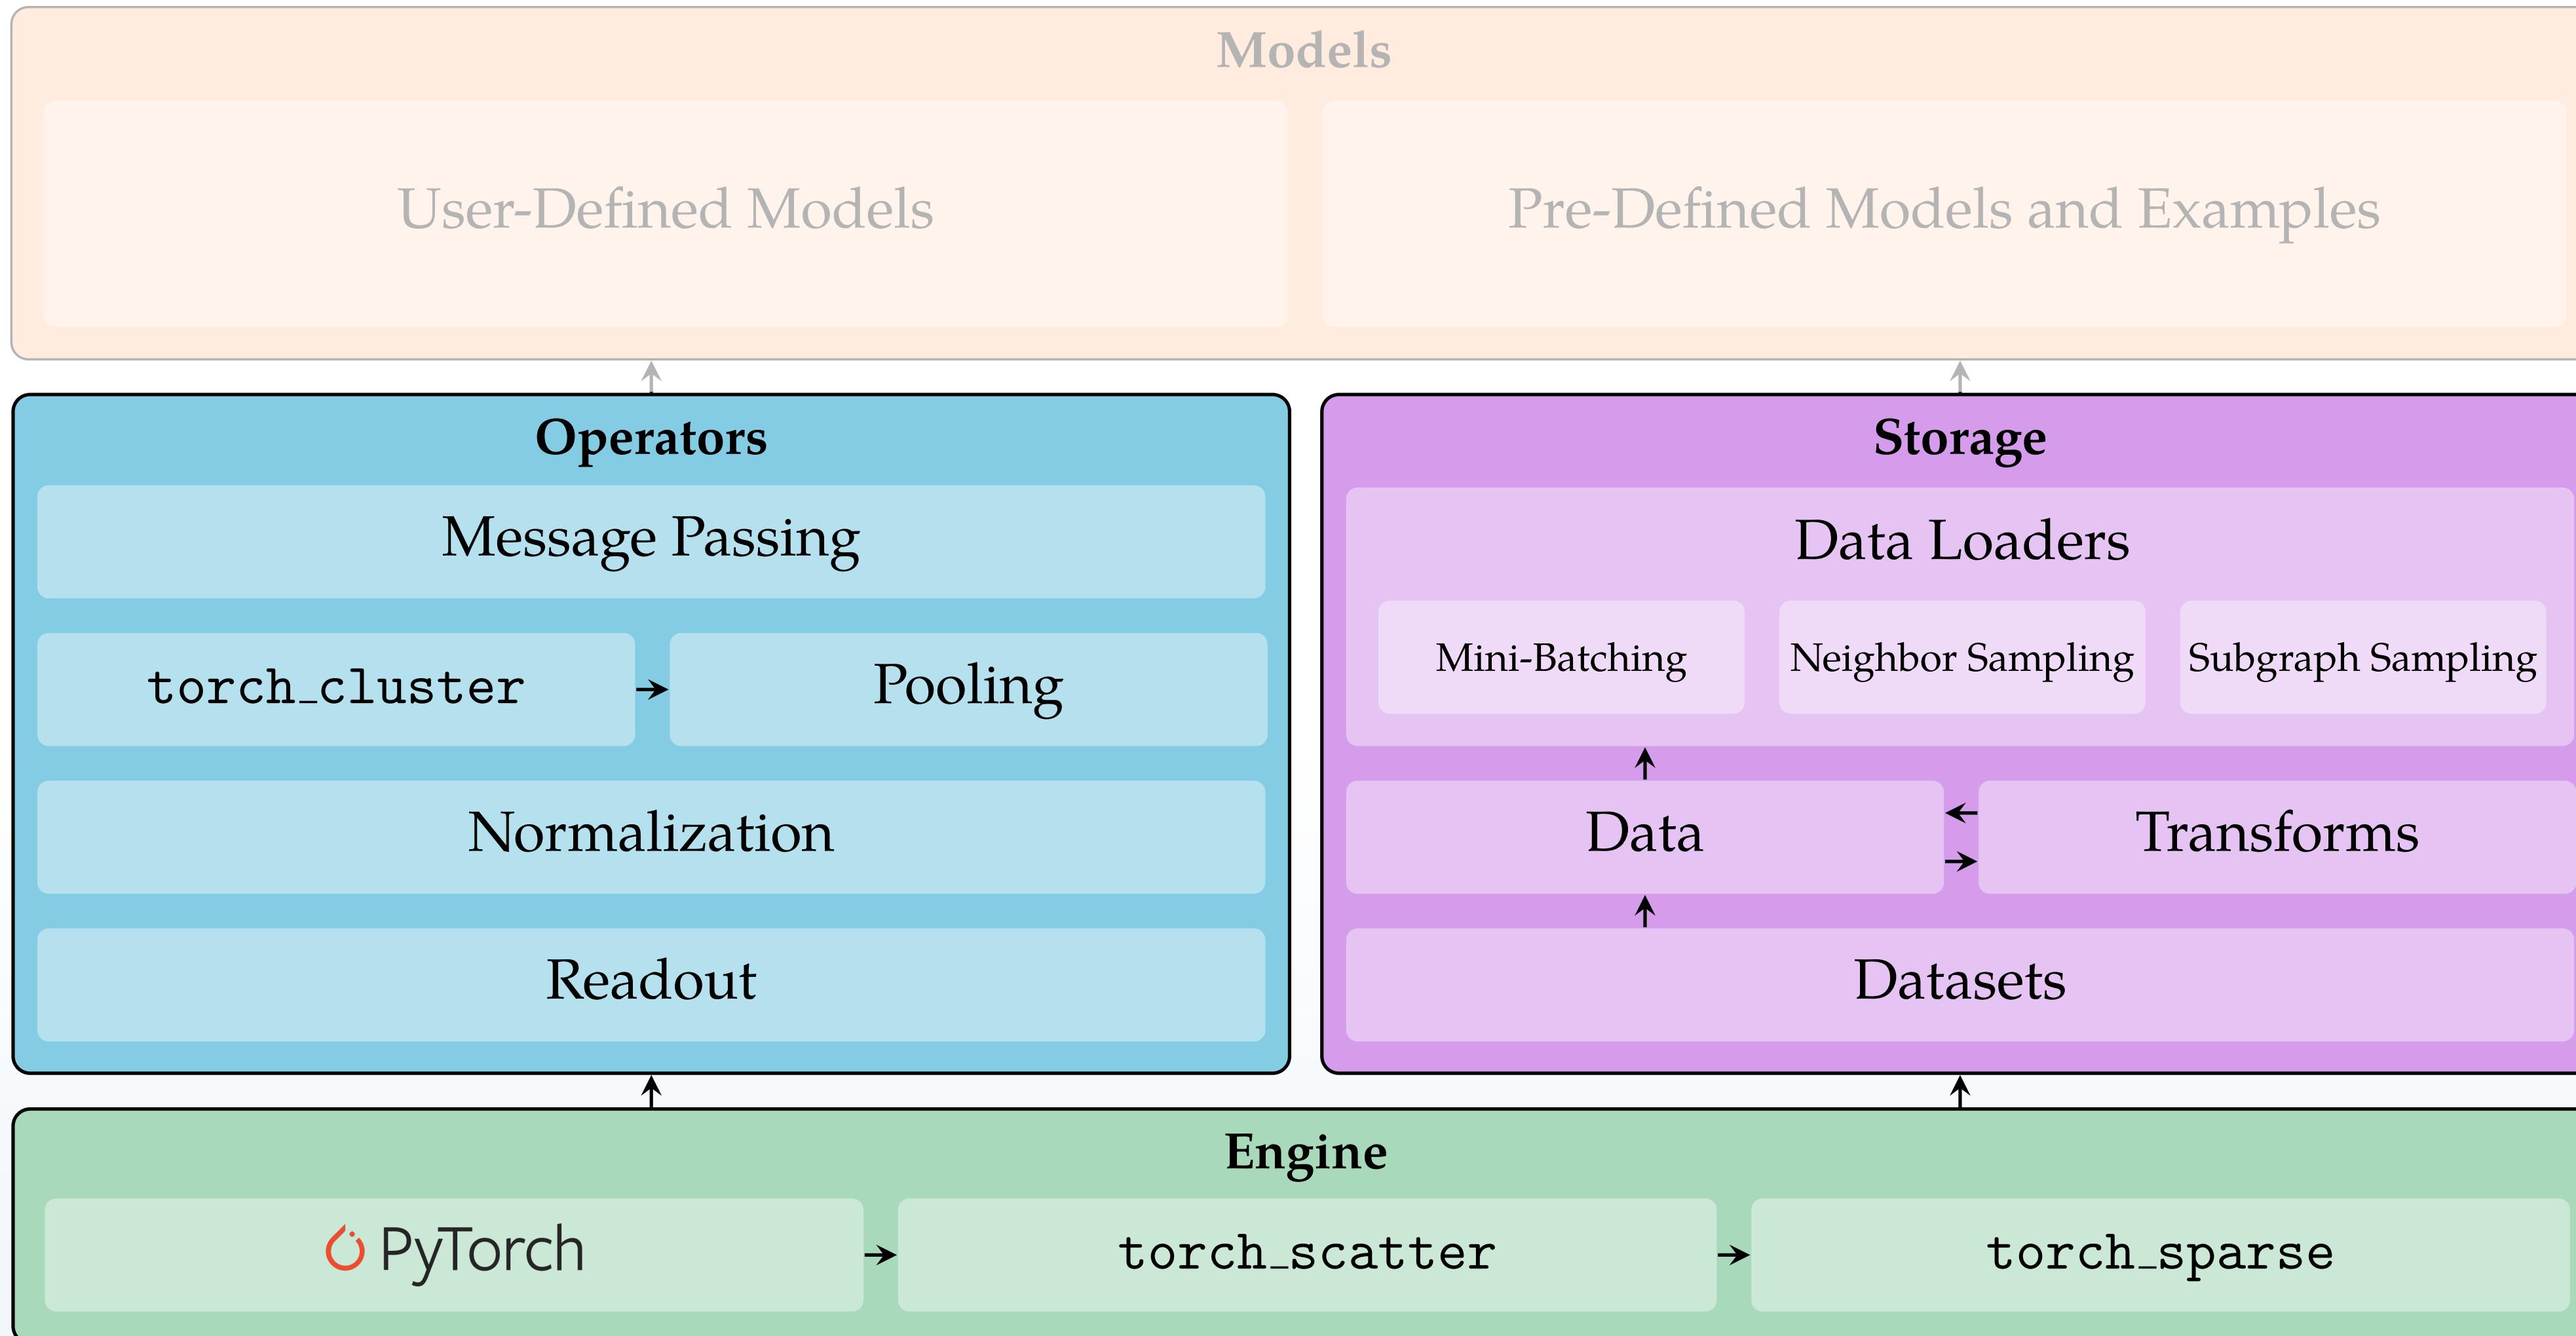


# Design Principles



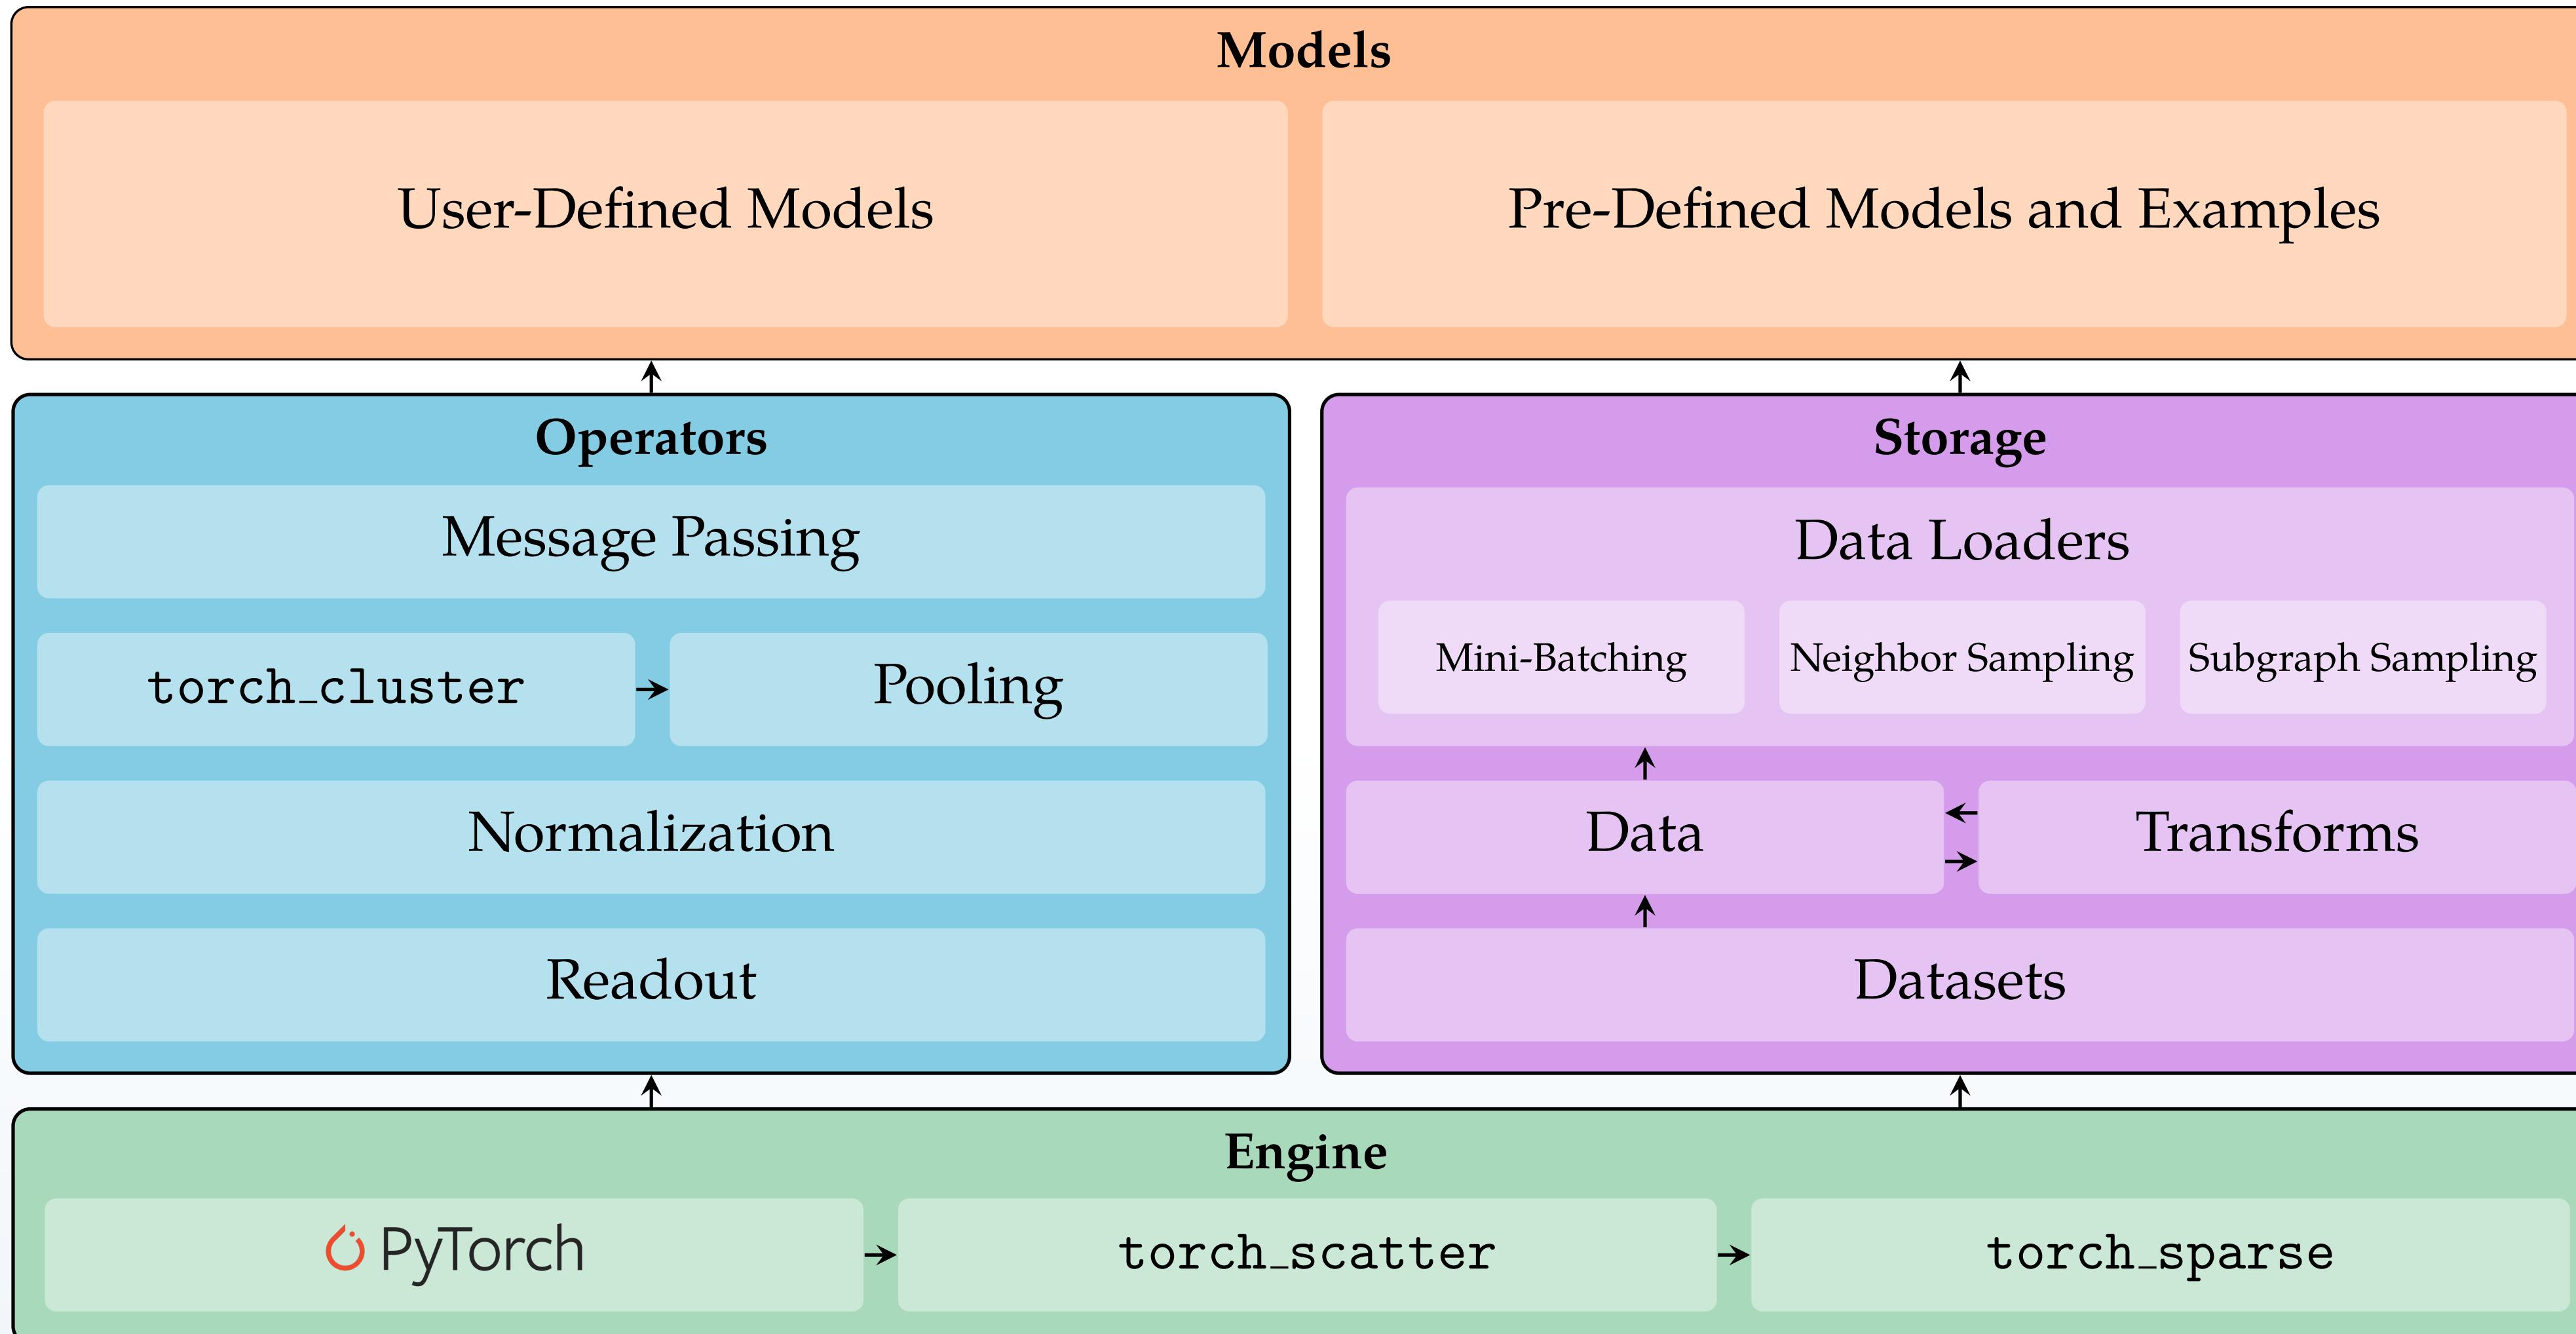


# Design Principles





# Design Principles

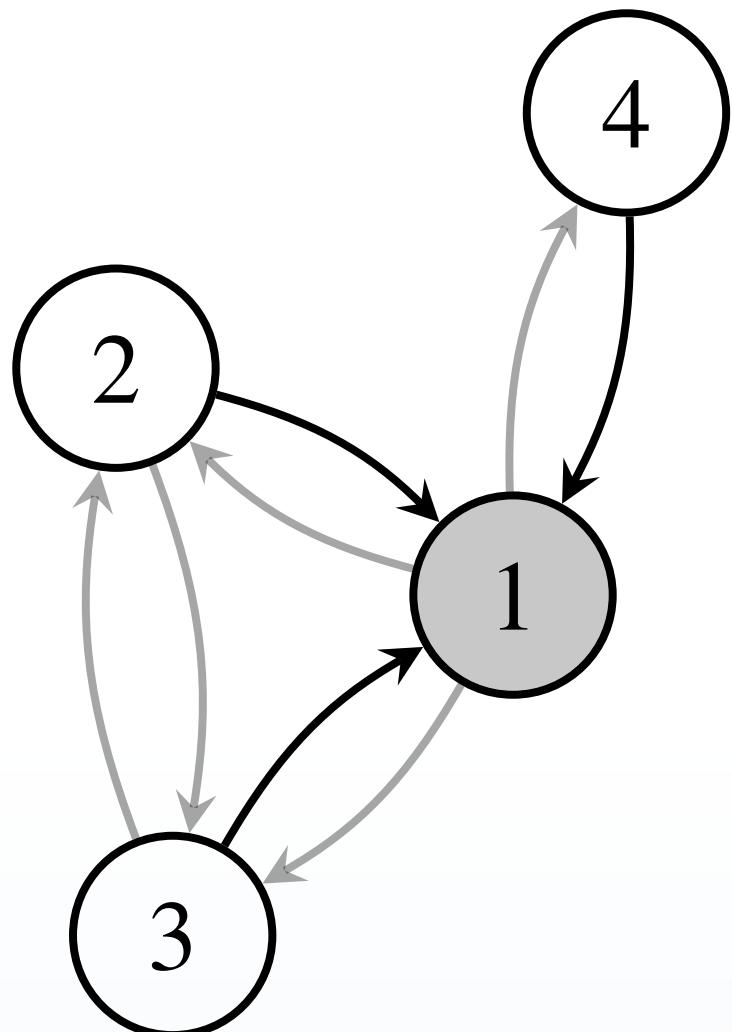




# Message Passing Graph Neural Networks

Given a *sparse* graph  $\mathcal{G} = (\mathbf{H}^{(0)}, (\mathbf{I}, \mathbf{E}))$  with

- **input node features**  $\mathbf{H}^{(0)} \in \mathbb{R}^{|\mathcal{V}| \times C}$
- **edge indices**  $\mathbf{I} \in \{1, \dots, |\mathcal{V}|\}^{2 \times |\mathcal{E}|}$
- *optional* **edge features**  $\mathbf{E} \in \mathbb{R}^{|\mathcal{E}| \times D}$

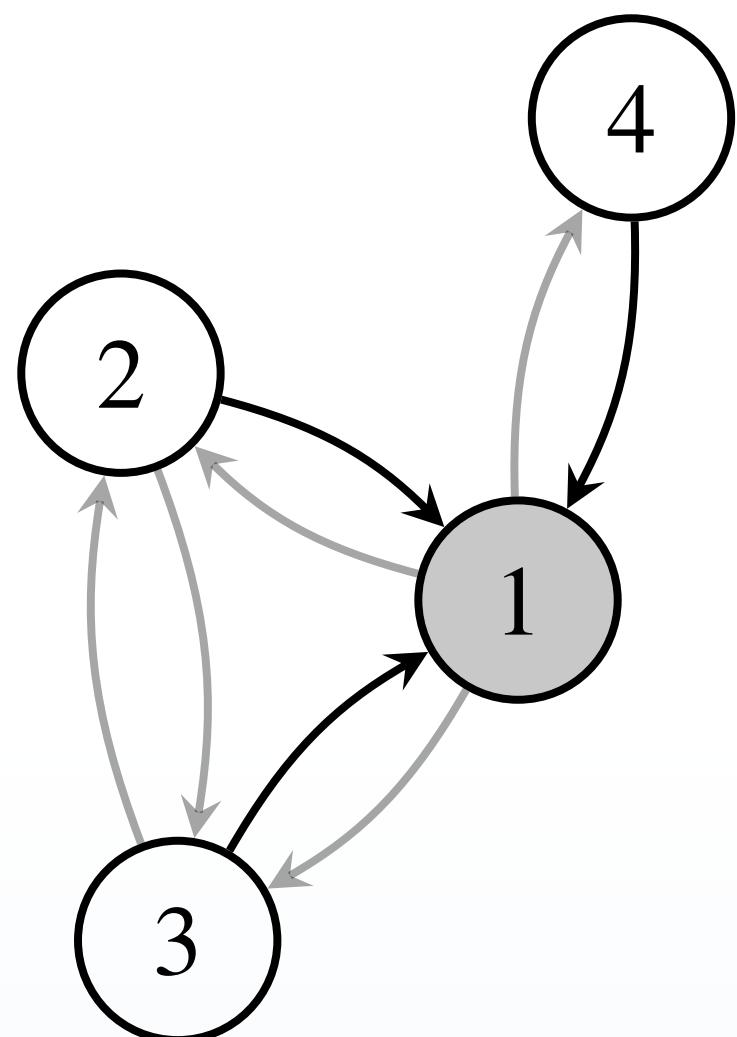




# Message Passing Graph Neural Networks

Given a *sparse graph*  $\mathcal{G} = (\mathbf{H}^{(0)}, (\mathbf{I}, \mathbf{E}))$  with

- **input node features**  $\mathbf{H}^{(0)} \in \mathbb{R}^{|\mathcal{V}| \times C}$
- **edge indices**  $\mathbf{I} \in \{1, \dots, |\mathcal{V}|\}^{2 \times |\mathcal{E}|}$
- *optional edge features*  $\mathbf{E} \in \mathbb{R}^{|\mathcal{E}| \times D}$



## Message Passing Scheme

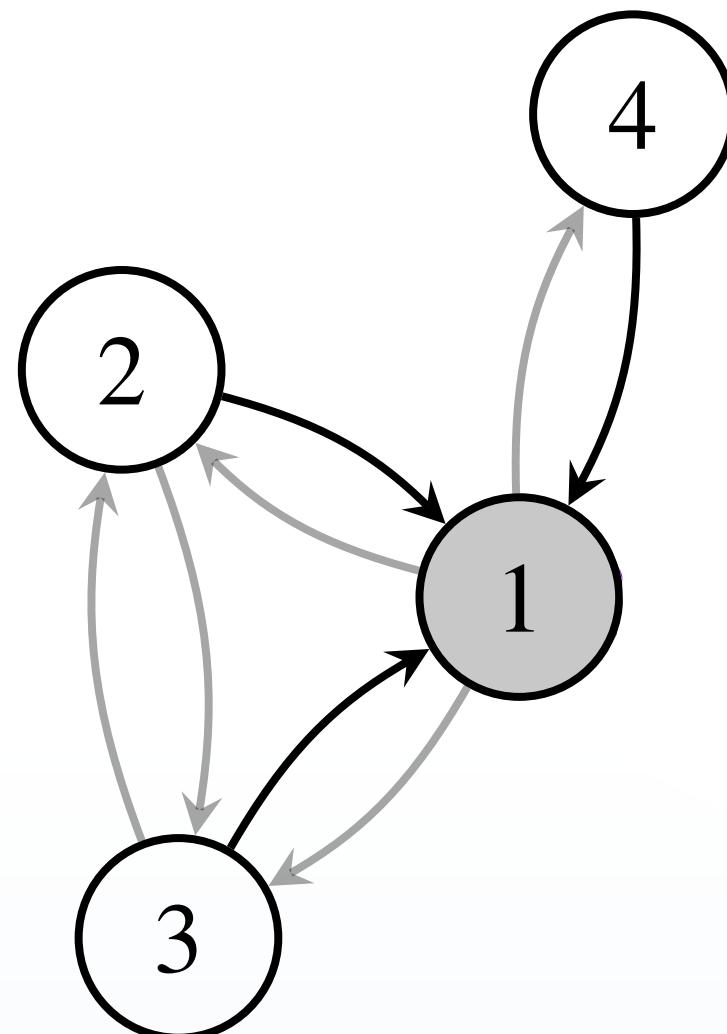
permutation-invariant aggregation operator, e.g., *sum*, *mean* or *max*

$$\mathbf{h}_i^{(\ell+1)} = \text{UPDATE}_{\theta} \left( \mathbf{h}_i^{(\ell)}, \downarrow \bigoplus_{j \in \mathcal{N}(i)} \text{MESSAGE}_{\theta} \left( \mathbf{h}_j^{(\ell)}, \mathbf{h}_i^{(\ell)}, \mathbf{e}_{j,i} \right) \right)$$



# Message Passing Graph Neural Networks

Flexible implementation via *parallelizable* gather and scatter operations  
condensed in a general **MessagePassing** interface

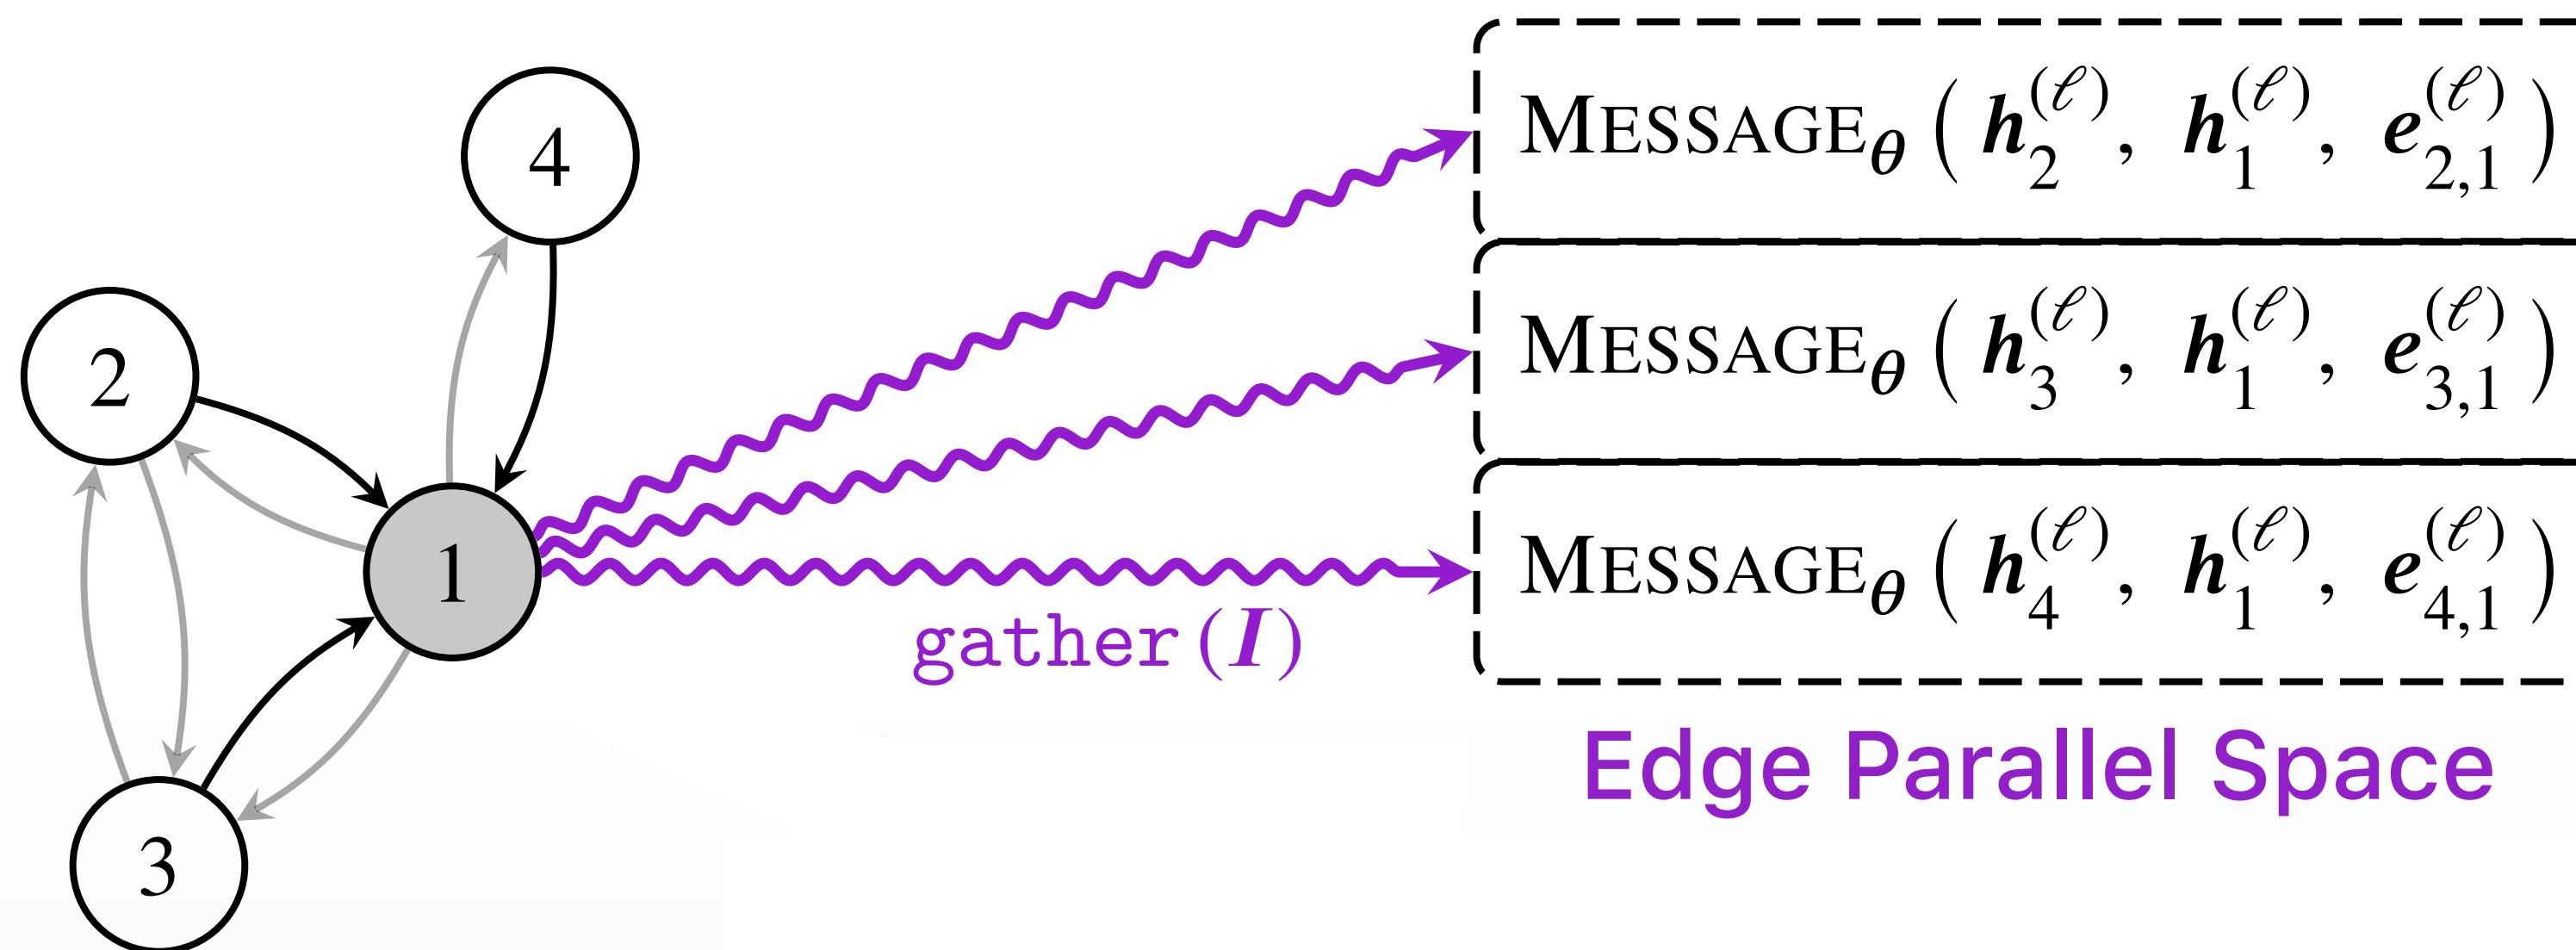


$$\mathbf{h}_i^{(\ell+1)} = \text{UPDATE}_{\theta} \left( \mathbf{h}_i^{(\ell)}, \bigoplus_{j \in \mathcal{N}(i)} \text{MESSAGE}_{\theta} \left( \mathbf{h}_j^{(\ell)}, \mathbf{h}_i^{(\ell)}, \mathbf{e}_{j,i} \right) \right)$$



# Message Passing Graph Neural Networks

Flexible implementation via *parallelizable* gather and scatter operations  
condensed in a general **MessagePassing** interface

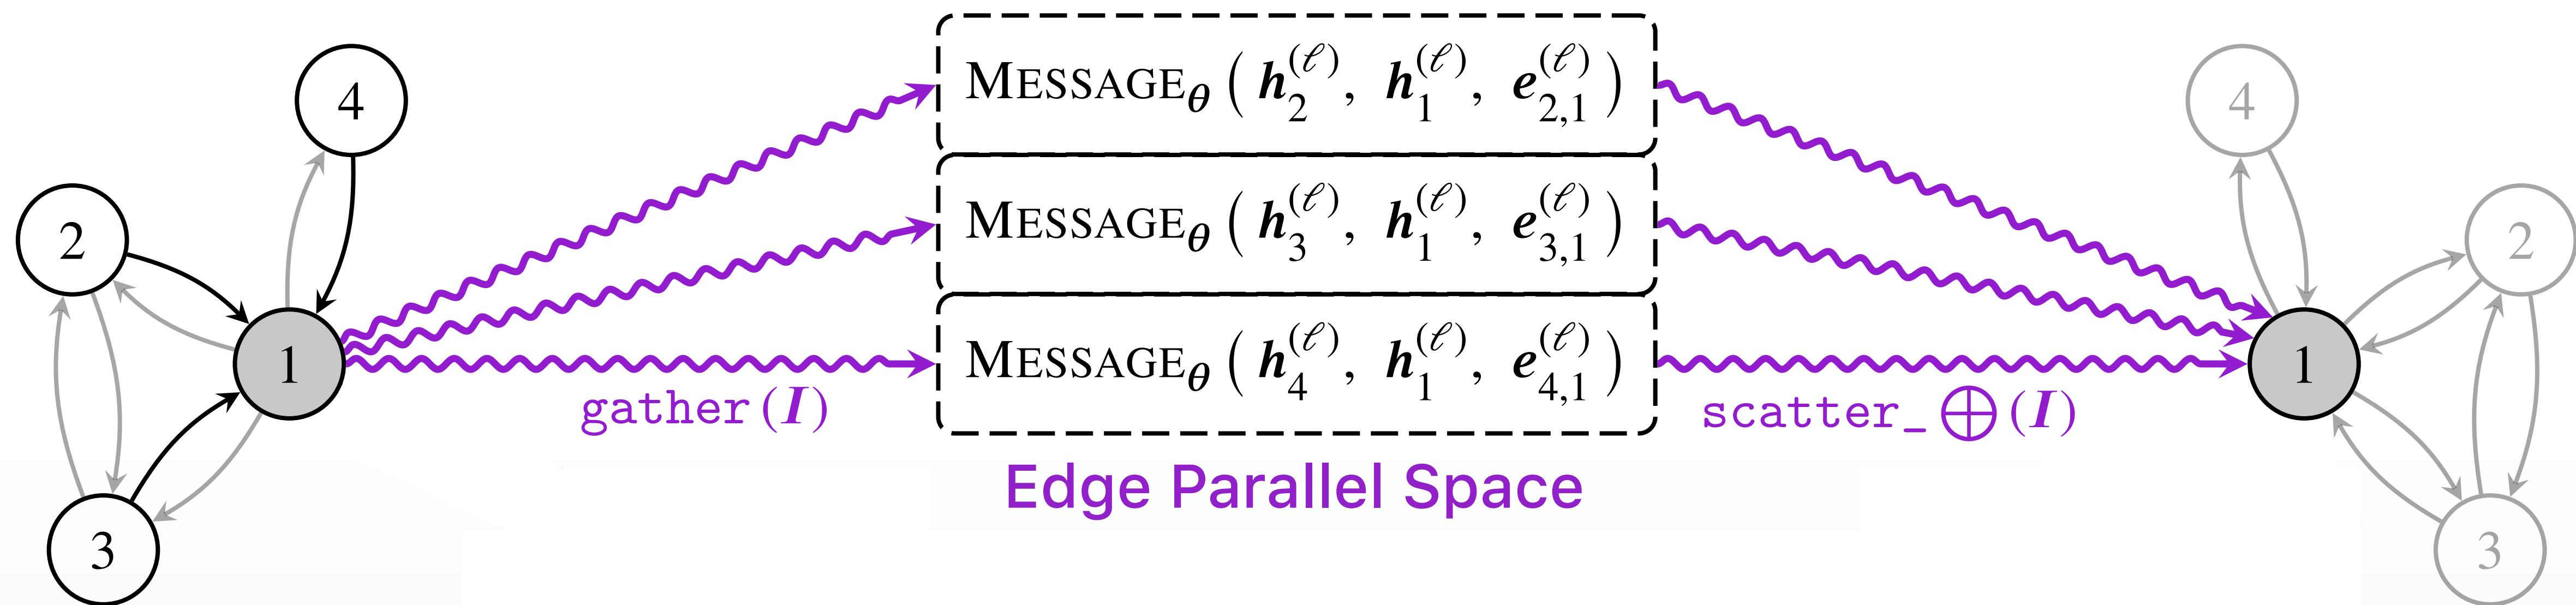


$$\mathbf{h}_i^{(\ell+1)} = \text{UPDATE}_{\theta} \left( \mathbf{h}_i^{(\ell)}, \bigoplus_{j \in \mathcal{N}(i)} \text{MESSAGE}_{\theta} \left( \mathbf{h}_j^{(\ell)}, \mathbf{h}_i^{(\ell)}, \mathbf{e}_{j,i} \right) \right)$$



# Message Passing Graph Neural Networks

Flexible implementation via *parallelizable* gather and scatter operations  
condensed in a general **MessagePassing** interface

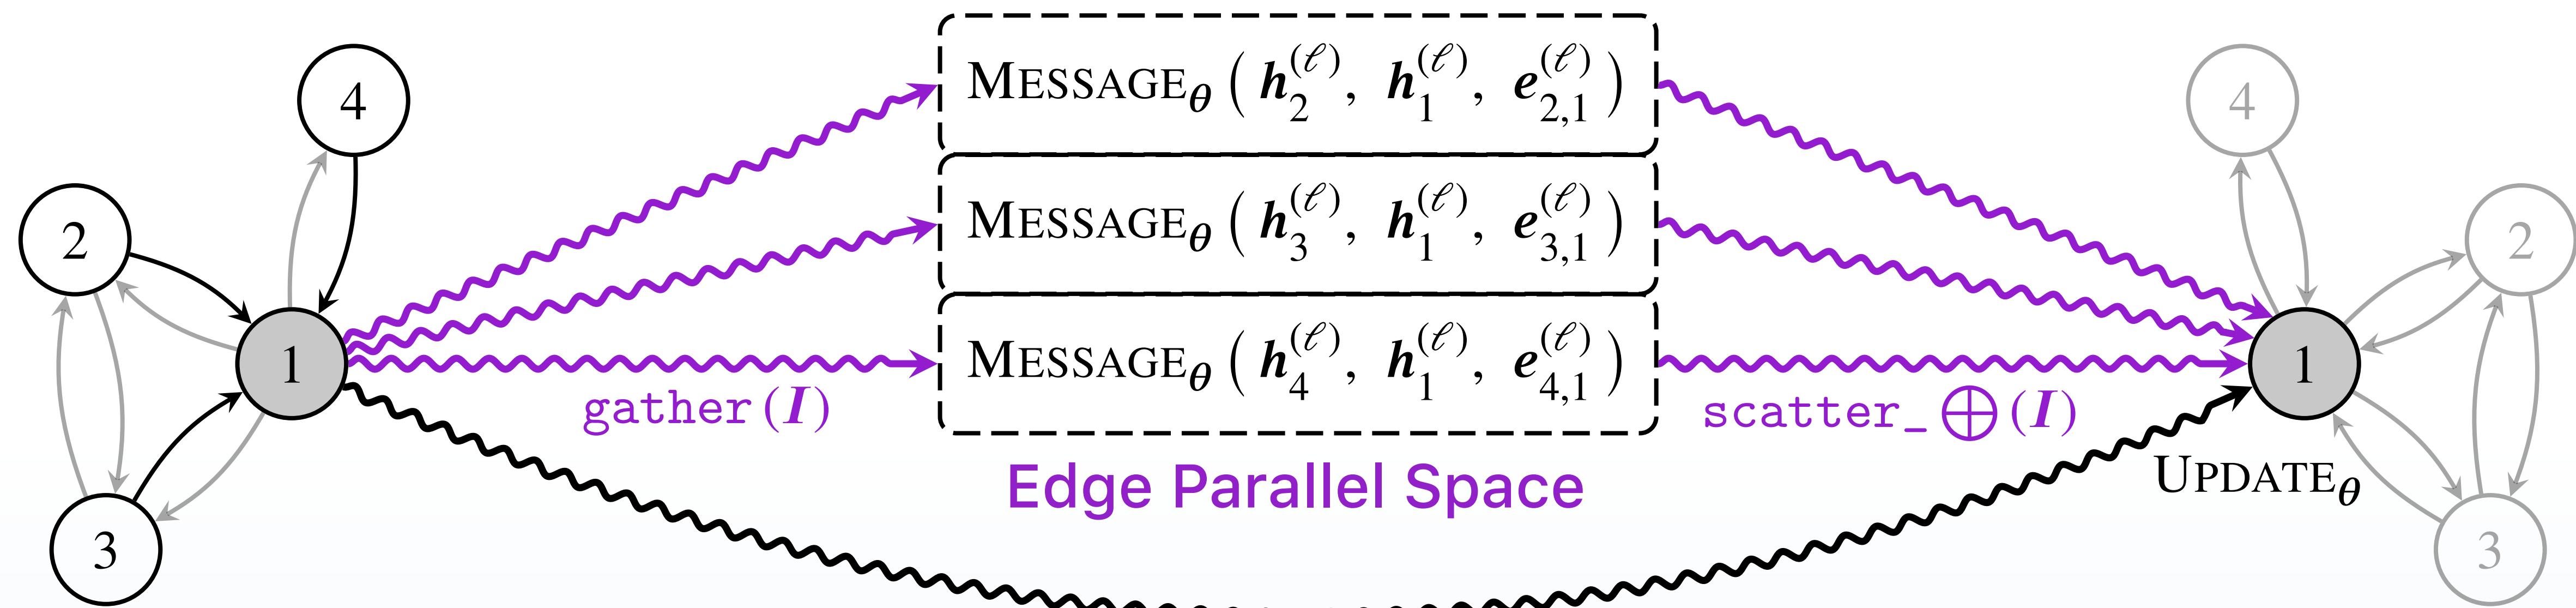


$$\mathbf{h}_i^{(\ell+1)} = \text{UPDATE}_{\theta} \left( \mathbf{h}_i^{(\ell)}, \bigoplus_{j \in \mathcal{N}(i)} \text{MESSAGE}_{\theta} \left( \mathbf{h}_j^{(\ell)}, \mathbf{h}_i^{(\ell)}, \mathbf{e}_{j,i} \right) \right)$$



# Message Passing Graph Neural Networks

Flexible implementation via *parallelizable* gather and scatter operations  
condensed in a general **MessagePassing** interface



$$\mathbf{h}_i^{(\ell+1)} = \text{UPDATE}_{\theta} \left( \mathbf{h}_i^{(\ell)}, \bigoplus_{j \in \mathcal{N}(i)} \text{MESSAGE}_{\theta} \left( \mathbf{h}_j^{(\ell)}, \mathbf{h}_i^{(\ell)}, \mathbf{e}_{j,i} \right) \right)$$



# Message Passing Graph Neural Networks

 **PyG** covers a large number of state-of-the-art GNN layers and architectures, and can easily be extended to fit to a specific use-case



# Message Passing Graph Neural Networks

 **PyG** covers a large number of state-of-the-art GNN layers and architectures, and can easily be extended to fit to a specific use-case

## GNN layers:

Cheby GCN SAGE PointNet MoNet MPNN GAT SuperGAT  
SplineCNN AGNN EdgeCNN S-GCN R-GCN PointCNN  
ARMA APPNP GIN GIN-E CG GatedGCN NMF TAG DNA  
Signed-GCN PPFNet FeaST Hyper-GCN GravNet PDN WL  
ResGatedGCN SparseTransformer DGCNN FeaST EG LeCNN  
PNA GEN GCN2 PAN FiLM SuperGAT FA HGT



# Message Passing Graph Neural Networks

**PyG** covers a large number of state-of-the-art GNN layers and architectures, and can easily be extended to fit to a specific use-case

**GNN layers:**

Cheby GCN SAGE PointNet MoNet MPNN GAT SuperGAT  
SplineCNN AGNN EdgeCNN S-GCN R-GCN PointCNN  
ARMA APPNP GIN GIN-E CG GatedGCN NMF TAG DNA  
Signed-GCN PPFNet FeaST Hyper-GCN GravNet PDN WL  
ResGatedGCN SparseTransformer DGCNN FeaST EG LeCNN  
PNA GEN GCN2 PAN FiLM SuperGAT FA HGT

**GNN models:**

GAE ARGA DGI Node2Vec Graph-UNet AFP C&S LP RECT  
GeniePath SchNet DimeNet MetaPath2Vec RENet TGN SEAL



# Message Passing Graph Neural Networks

Not all GNN operators need to explicitly materialize the edge parallel space

In this case,  **PyG** can fuse message and aggregation computation  
for better memory efficiency:  $\mathcal{O}(|\mathcal{V}|)$

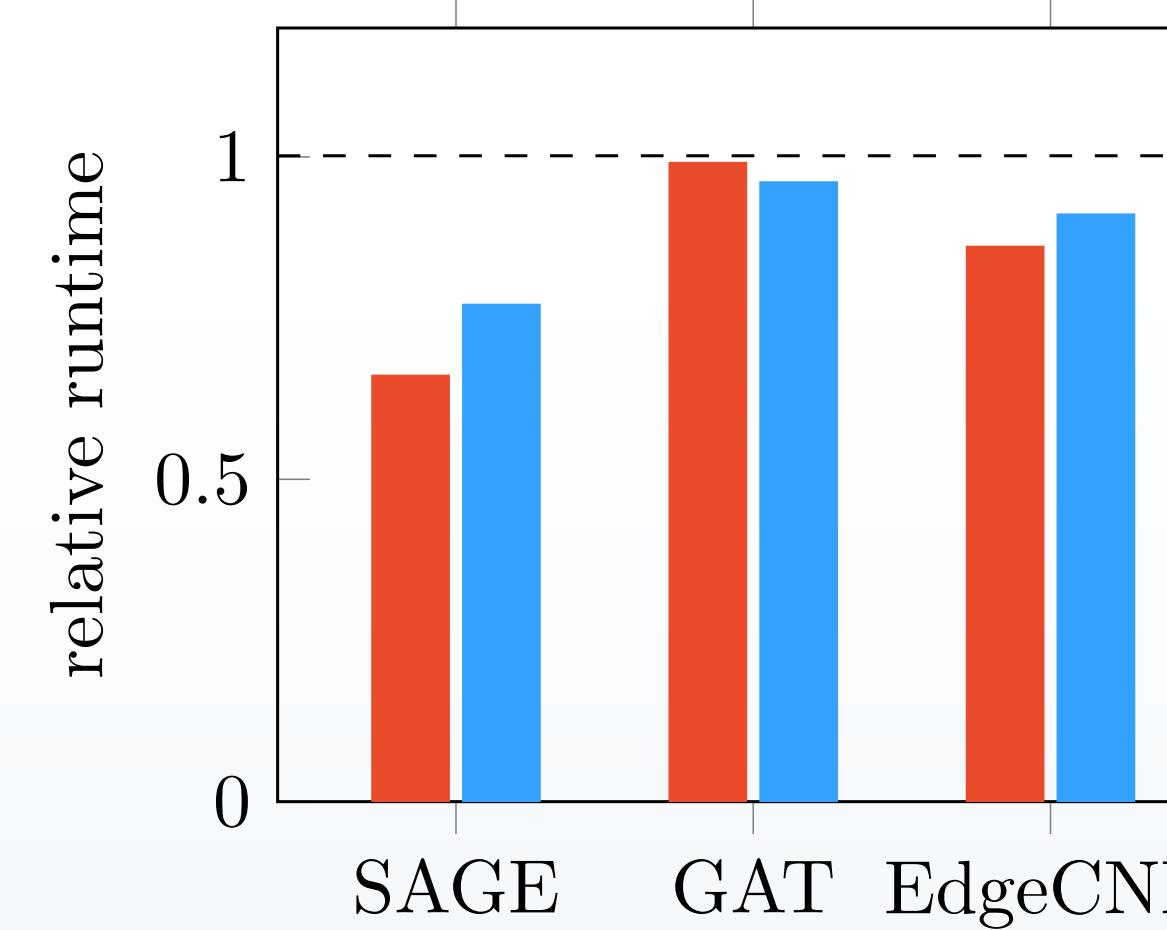


# Message Passing Graph Neural Networks

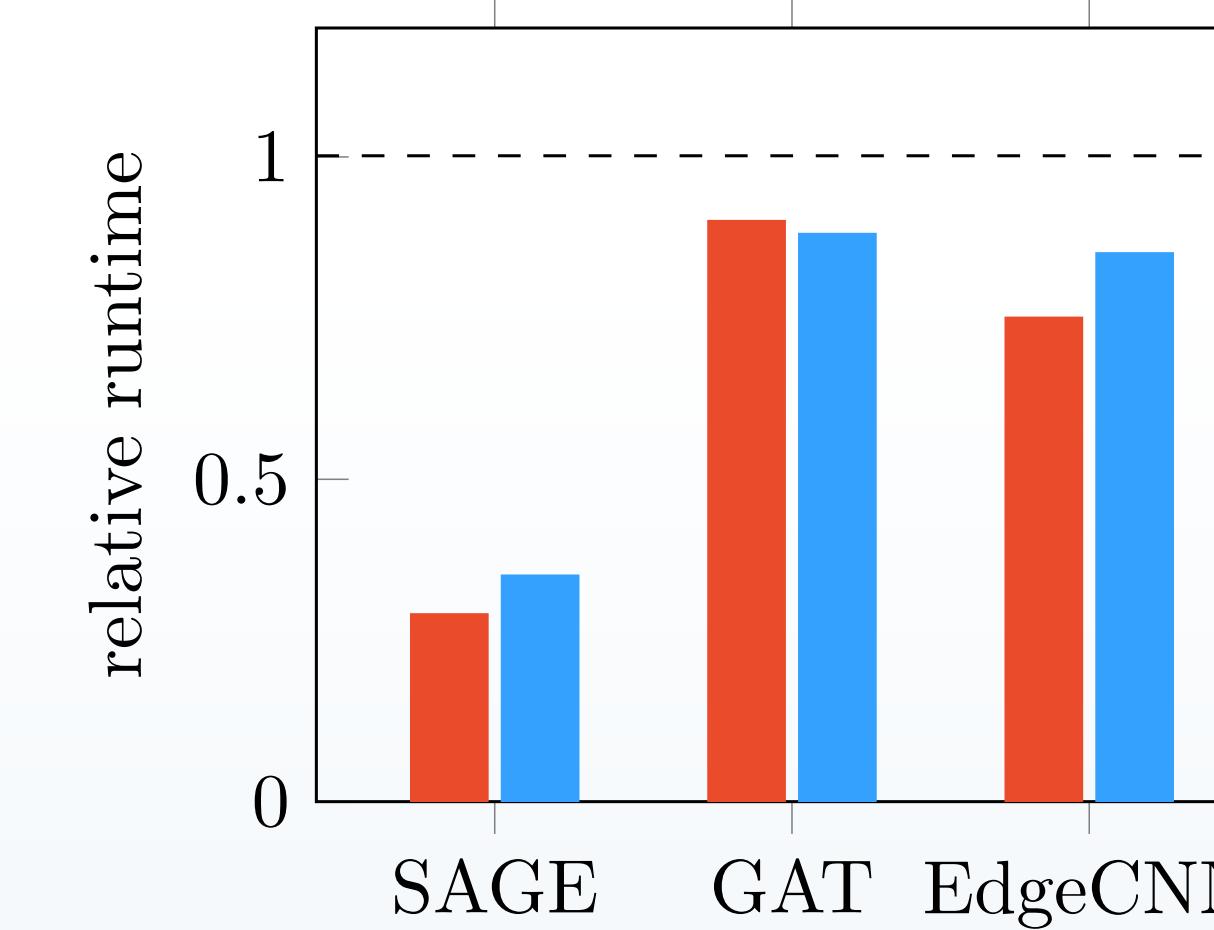
Not all GNN operators need to explicitly materialize the edge parallel space

In this case, PyG can fuse message and aggregation computation  
for better memory efficiency:  $\mathcal{O}(|\mathcal{V}|)$

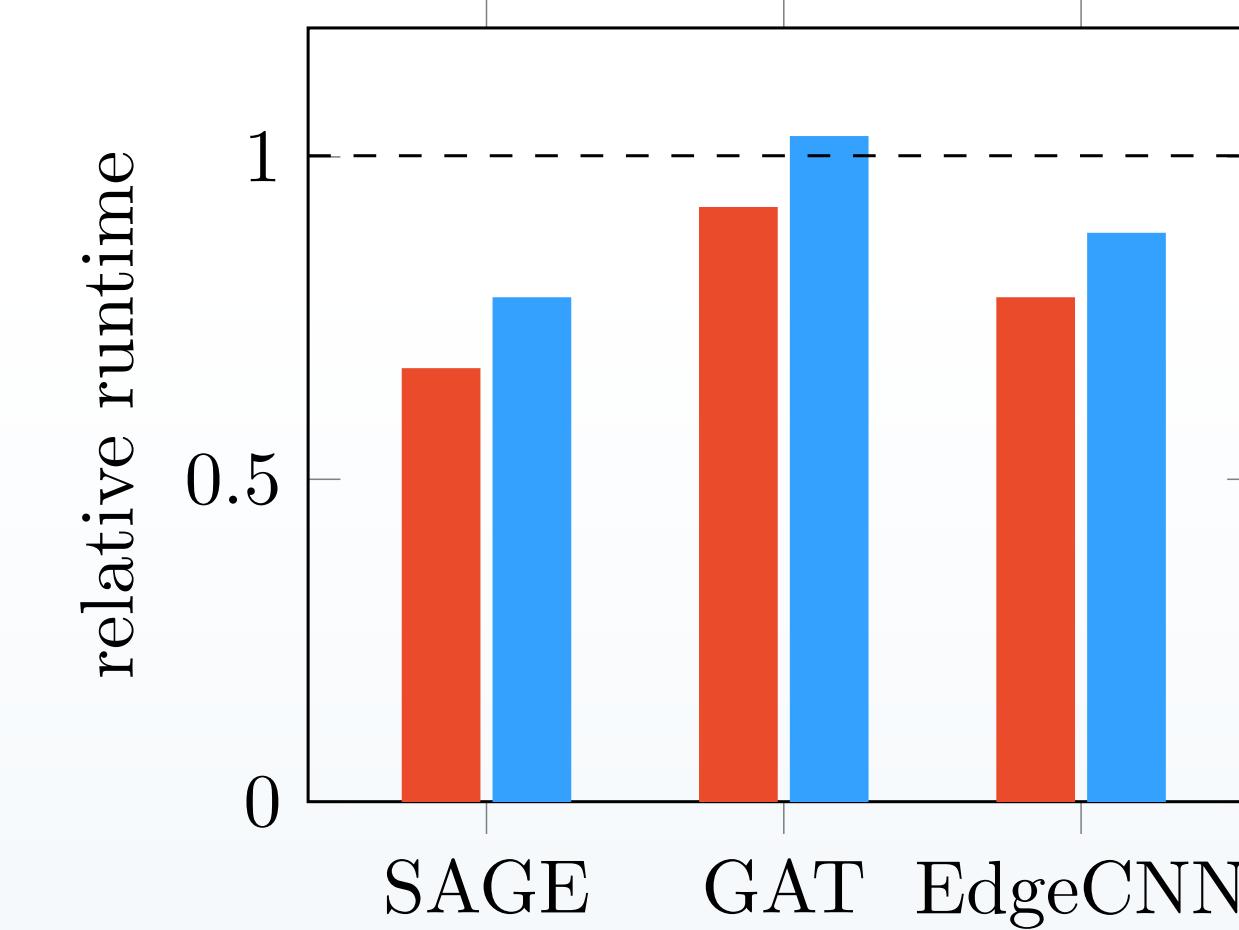
PubMed



PPI



DD



■ Fused

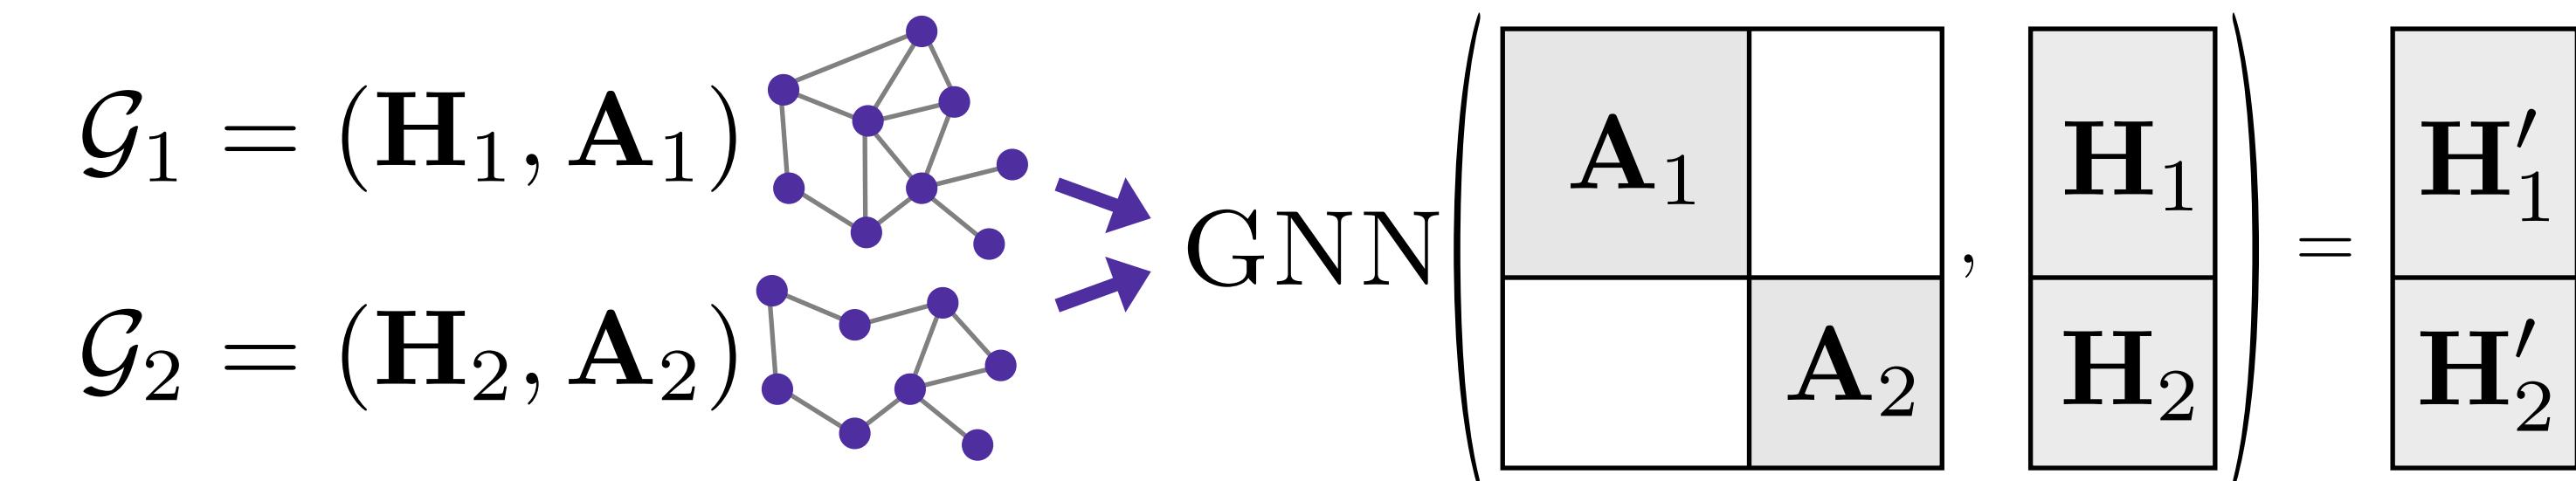
■ DGL  
DEEP GRAPH LIBRARY



# Mini-Batching in GNNs



**PyG** supports mini-batching on many small graphs

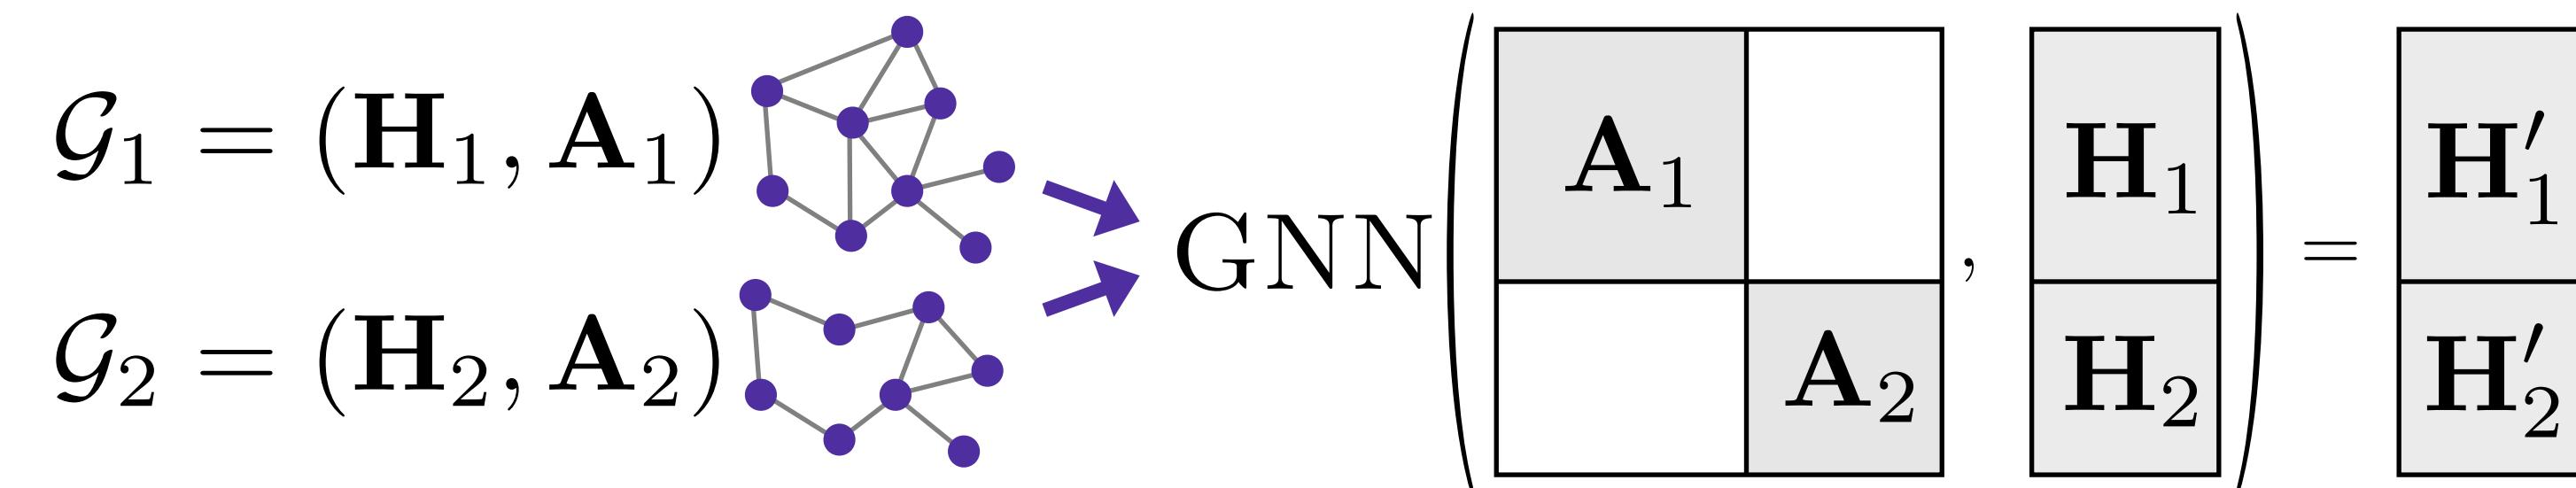




# Mini-Batching in GNNs



**PyG** supports mini-batching on many small graphs



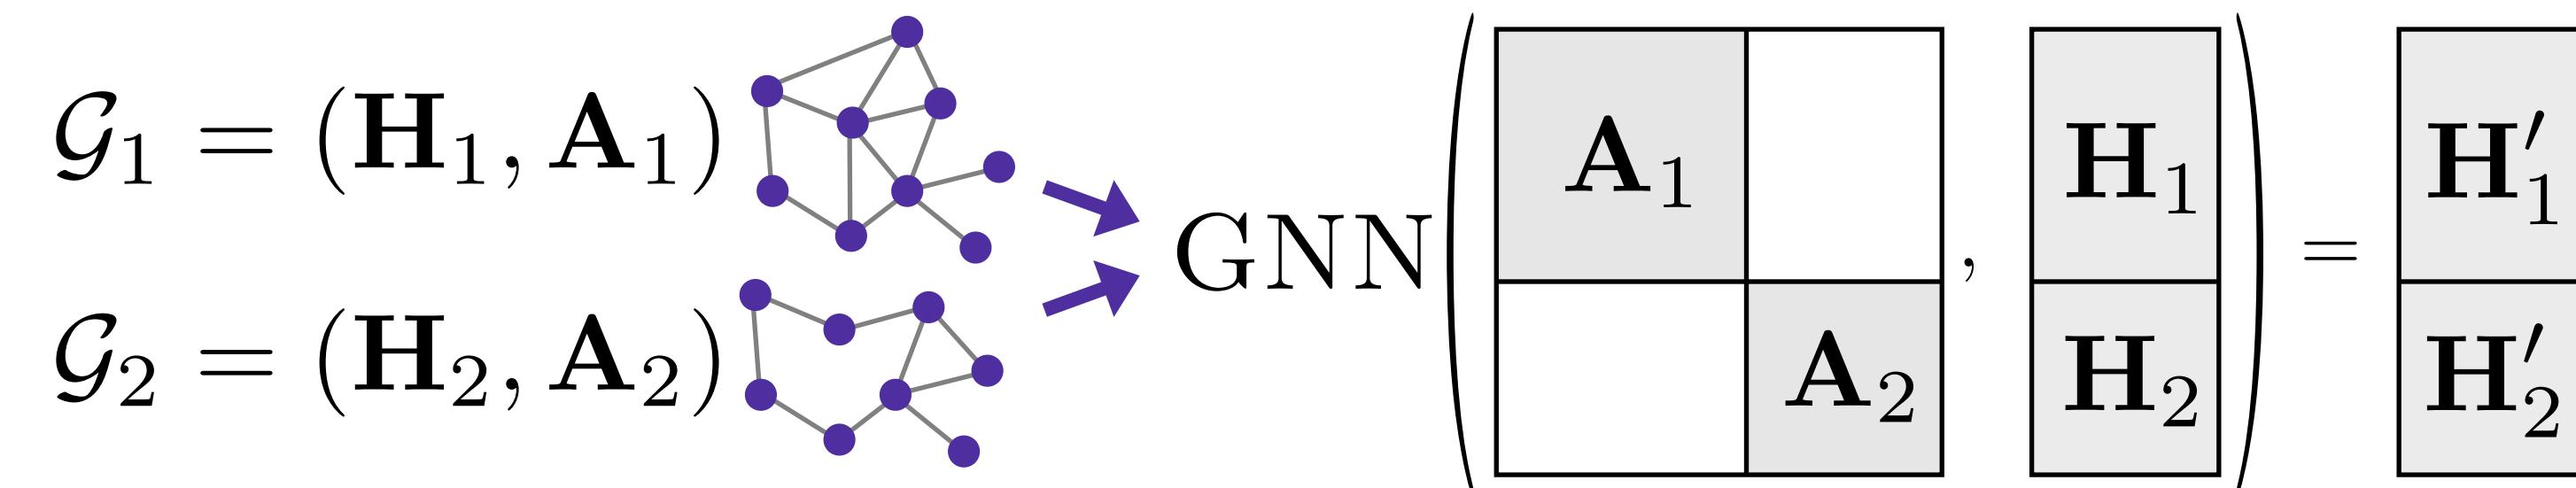
- ✓ no GNN modifications needed
- ✓ no memory/computation overhead
- ✓ supports examples of varying size



# Mini-Batching in GNNs



**PyG** supports mini-batching on many small graphs



- ✓ no GNN modifications needed
- ✓ no memory/computation overhead
- ✓ supports examples of varying size



```
from torch_geometric.datasets import TUDataset  
from torch_geometric.loader import DataLoader  
  
dataset = TUDataset(name='IMDB-BINARY')  
loader = DataLoader(dataset, batch_size=128)
```



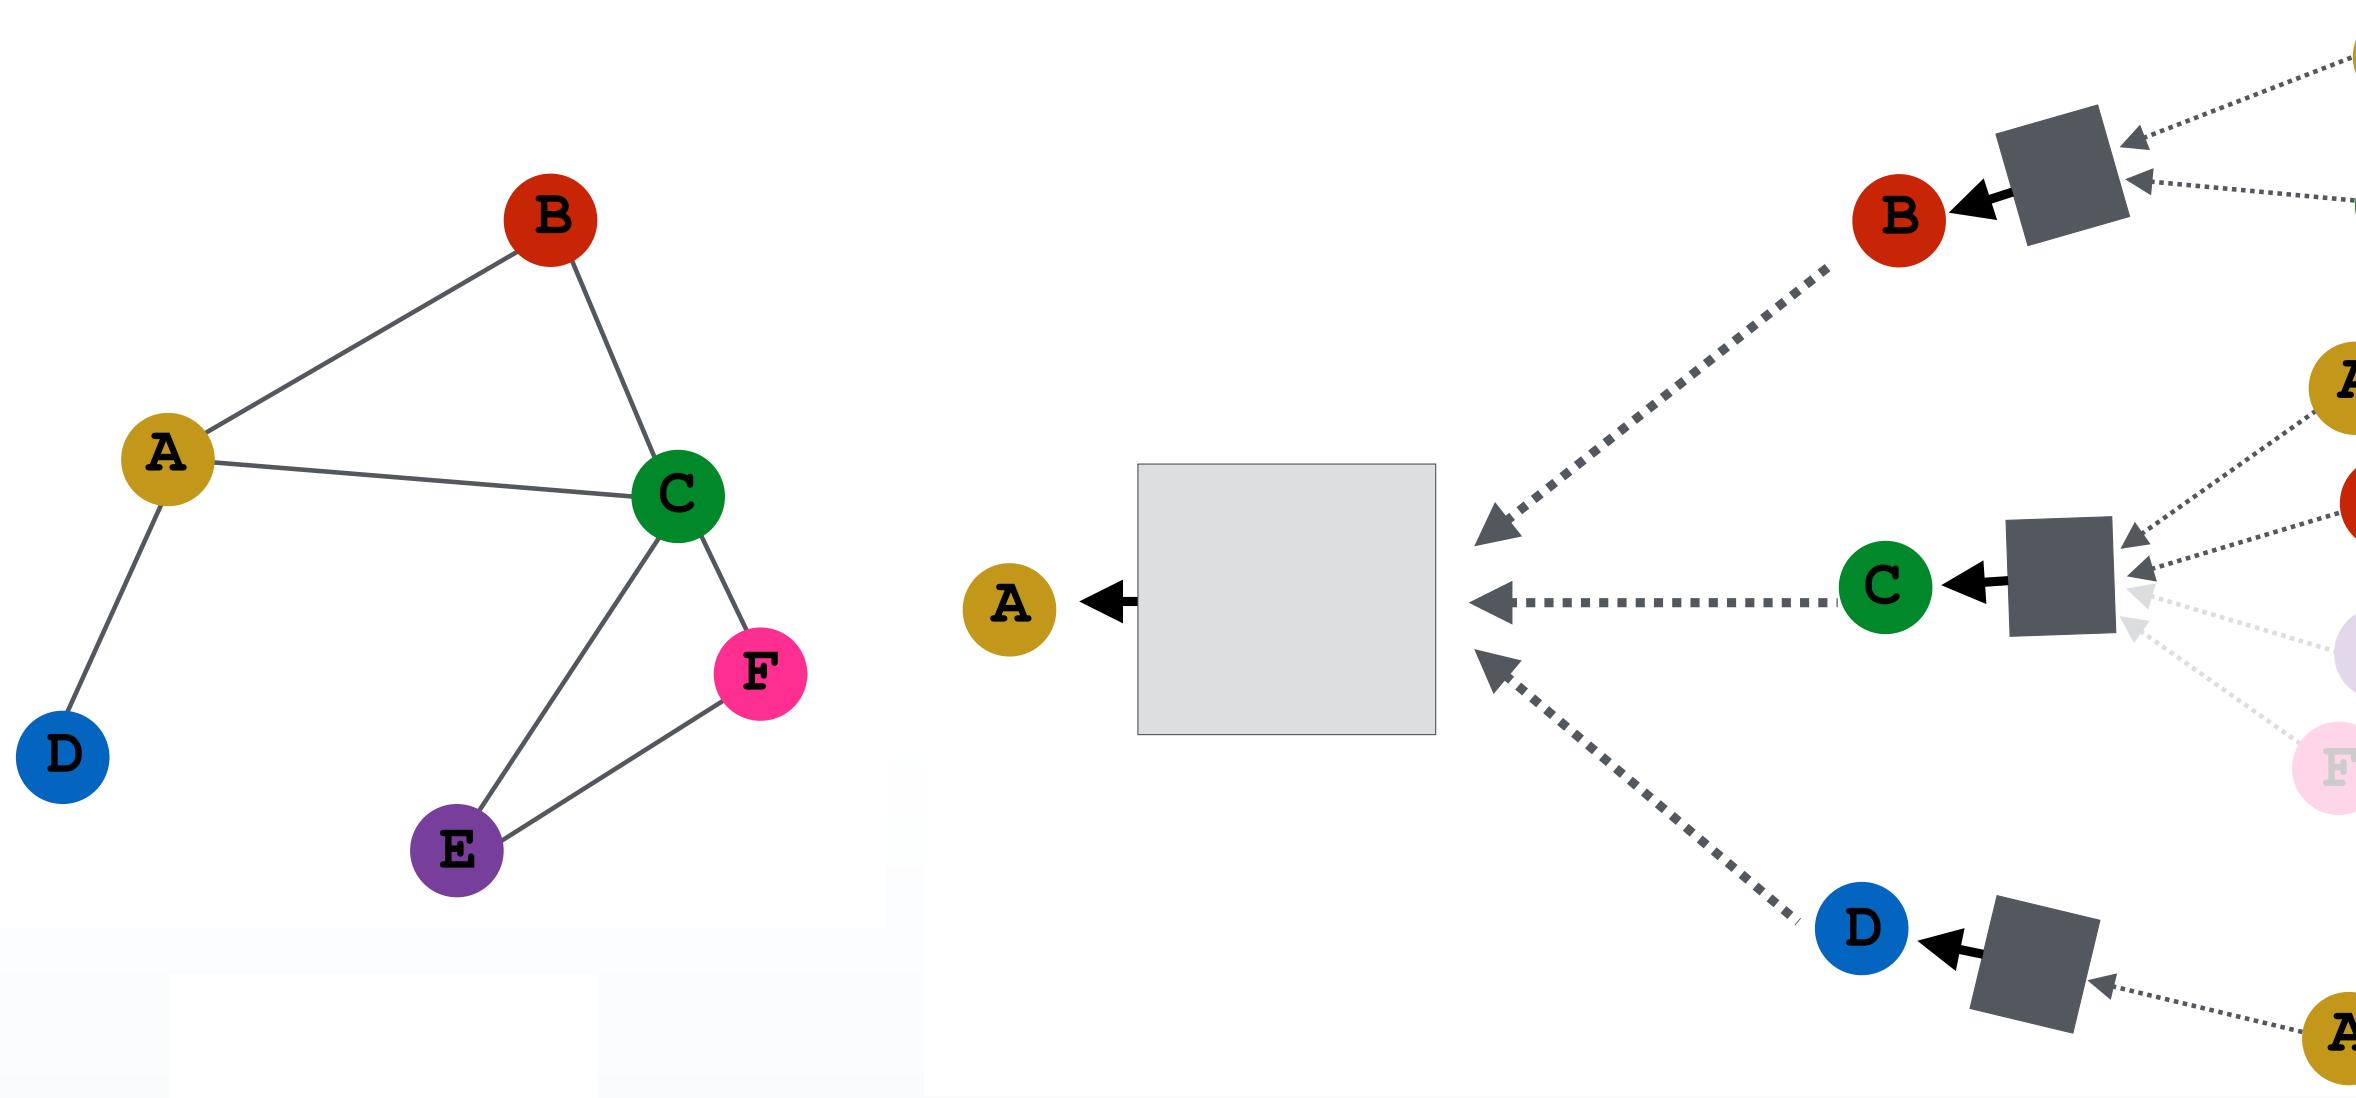
# Mini-Batching in GNNs

 **PyG** supports mini-batching on single giant graphs



# Mini-Batching in GNNs

 PyG supports mini-batching on single giant graphs

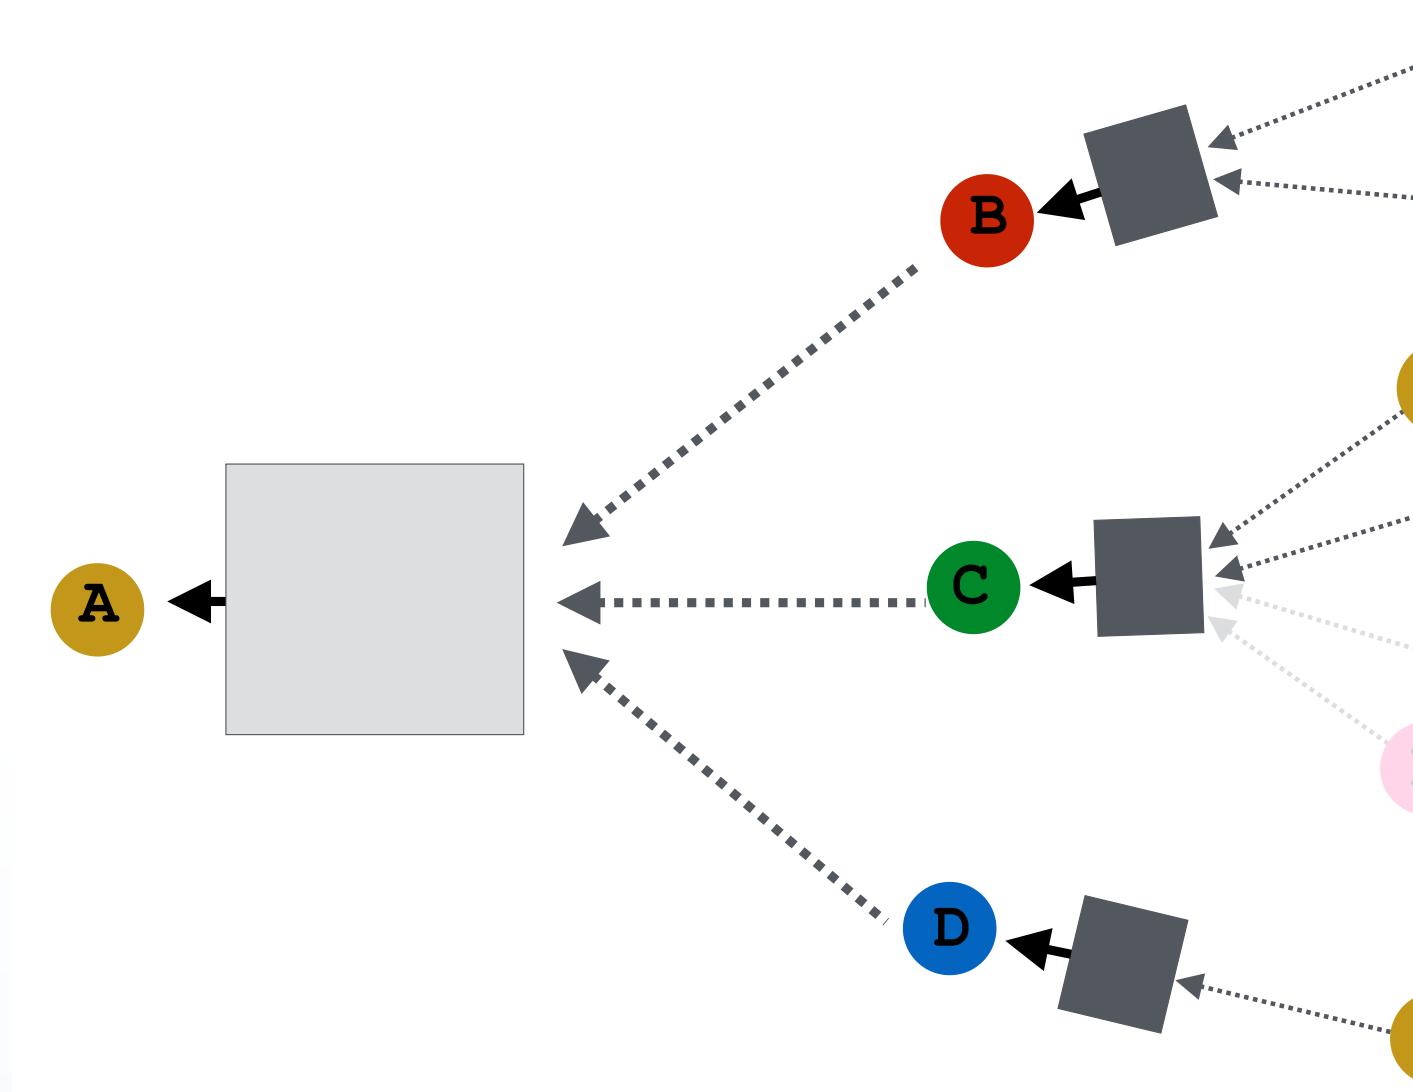
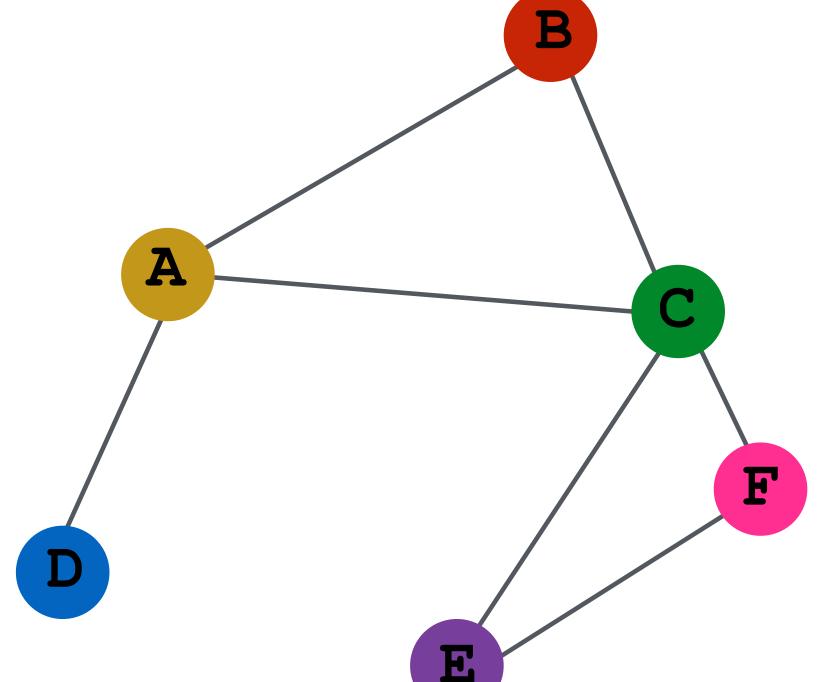




# Mini-Batching in GNNs



**PyG** supports mini-batching on single giant graphs



```
from torch_geometric.datasets import Reddit
from torch_geometric.loader import NeighborLoader

data = Reddit('data/Reddit')[0]
loader = NeighborLoader(data, batch_size=128,
                        num_neighbors=[25, 10])

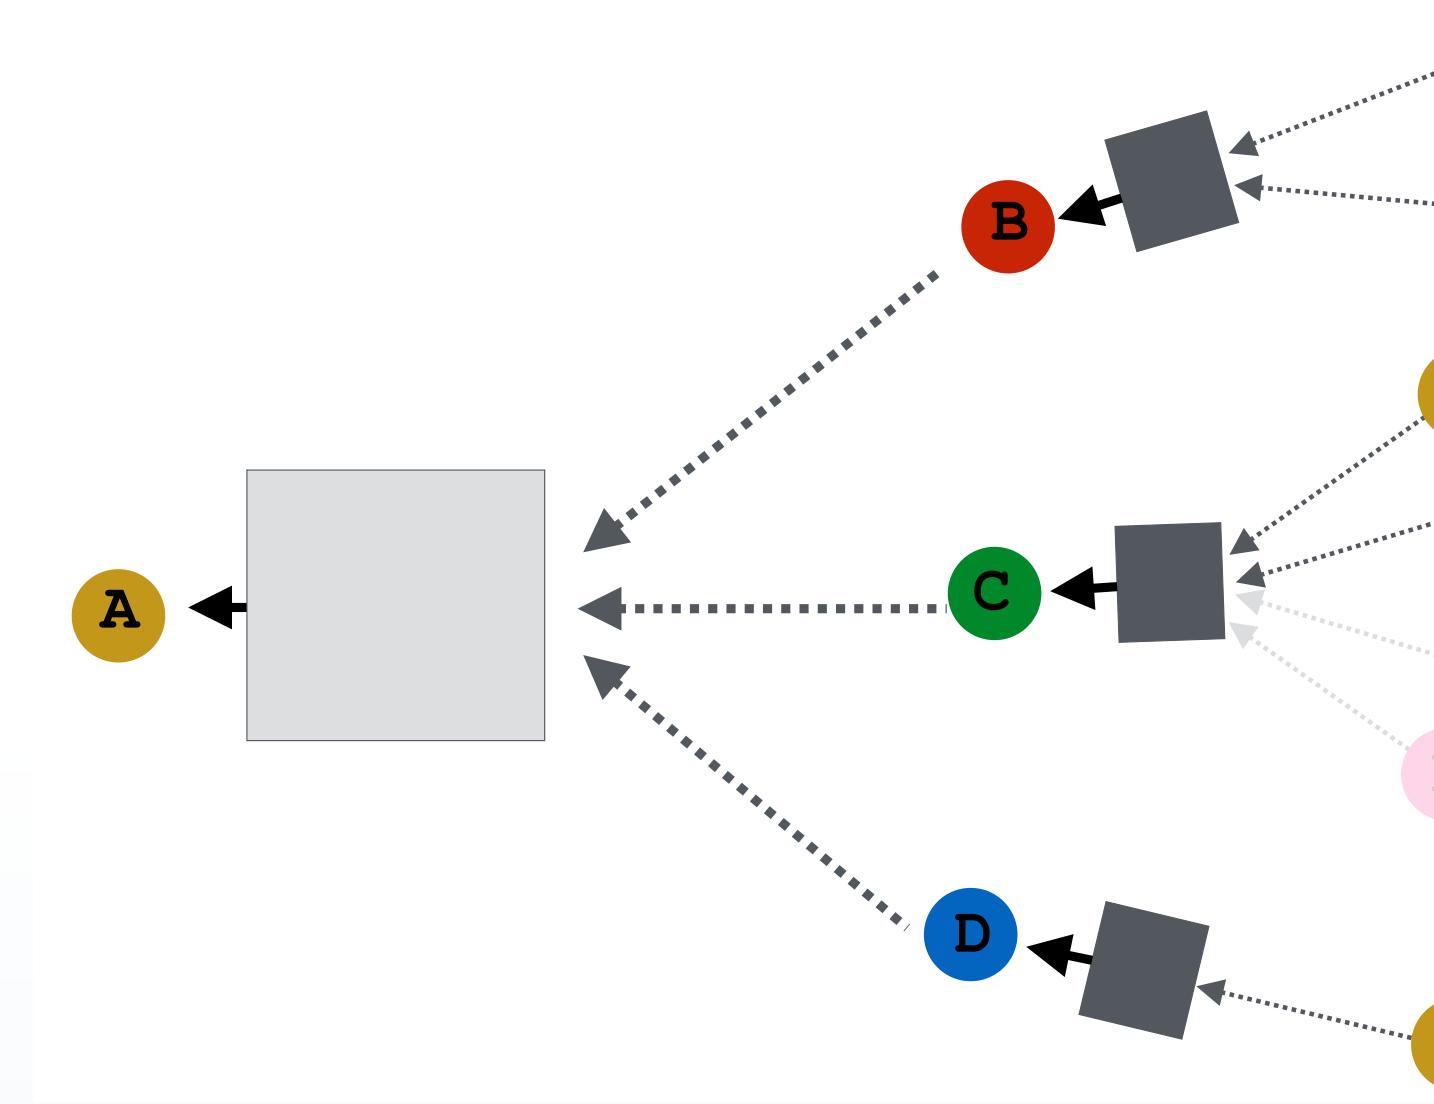
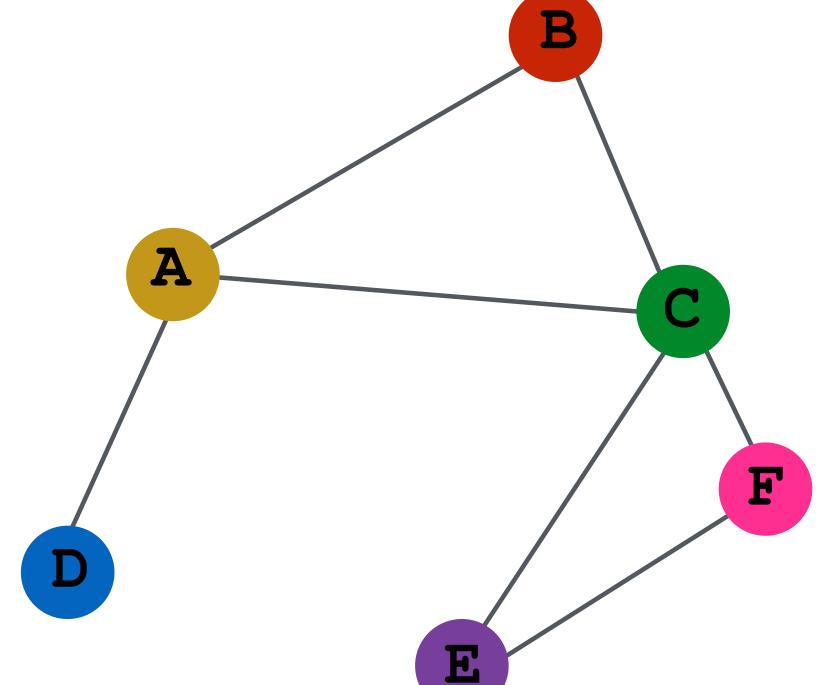
for batch in loader:
    model(batch.inputs, batch.edge_index)
```



# Mini-Batching in GNNs



**PyG** supports mini-batching on single giant graphs



```
from torch_geometric.datasets import Reddit
from torch_geometric.loader import NeighborLoader

data = Reddit('data/Reddit')[0]
loader = NeighborLoader(data, batch_size=128,
                        num_neighbors=[25, 10])

for batch in loader:
    model(batch.inputs, batch.edge_index)
```

Scalability support:

NeighborSampling  
SGC      ClusterGCN  
SIGN      Correct&Smooth      GraphSAINT  
ShaDow  
GNNAutoScale



# Graph Machine Learning Toolkit

 **PyG** can handle standard graph learning tasks with ease ...



# Graph Machine Learning Toolkit

PyG can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$



# Graph Machine Learning Toolkit

 PyG can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$

## Link-level Predictions

*Predict the class of a link*

$$\phi(\mathcal{G}, v, w) \in [0, 1]^C$$



# Graph Machine Learning Toolkit

 **PyG** can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$

## Link-level Predictions

*Predict the class of a link*

$$\phi(\mathcal{G}, v, w) \in [0, 1]^C$$

## Graph-level Predictions

*Predict the class of a graph*

$$\phi(\mathcal{G}) \in [0, 1]^C$$



# Graph Machine Learning Toolkit

 PyG can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$

## Link-level Predictions

*Predict the class of a link*

$$\phi(\mathcal{G}, v, w) \in [0, 1]^C$$

## Graph-level Predictions

*Predict the class of a graph*

$$\phi(\mathcal{G}) \in [0, 1]^C$$

... but is not limited to those:



# Graph Machine Learning Toolkit

 **PyG** can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$

## Link-level Predictions

*Predict the class of a link*

$$\phi(\mathcal{G}, v, w) \in [0, 1]^C$$

## Graph-level Predictions

*Predict the class of a graph*

$$\phi(\mathcal{G}) \in [0, 1]^C$$

... but is not limited to those:

Unsupervised Learning

Few/Zero-Shot Learning

Self-Supervised Learning

Pre-Training

Explainability

...



# Graph Machine Learning Toolkit

 PyG can handle standard graph learning tasks with ease ...

## Node-level Predictions

*Predict the class of a node*

$$\phi(\mathcal{G}, v) \in [0, 1]^C$$

## Link-level Predictions

*Predict the class of a link*

$$\phi(\mathcal{G}, v, w) \in [0, 1]^C$$

## Graph-level Predictions

*Predict the class of a graph*

$$\phi(\mathcal{G}) \in [0, 1]^C$$

... but is not limited to those:

Unsupervised Learning    Self-Supervised Learning  
Few/Zero-Shot Learning    Pre-Training    Explainability    ...

 PyG provides over 80 examples with access to over 200 benchmark datasets to get familiar with the latest trends in graph machine learning



# Additional Features



# Additional Features

## TorchScript

Convert pure Python GNN  
model to an optimized and  
standalone program



```
class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = GCNConv( ... ).jittable()
        self.conv2 = GCNConv( ... ).jittable()

    model = torch.jit.script(GNN())
```



# Additional Features

## TorchScript

Convert pure Python GNN model to an optimized and standalone program



```
class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = GCNConv( ... ).jittable()
        self.conv2 = GCNConv( ... ).jittable()

model = torch.jit.script(GNN())
```



## PyTorch Lightning

1. Choose a GNN model
2. Setup a trainer (#GPUs, accelerator type)
3. Call `trainer.fit()`

```
data = Reddit('data/Reddit')
model = GraphSAGE(in_channels, out_channels)

trainer = Trainer(gpus=2, accelerator='ddp')

trainer.fit(model, data)
trainer.test()
```



# Additional Features

## TorchScript

Convert pure Python GNN model to an optimized and standalone program



```
class GNN(torch.nn.Module):
    def __init__(self):
        self.conv1 = GCNConv( ... ).jittable()
        self.conv2 = GCNConv( ... ).jittable()

    model = torch.jit.script(GNN())
```



## PyTorch Lightning

1. Choose a GNN model
2. Setup a trainer (#GPUs, accelerator type)
3. Call `trainer.fit()`



```
data = Reddit('data/Reddit')
model = GraphSAGE(in_channels, out_channels)

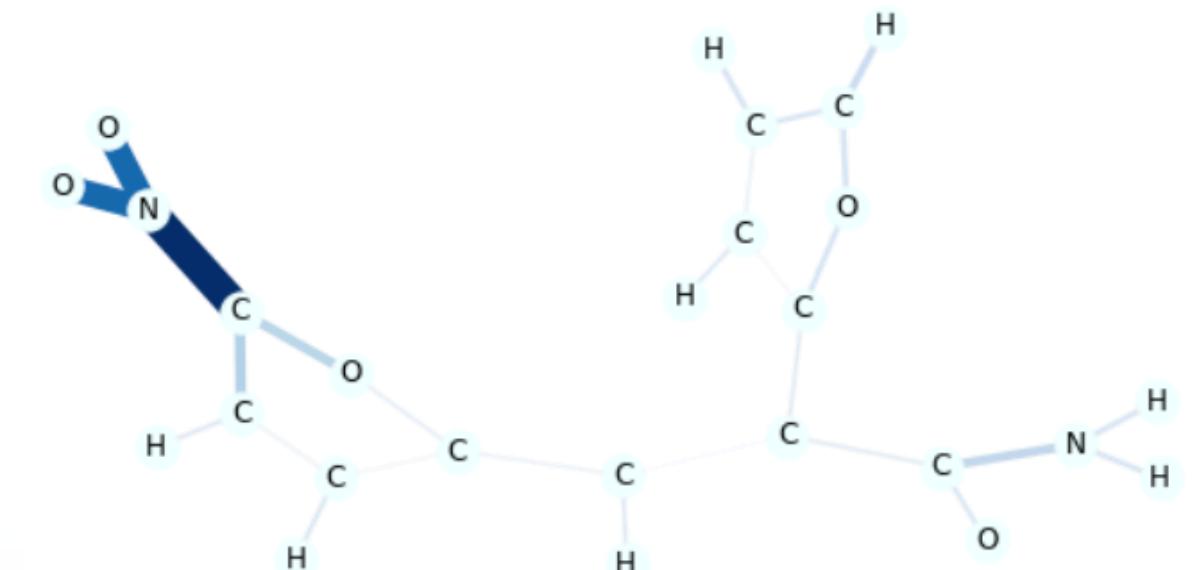
trainer = Trainer(gpus=2, accelerator='ddp')

trainer.fit(model, data)
trainer.test()
```



## Captum

Explain GNN predictions out-of-the-box



```
from captum.attr import IntegratedGradients

ig = IntegratedGradients(model)
ig.attribute(input, edge_index, target=0)
```



# Success Stories

stars 11k

forks 1.9k

issues 1.2k closed

pull requests 318 closed

- ~700 research papers using  PyG
- ~2M weekly wheel downloads
- ~180 external contributors
- ~600 members on Slack



**Yann LeCun**  
@ylecun

A fast & nice-looking PyTorch library for geometric deep learning (NN on graphs and other irregular structures).



**Thomas Kipf**  
@thomaskipf

PyTorch Geometric has been growing into a fantastic library for graph neural nets and related methods.



# Success Stories

stars 11k

forks 1.9k

issues 1.2k closed

pull requests 318 closed

- ~700 research papers using  **PyG**
- ~2M weekly wheel downloads
- ~180 external contributors
- ~600 members on Slack



**Yann LeCun**  
@ylecun

A fast & nice-looking PyTorch library for geometric deep learning (NN on graphs and other irregular structures).

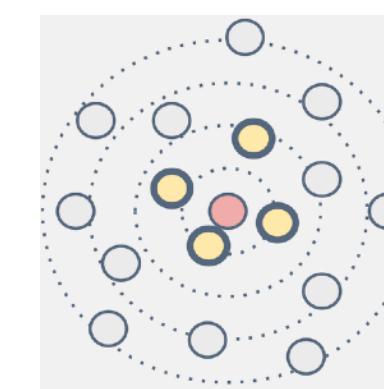
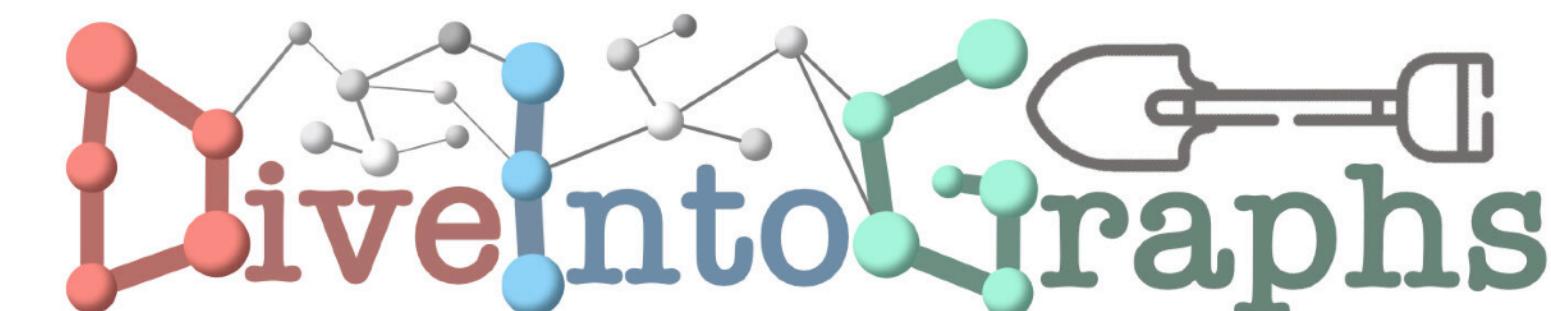
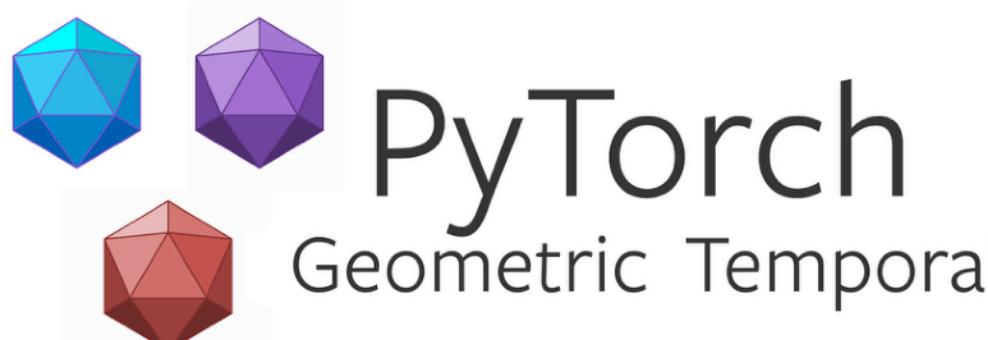


**Thomas Kipf**  
@thomaskipf

PyTorch Geometric has been growing into a fantastic library for graph neural nets and related methods.



**PyG** already has its own ecosystem:

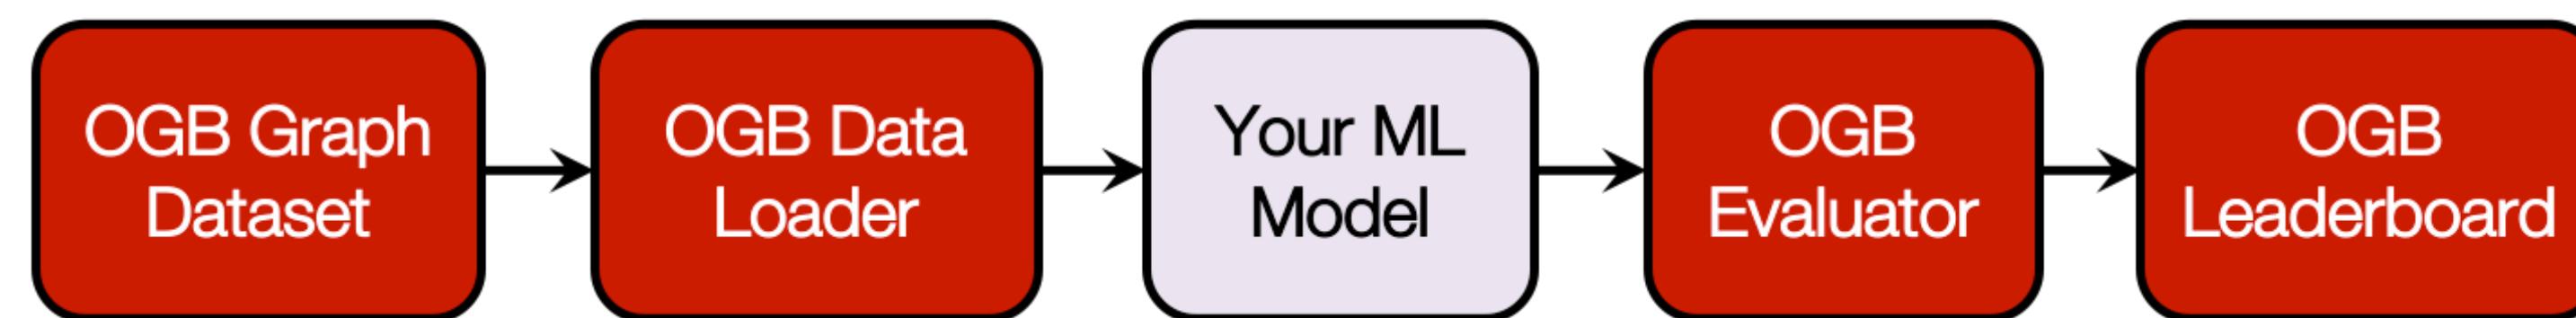


**DeepGCNs**

**Open Catalyst Project**  
FACEBOOK AI



# Success Stories



[https://ogb.stanford.edu  
/snap-stanford/ogb](https://ogb.stanford.edu/snap-stanford/ogb)

- OGB provides a variety of **realistic** and **large-scale** graph benchmark datasets
- Data loaders are compatible with PyG
- Dive-in examples are utilizing PyG



```
from ogb import PygGraphPropPredDataset
from torch_geometric.loader import DataLoader

dataset = PygGraphPropPredDataset('ogbg-molhiv')
loader = DataLoader(dataset, batch_size=128)
```

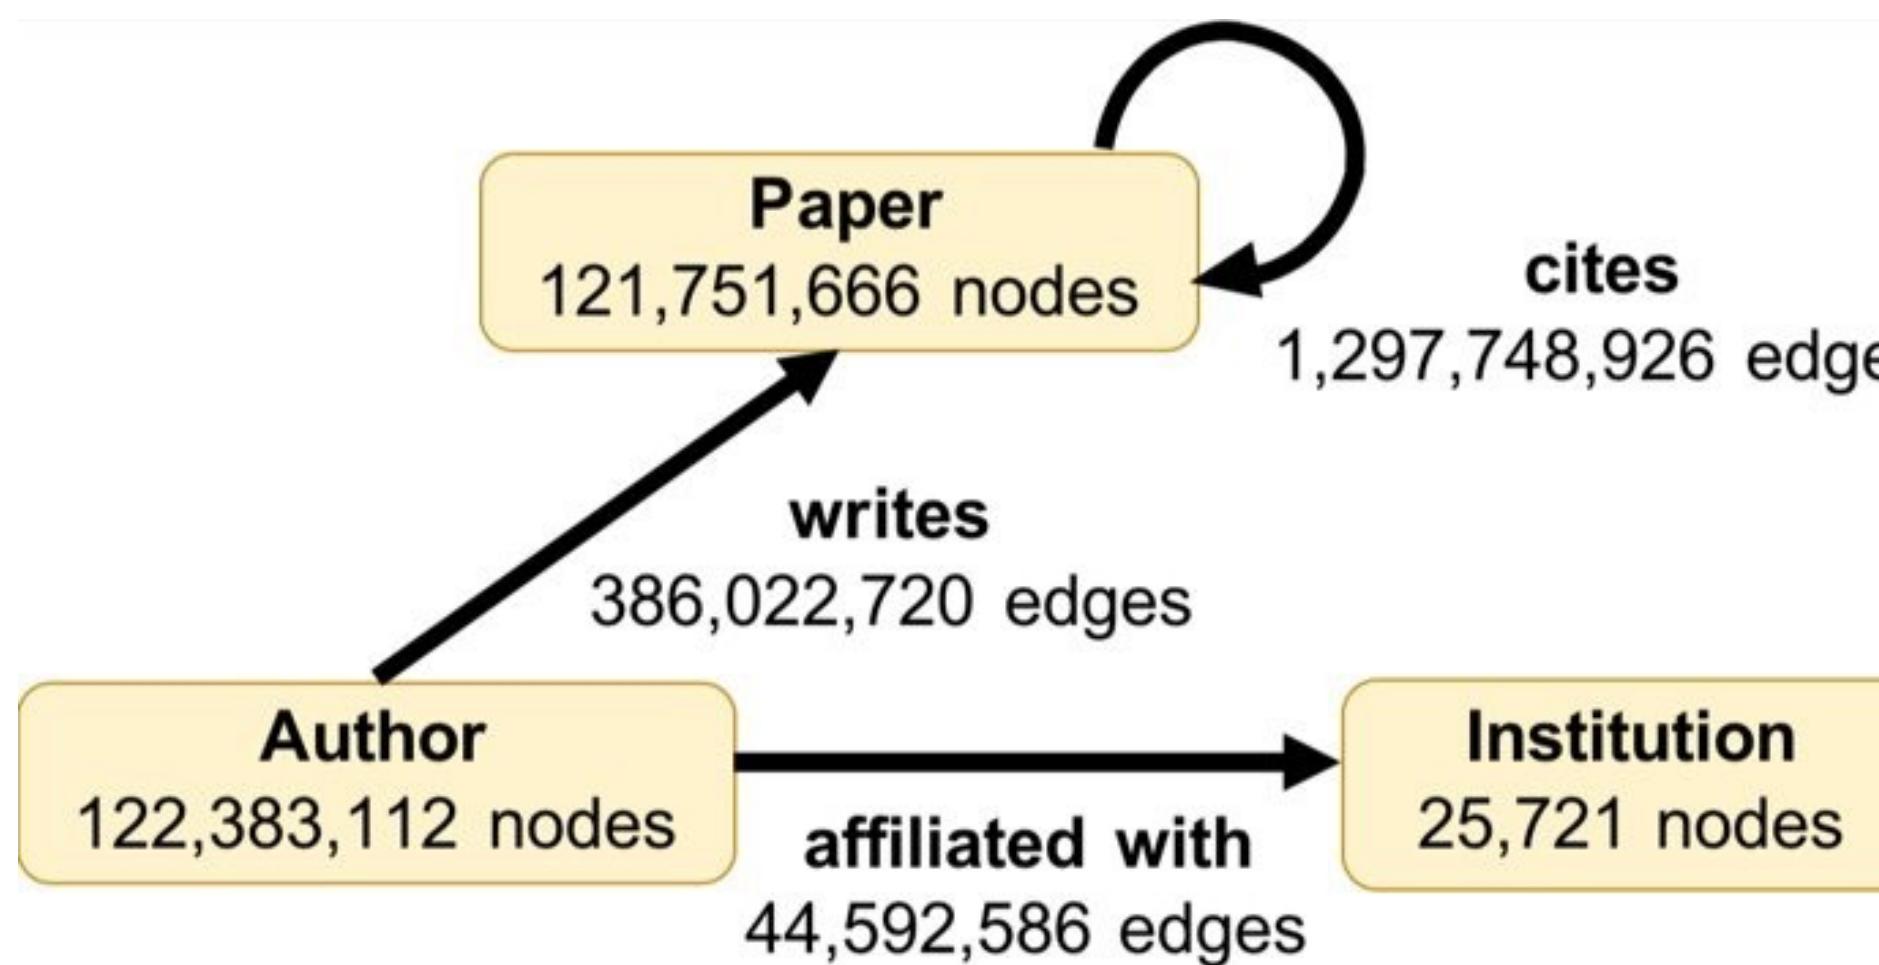


# Success Stories

## OGB-LSC @ KDD Cup 2021

— Large - Scale Challenge —

Competition on three large-scale graph datasets, e.g.:



- ~240M nodes across three node types
- ~1.7B edges across three edge types

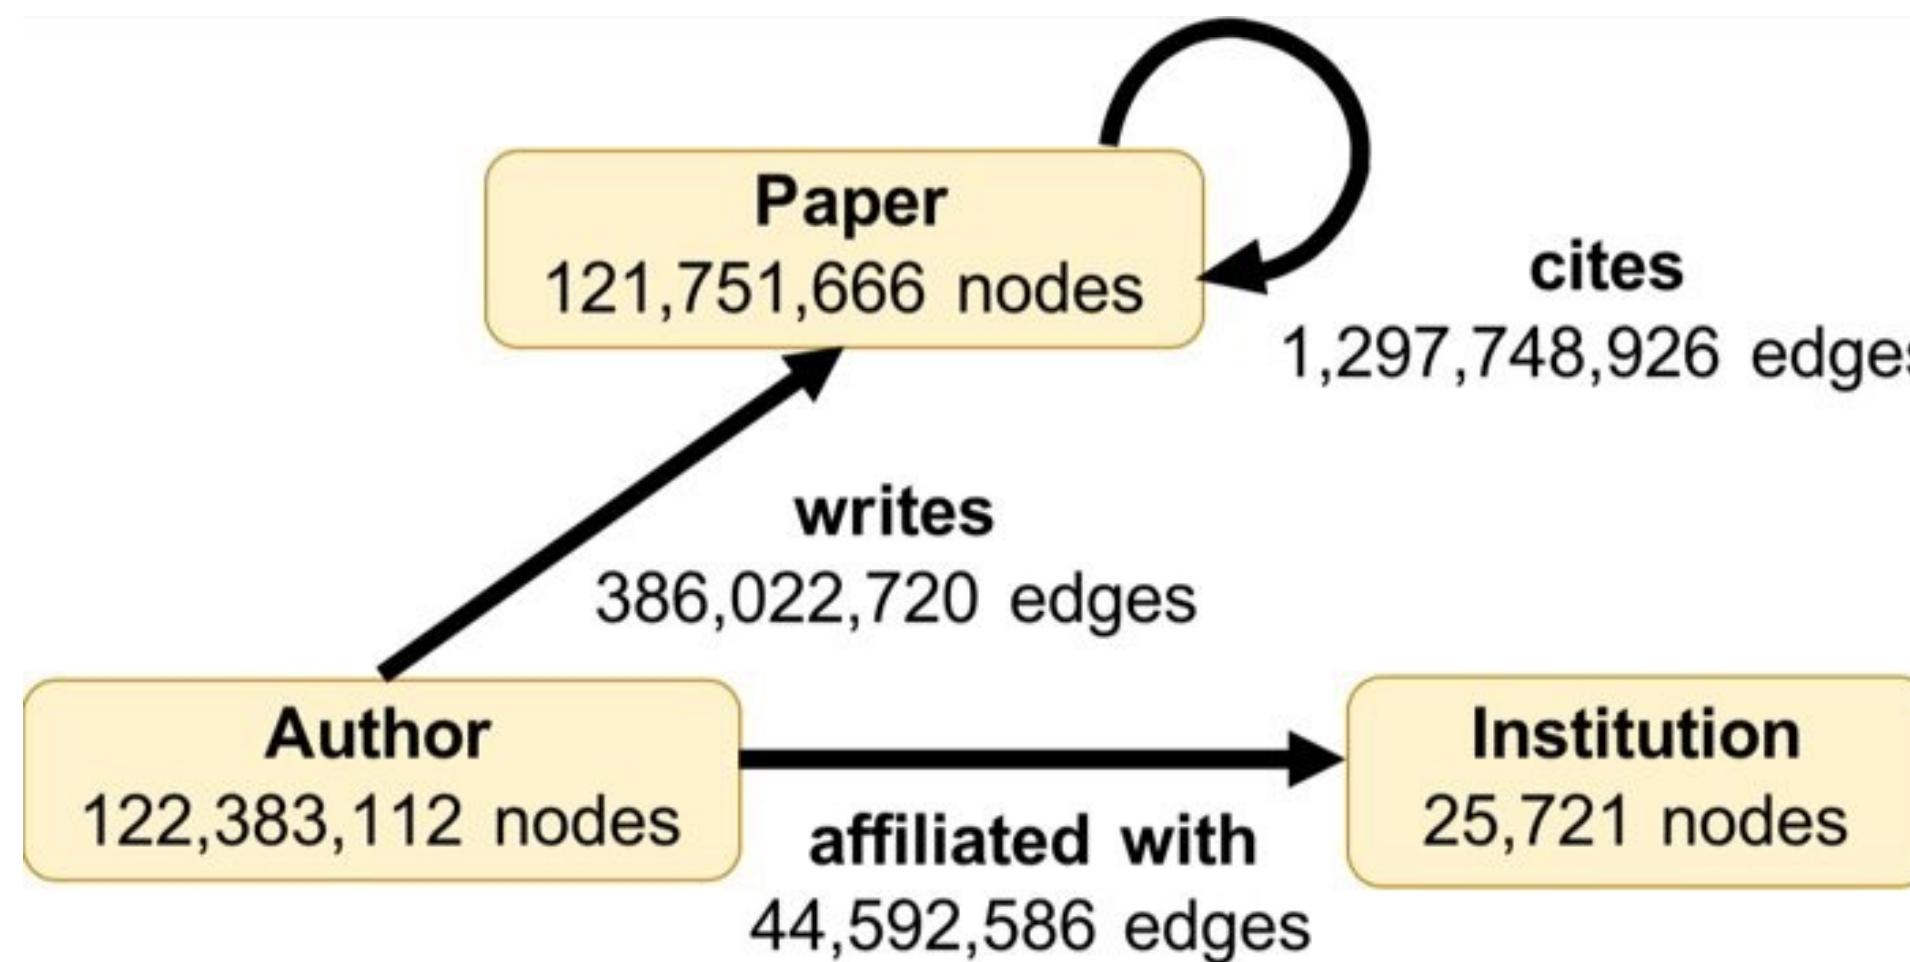


# Success Stories

## OGB-LSC @ KDD Cup 2021

— Large - Scale Challenge —

Competition on three large-scale graph datasets, e.g.:



- ~240M nodes across three node types
- ~1.7B edges across three edge types

We used PyG to benchmark GNNs:

Model	#Params	Validation	Test
MLP	0.5M	52.67	52.73
LABELPROP	0	58.44	56.29
SGC	0.7M	65.82	65.29
SIGN	3.8M	66.64	66.09
MLP+C&S	0.5M	66.98	66.18
GRAPH SAGE (NS)	4.9M	66.79	66.28
GAT (NS)	4.9M	67.15	66.80
R-GRAPH SAGE (NS)	12.2M	69.86	68.94
R-GAT (NS)	12.3M	<b>70.02</b>	<b>69.42</b>



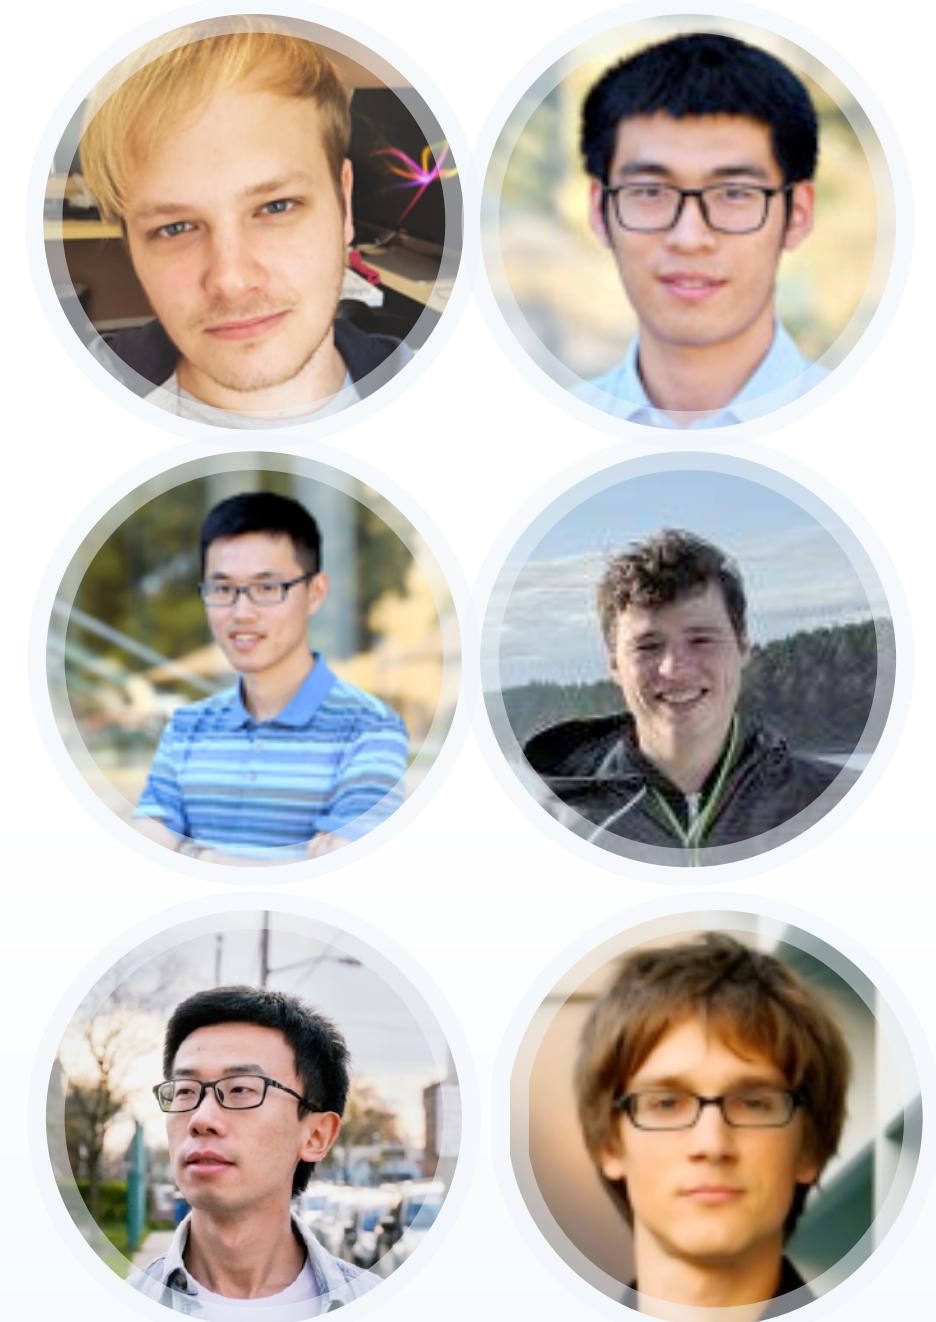
# Joining Forces

We have build a larger team for the future development of PyG, in close collaboration with



- ✓ ensure longevity of PyG
- ✓ integrate tools from both parties into a unified package
- ✓ keep up with integrating latest trends in academic research
- ✓ extend its scope to better support real-world use-cases
- ✓ make it (even) easier to use for both academia and industry
- ✓ make it (even) more scalable

<https://pyg.org>





# Joining Forces

We have build a larger team for the future development of PyG,  
in close collaboration with



don't  
un

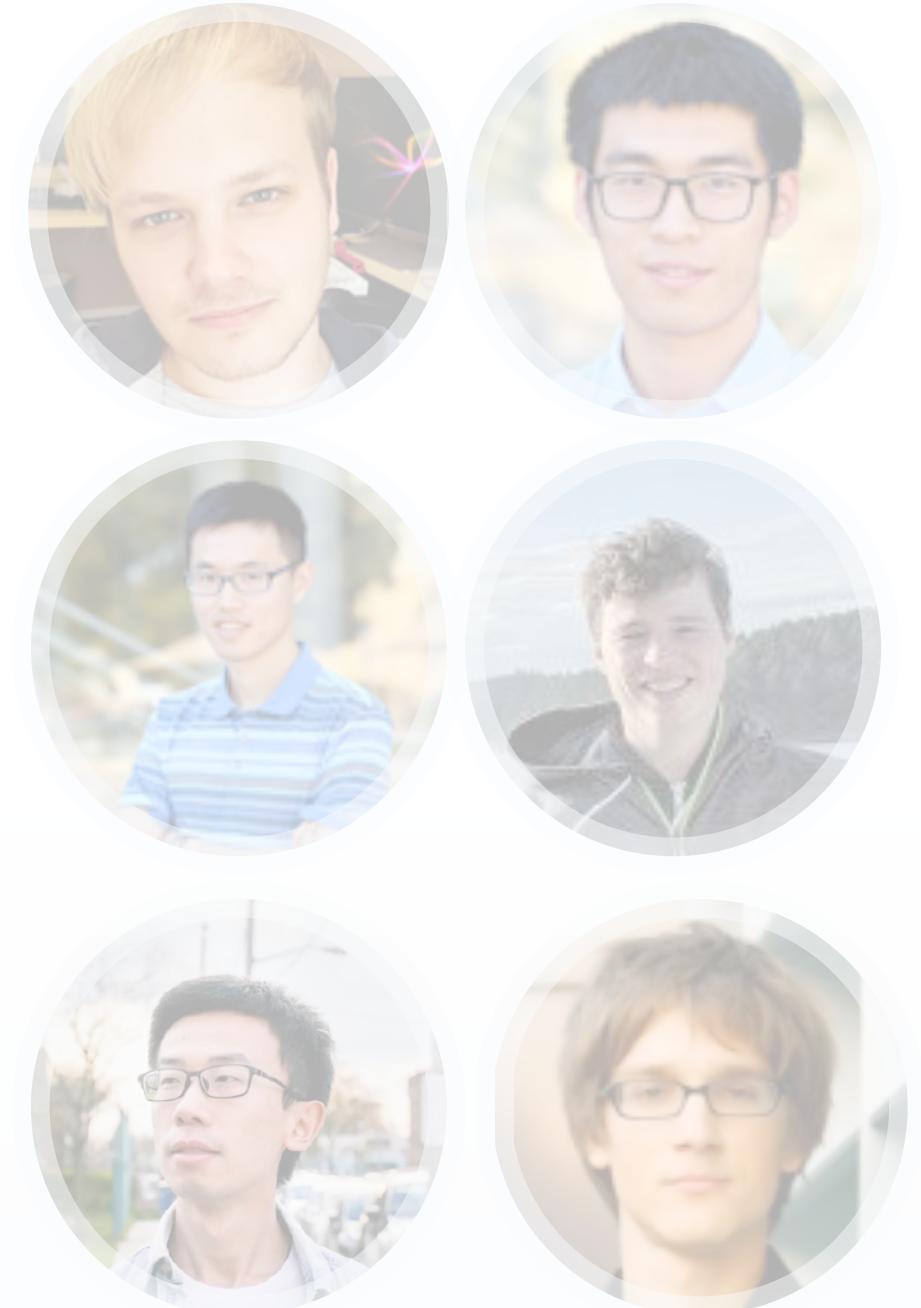
do it  
Crea

1

We start with this by introducing ...

- ✓ ensure longevity
- ✓ integrate tools from both parties into a unified package
- ✓ keep up with integrating latest trends in academic research
- ✓ extend its scope to better support real-world use-cases
- ✓ make it (even) easier to use for both academia and industry
- ✓ make it (even) more scalable

<https://pyg.org>



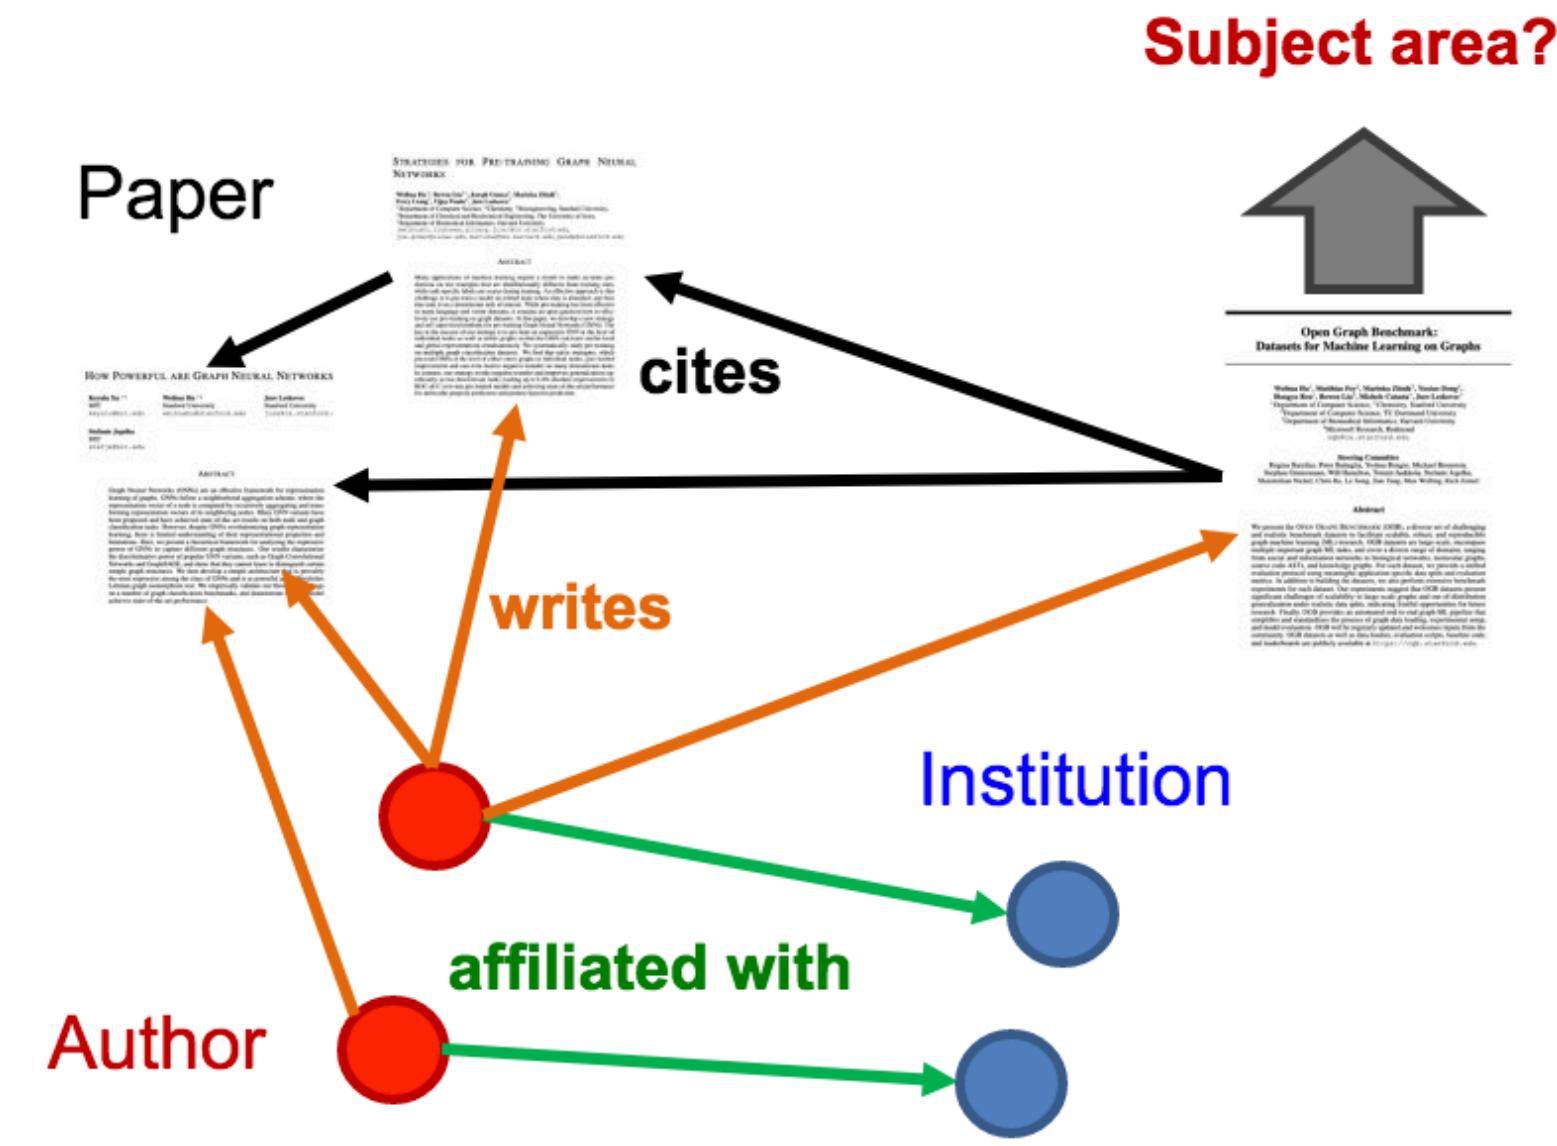


# **PyG 2.0**

## Advanced Representation Learning on Graphs



# Overview

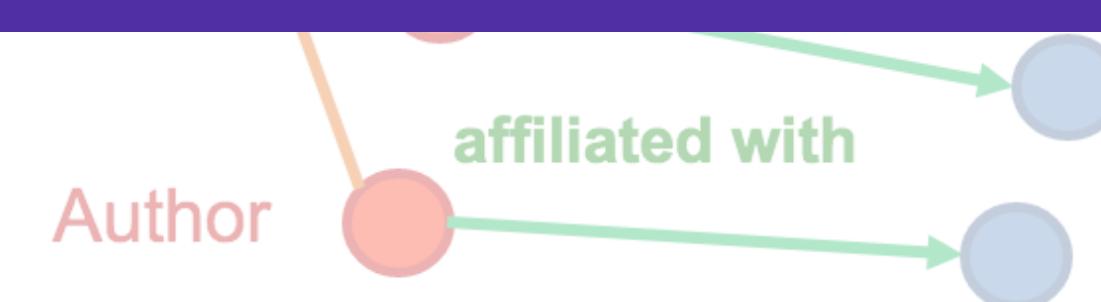


## Heterogeneous Graphs



# Overview

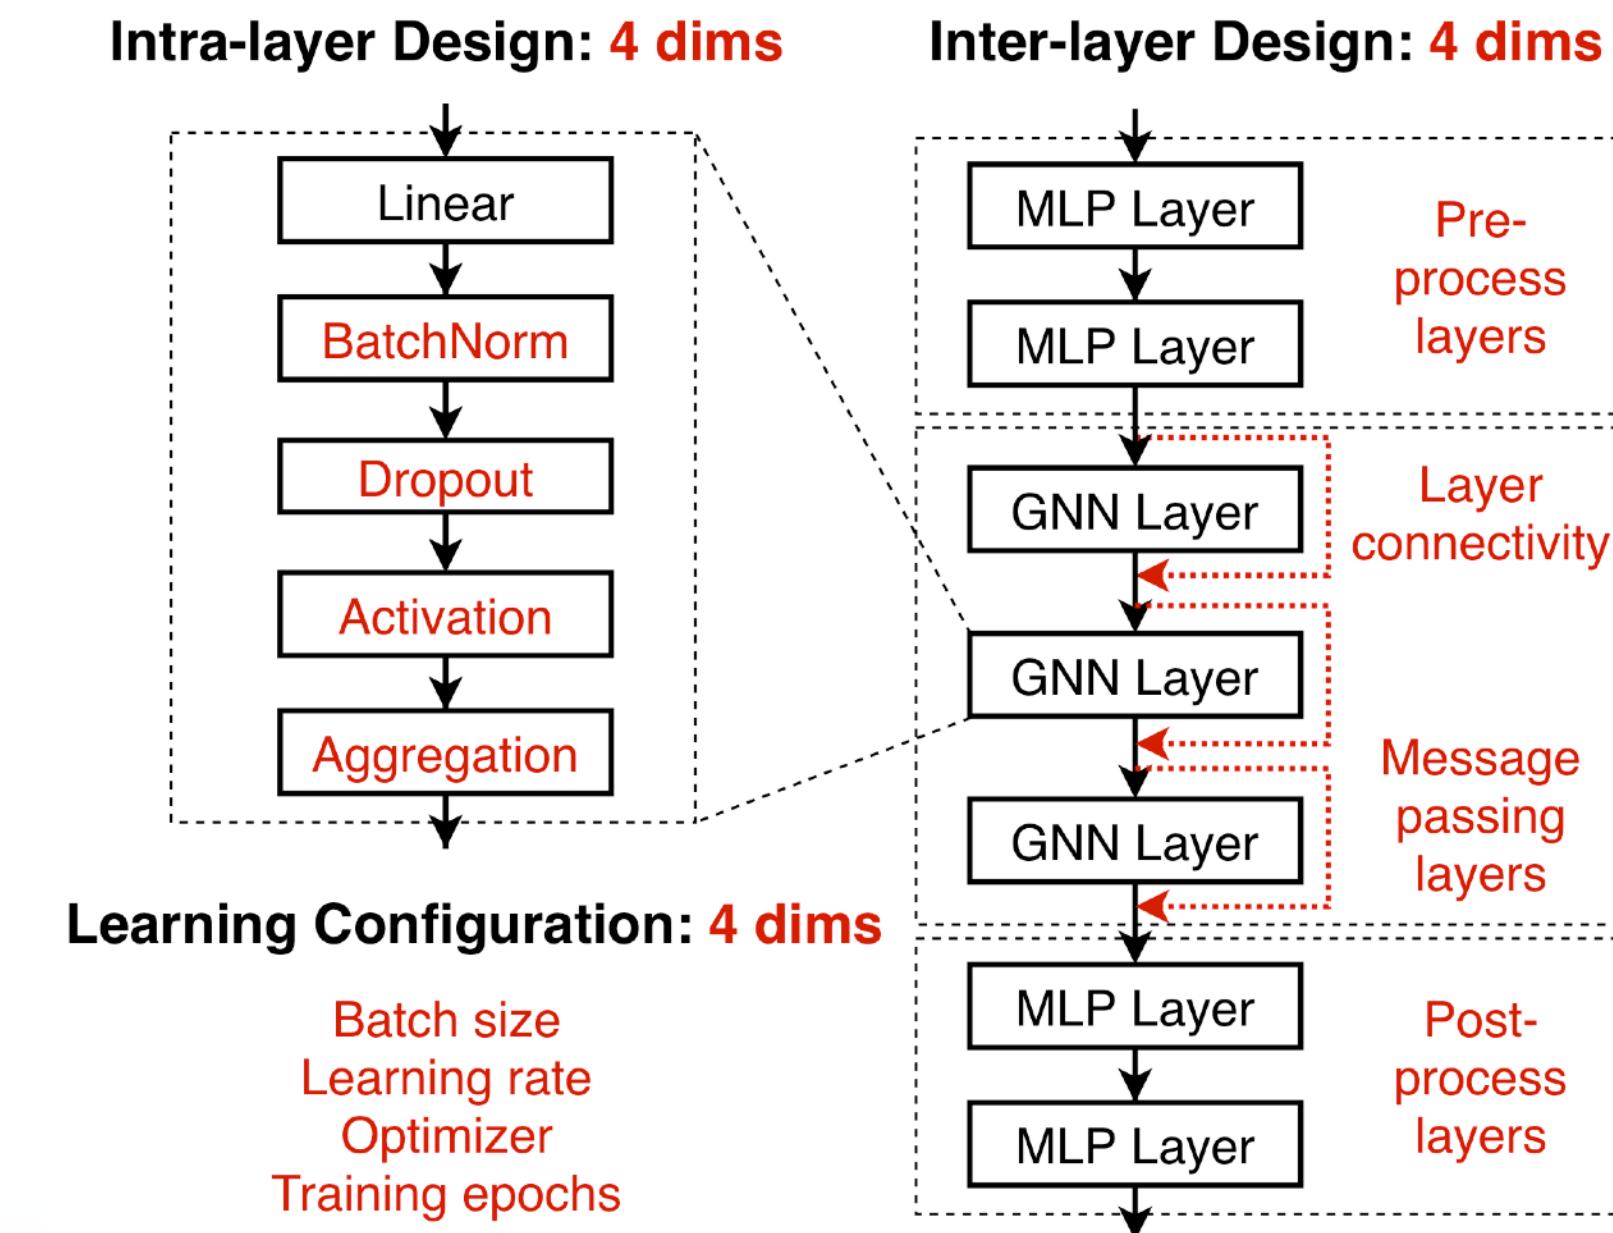
(Nearly) all real-world graphs  
are heterogeneous!



## Heterogeneous Graphs



# Overview

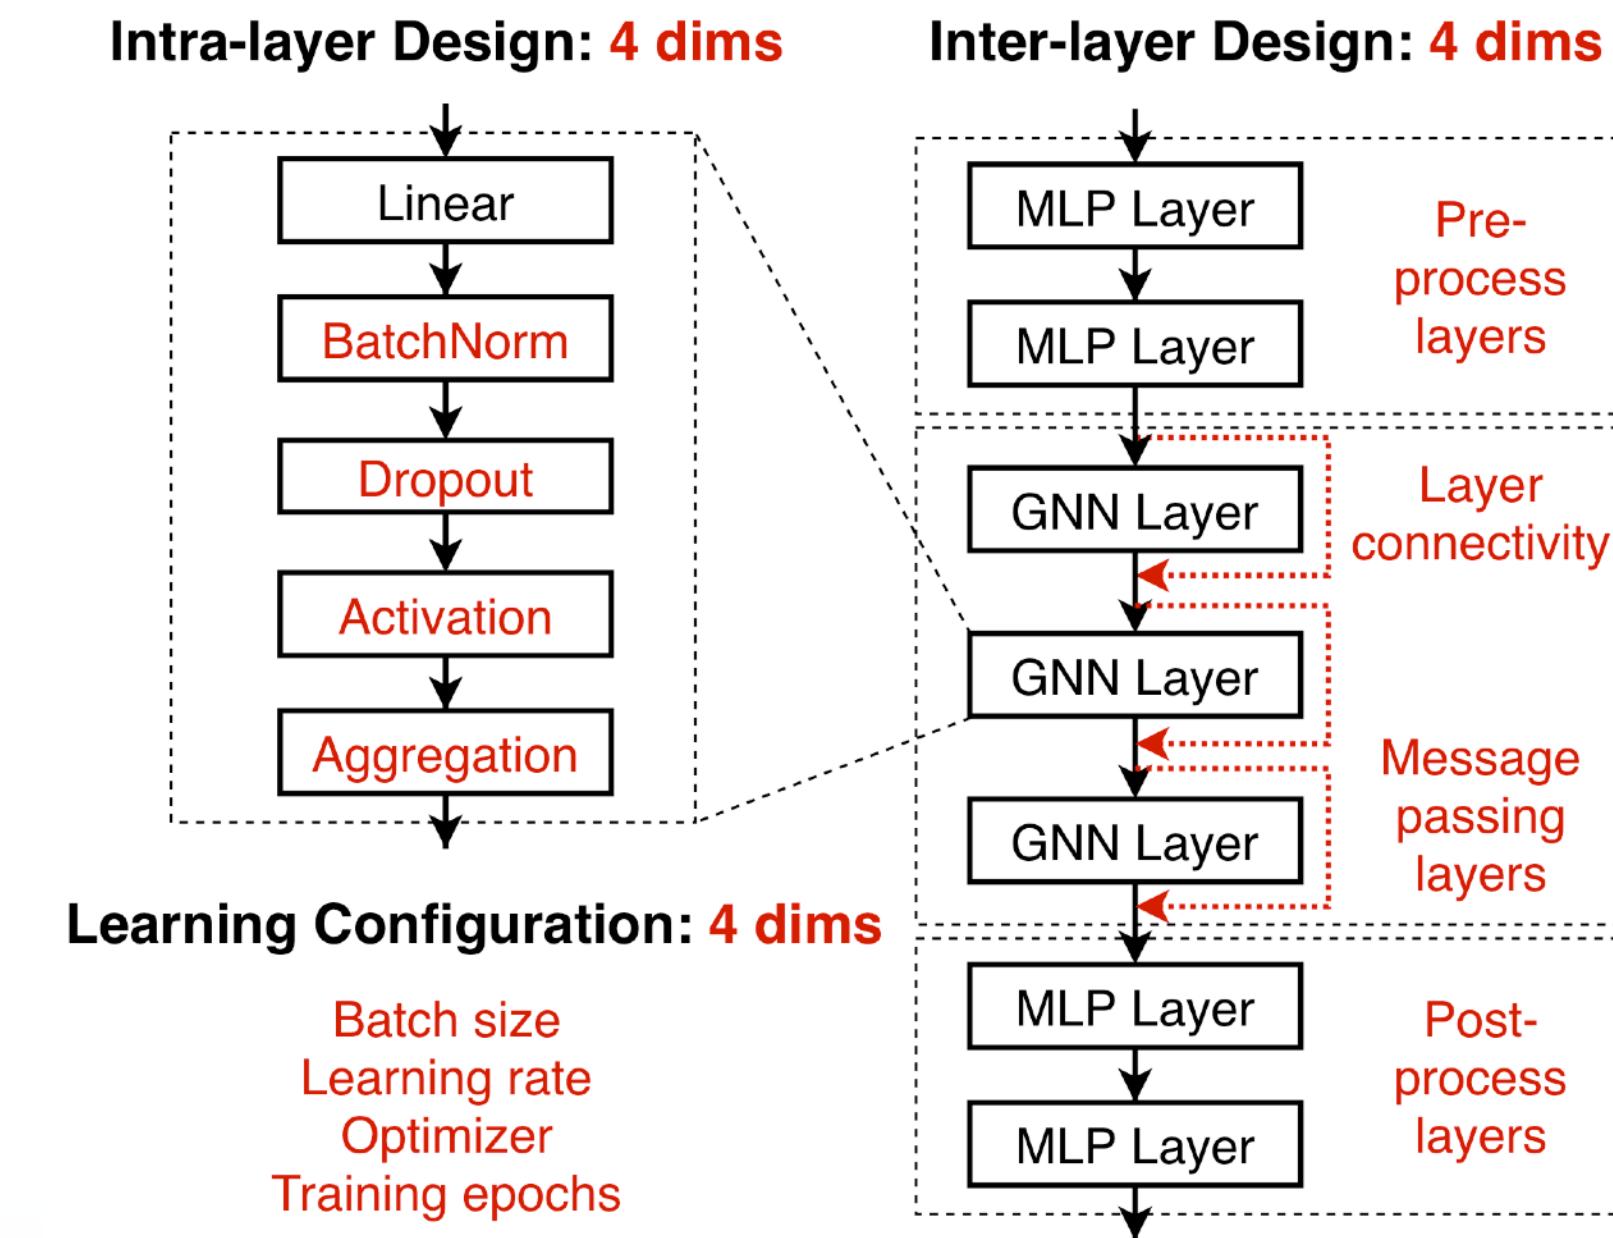


## Heterogeneous Graphs

## GraphGym



# Overview



## Heterogeneous Graphs

- ✓ Half-Precision Support
- ✓ \*.profile for benchmarking runtimes and memory consumptions of GNNs
- ✓ A variety of new operators, models, datasets and examples

## GraphGym



# Heterogeneous Graph Support

Heterogeneous graph learning is notoriously challenging



# Heterogeneous Graph Support

Heterogeneous graph learning is notoriously challenging

- Different input feature distributions across node and edge types
- Necessity of learning node/edge type dependent representations  
*non-shared weights across different node and edge types*  
*bipartite message passing*
- Heterogeneous scalability approaches  
*Relational-wise neighborhood sampling*
- Complicated implementation  
*requires sequentially iterating over different node and edge types*  
*involves keeping track of different input feature dimensionalities*



# Heterogeneous Graph Support

Heterogeneous graph learning is notoriously challenging

- Different input feature distributions across node and edge types
- Necessity of learning node/edge type dependent representations  
*non-shared weights across different node and edge types*  
*bipartite message passing*
- Heterogeneous scalability approaches  
*Relational-wise neighborhood sampling*
- Complicated implementation  
*requires sequentially iterating over different node and edge types*  
*involves keeping track of different input feature dimensionalities*



**PyG** makes working with heterogeneous graphs a breeze



# Heterogeneous Graph Support

## Data Storage

Holds information about different node and edge types in individual containers

Edge types are described by a triplet of source node, relation and destination node type

Transformations enhance the graph for message passing, e.g., by *adding reverse edges*



```
from torch_geometric.data import HeteroData

data = HeteroData()

data['user'].x = ... # User node feature matrix
data['product'].x = ... # Product node feature matrix

# Connecting user and product nodes via "buys" relation:
data['user', 'buys', 'product'].edge_index = ... # [2, num_edges]

# Adding reverse edges:
from torch_geometric.transforms import ToUndirected

data = ToUndirected()(data)
```



# Heterogeneous Graph Support

## Data Storage

Holds information about different node and edge types in individual containers

Edge types are described by a triplet of source node, relation and destination node type

Transformations enhance the graph for message passing, e.g., by *adding reverse edges*

```
● ● ●  
from torch_geometric.data import HeteroData  
  
data = HeteroData()  
  
data['user'].x = ... # User node feature matrix  
data['product'].x = ... # Product node feature matrix  
  
# Connecting user and product nodes via "buys" relation:  
data['user', 'buys', 'product'].edge_index = ... # [2, num_edges]  
  
# Adding reverse edges:  
from torch_geometric.transforms import ToUndirected  
  
data = ToUndirected()(data)
```

We provide a full example for loading raw \*.csv files in the documentation



# Heterogeneous Graph Support

## Heterogeneous Graph Neural Networks

A homogeneous GNN can be converted to a heterogeneous one by learning distinct parameters for each individual edge type:

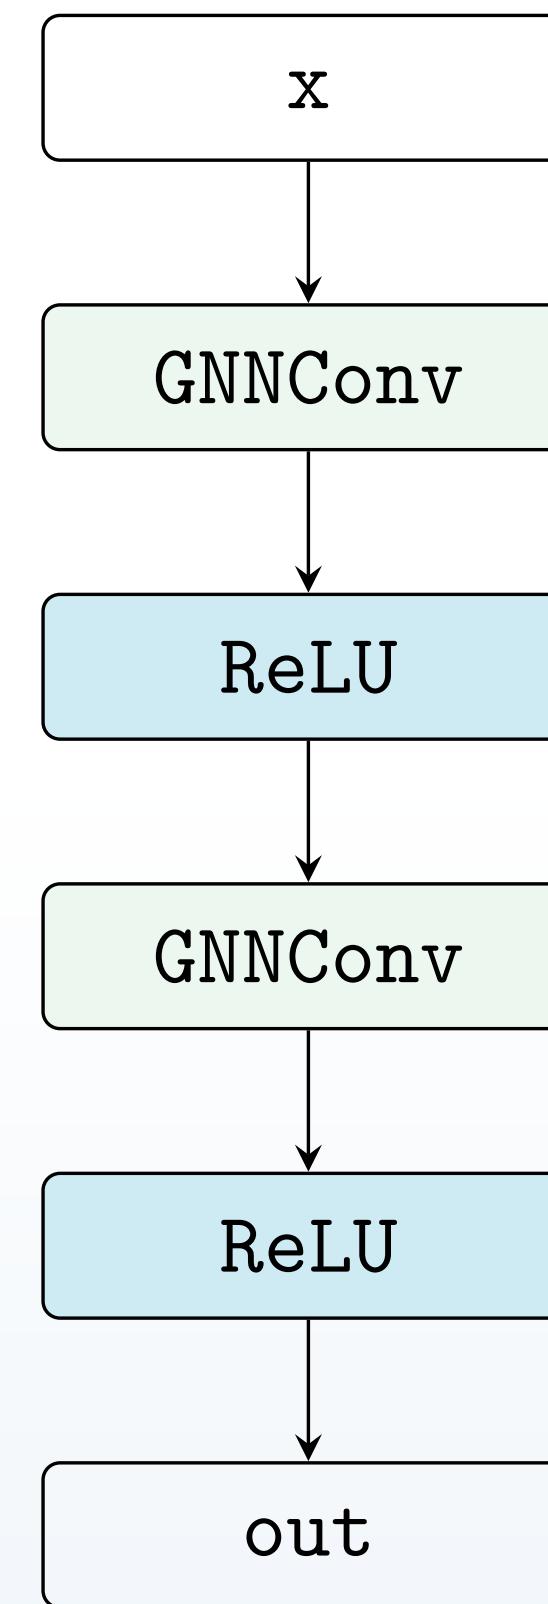
A custom GNN for each relation

$$\mathbf{h}_i^{(\ell+1)} = \sum_{r \in \mathcal{R}} \text{GNN}_{\theta}^{(r)} \left( \mathbf{h}_i^{(\ell)}, \{\mathbf{h}_j^{(\ell)} : j \in \mathcal{N}^{(r)}(i)\} \right)$$

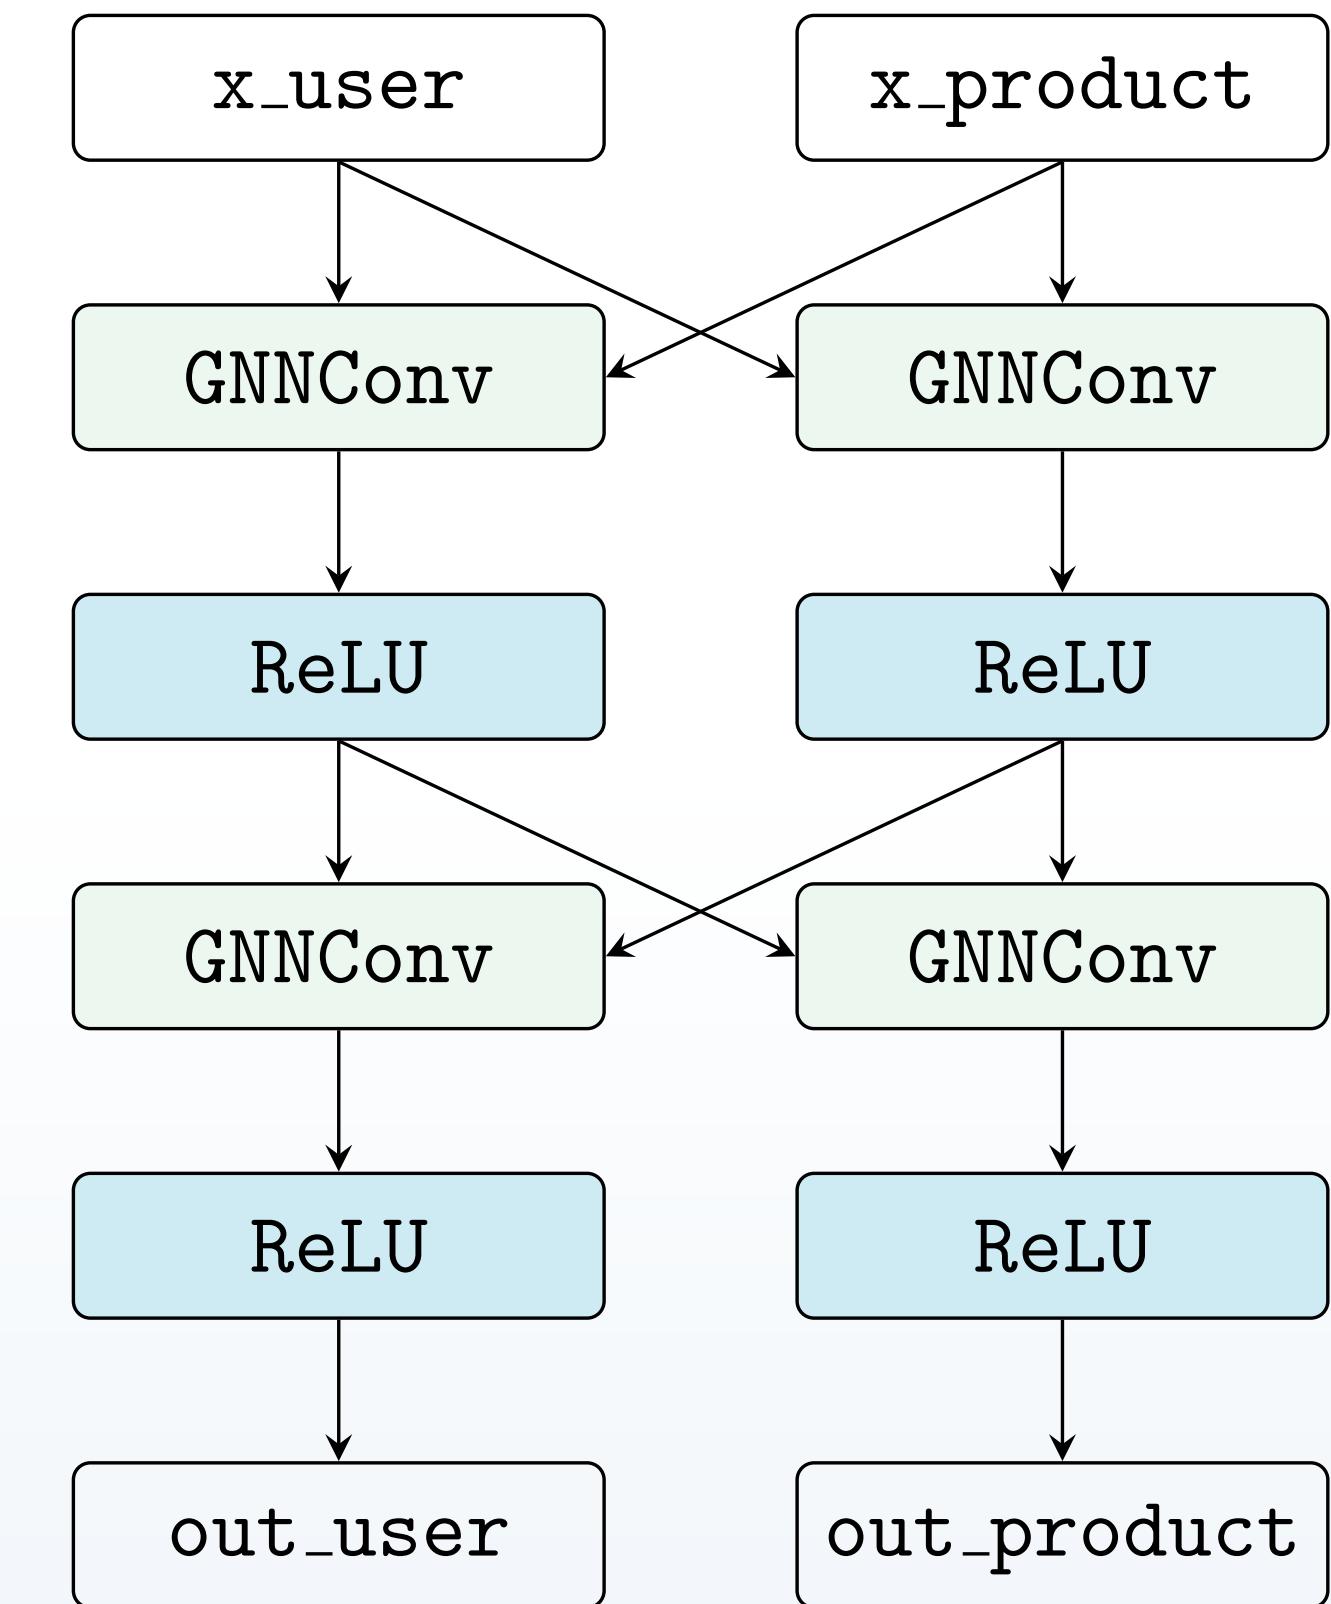
The number of relations

↑  
Relational-wise neighborhood

Homogeneous Model



Heterogeneous Model





# Heterogeneous Graph Support

## Heterogeneous Graph Neural Networks

Rapid growth in the number of parameters w.r.t number of relations may lead to overfitting on rare relations

## Basis-decomposition for regularization

A custom GNN for each basis

$$\mathbf{h}_i^{(\ell+1)} = \sum_{r \in \mathcal{R}} \sum_{b=1}^B \text{GNN}_{\theta}^{(b)} \left( \mathbf{h}_i^{(\ell)}, \{a_{r,b}^{(\ell)} \cdot \mathbf{h}_j^{(\ell)} : j \in \mathcal{N}(i)\} \right)$$

↑   ↓  
The number of bases                      Relational-depend  
    trainable coefficients



# Heterogeneous Graph Support

PyG can automatically convert homogeneous GNNs to heterogeneous ones

```
to_hetero(model (node_types, edge_types)):  
to_hetero_with_bases(model (node_types, edge_types)):
```

- 1.Duplicates message passing modules for each edge type
- 2.Transforms the underlying computation graph so that messages are exchanged along different edge types
- 3.Uses lazy initialization (-1) to handle different input feature dimensionalities

```
from torch_geometric.nn import GAT, to_hetero  
  
model = GAT(in_channels=-1, hidden_channels=64,  
            out_channels=72, num_layers=2)  
  
model = to_hetero(model, (node_types, edge_types))  
  
out = model(data.x_dict, data.edge_index_dict)
```



# Heterogeneous Graph Support

## Heterogeneous Graph Samplers

Scaling up heterogeneous GNNs  
to large-scale graphs with ease  
via relational neighbor sampling

Only requires a few lines  
of code change!



```
from torch_geometric.datasets import OGB_MAG
from torch_geometric.loader import NeighborLoader

data = OGB_MAG(path)[0]

loader = NeighborLoader(
    data,
    num_neighbors={key: [15, 10] for key in data.edge_types},
    batch_size=128,
    input_nodes='paper')

for batch in loader:
    model(batch.x_dict, batch.edge_index_dict)
```

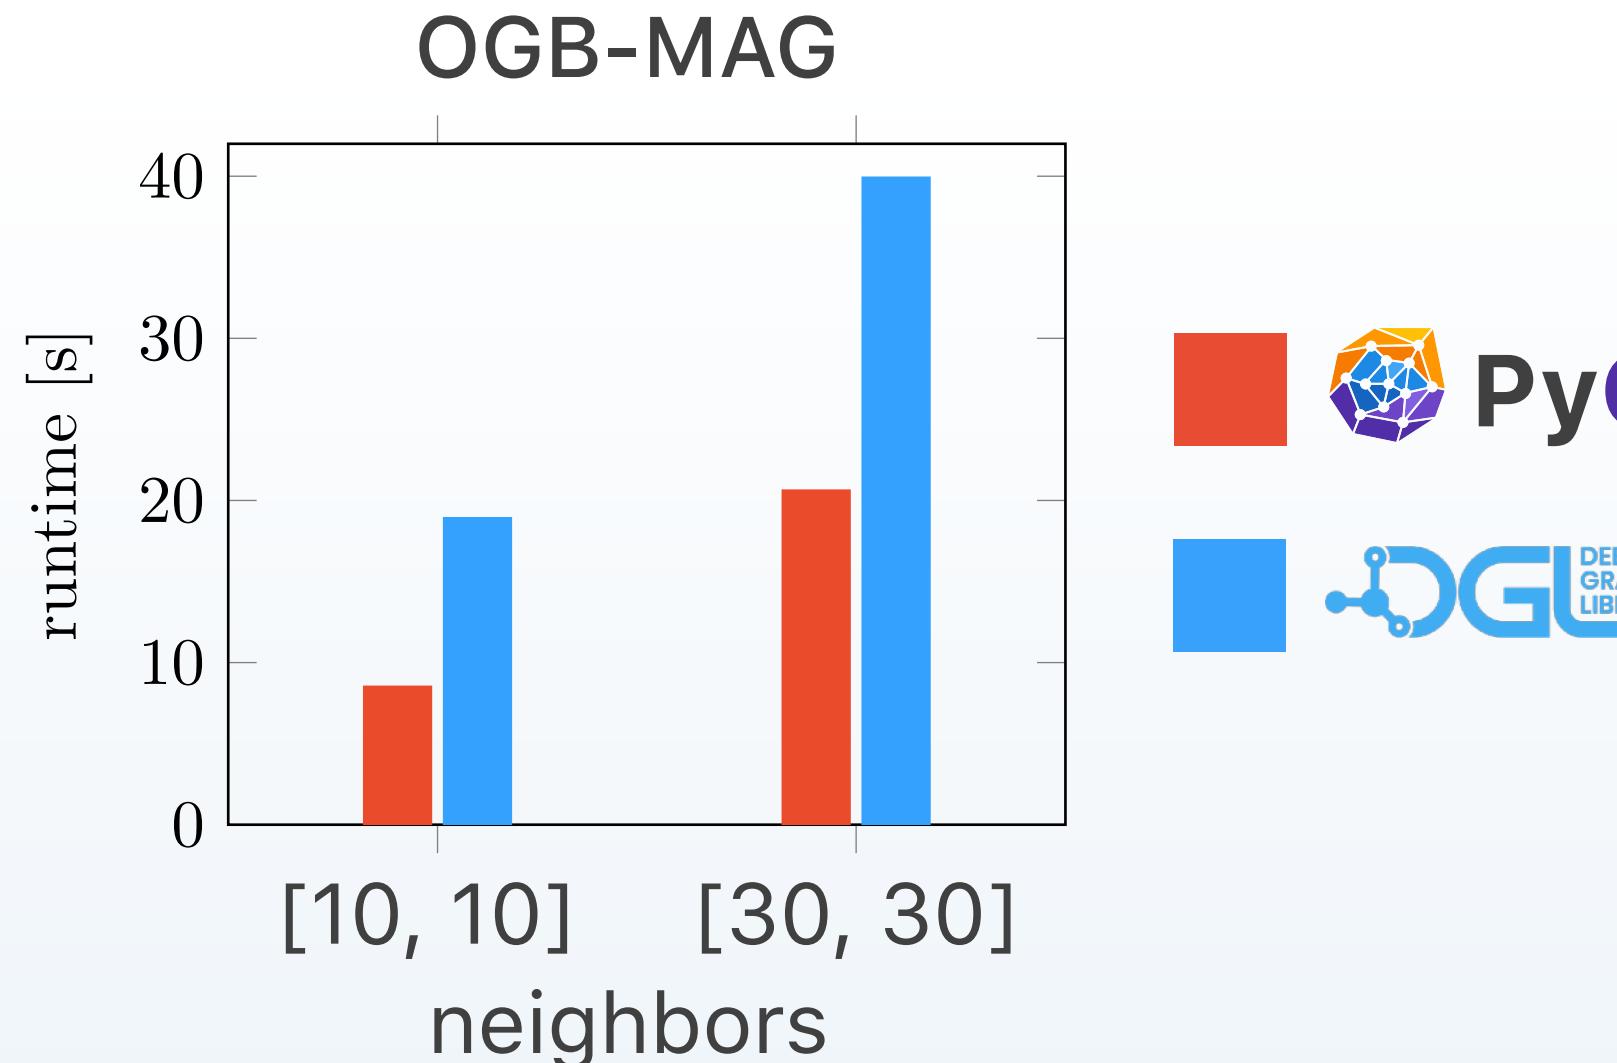


# Heterogeneous Graph Support

## Heterogeneous Graph Samplers

Scaling up heterogeneous GNNs  
to large-scale graphs with ease  
via relational neighbor sampling

Only requires a few lines  
of code change!



```
from torch_geometric.datasets import OGB_MAG
from torch_geometric.loader import NeighborLoader

data = OGB_MAG(path)[0]

loader = NeighborLoader(
    data,
    num_neighbors={key: [15, 10] for key in data.edge_types},
    batch_size=128,
    input_nodes='paper')

for batch in loader:
    model(batch.x_dict, batch.edge_index_dict)
```



# Heterogeneous Graph Support

- ✓ A tutorial introducing all newly released heterogeneous graph features
- ✓ A tutorial showcasing how to load heterogeneous graphs from raw \*.csv files
- ✓ Heterogeneous graph transformations
- ✓ Conversion from heterogeneous graphs to homogeneous "typed" graphs
- ✓ A generic wrapper (HeteroConv) for computing heterogeneous graph convolution via different message passing operators
- ✓ Lazy initialization (-1) for all message passing operators in  PyG
- ✓ A variety of heterogeneous GNN examples, including an example for scaling heterogeneous graph models via  PyTorch Lightning
- ✓ Dedicated heterogeneous graph operators (HGConv) and samplers (HGTLoder)

[https://pytorch-geometric.readthedocs.io  
en/latest/releasenotes/0.4.0.html#heterogeneous-graph-support](https://pytorch-geometric.readthedocs.io/en/latest/releasenotes/0.4.0.html#heterogeneous-graph-support)



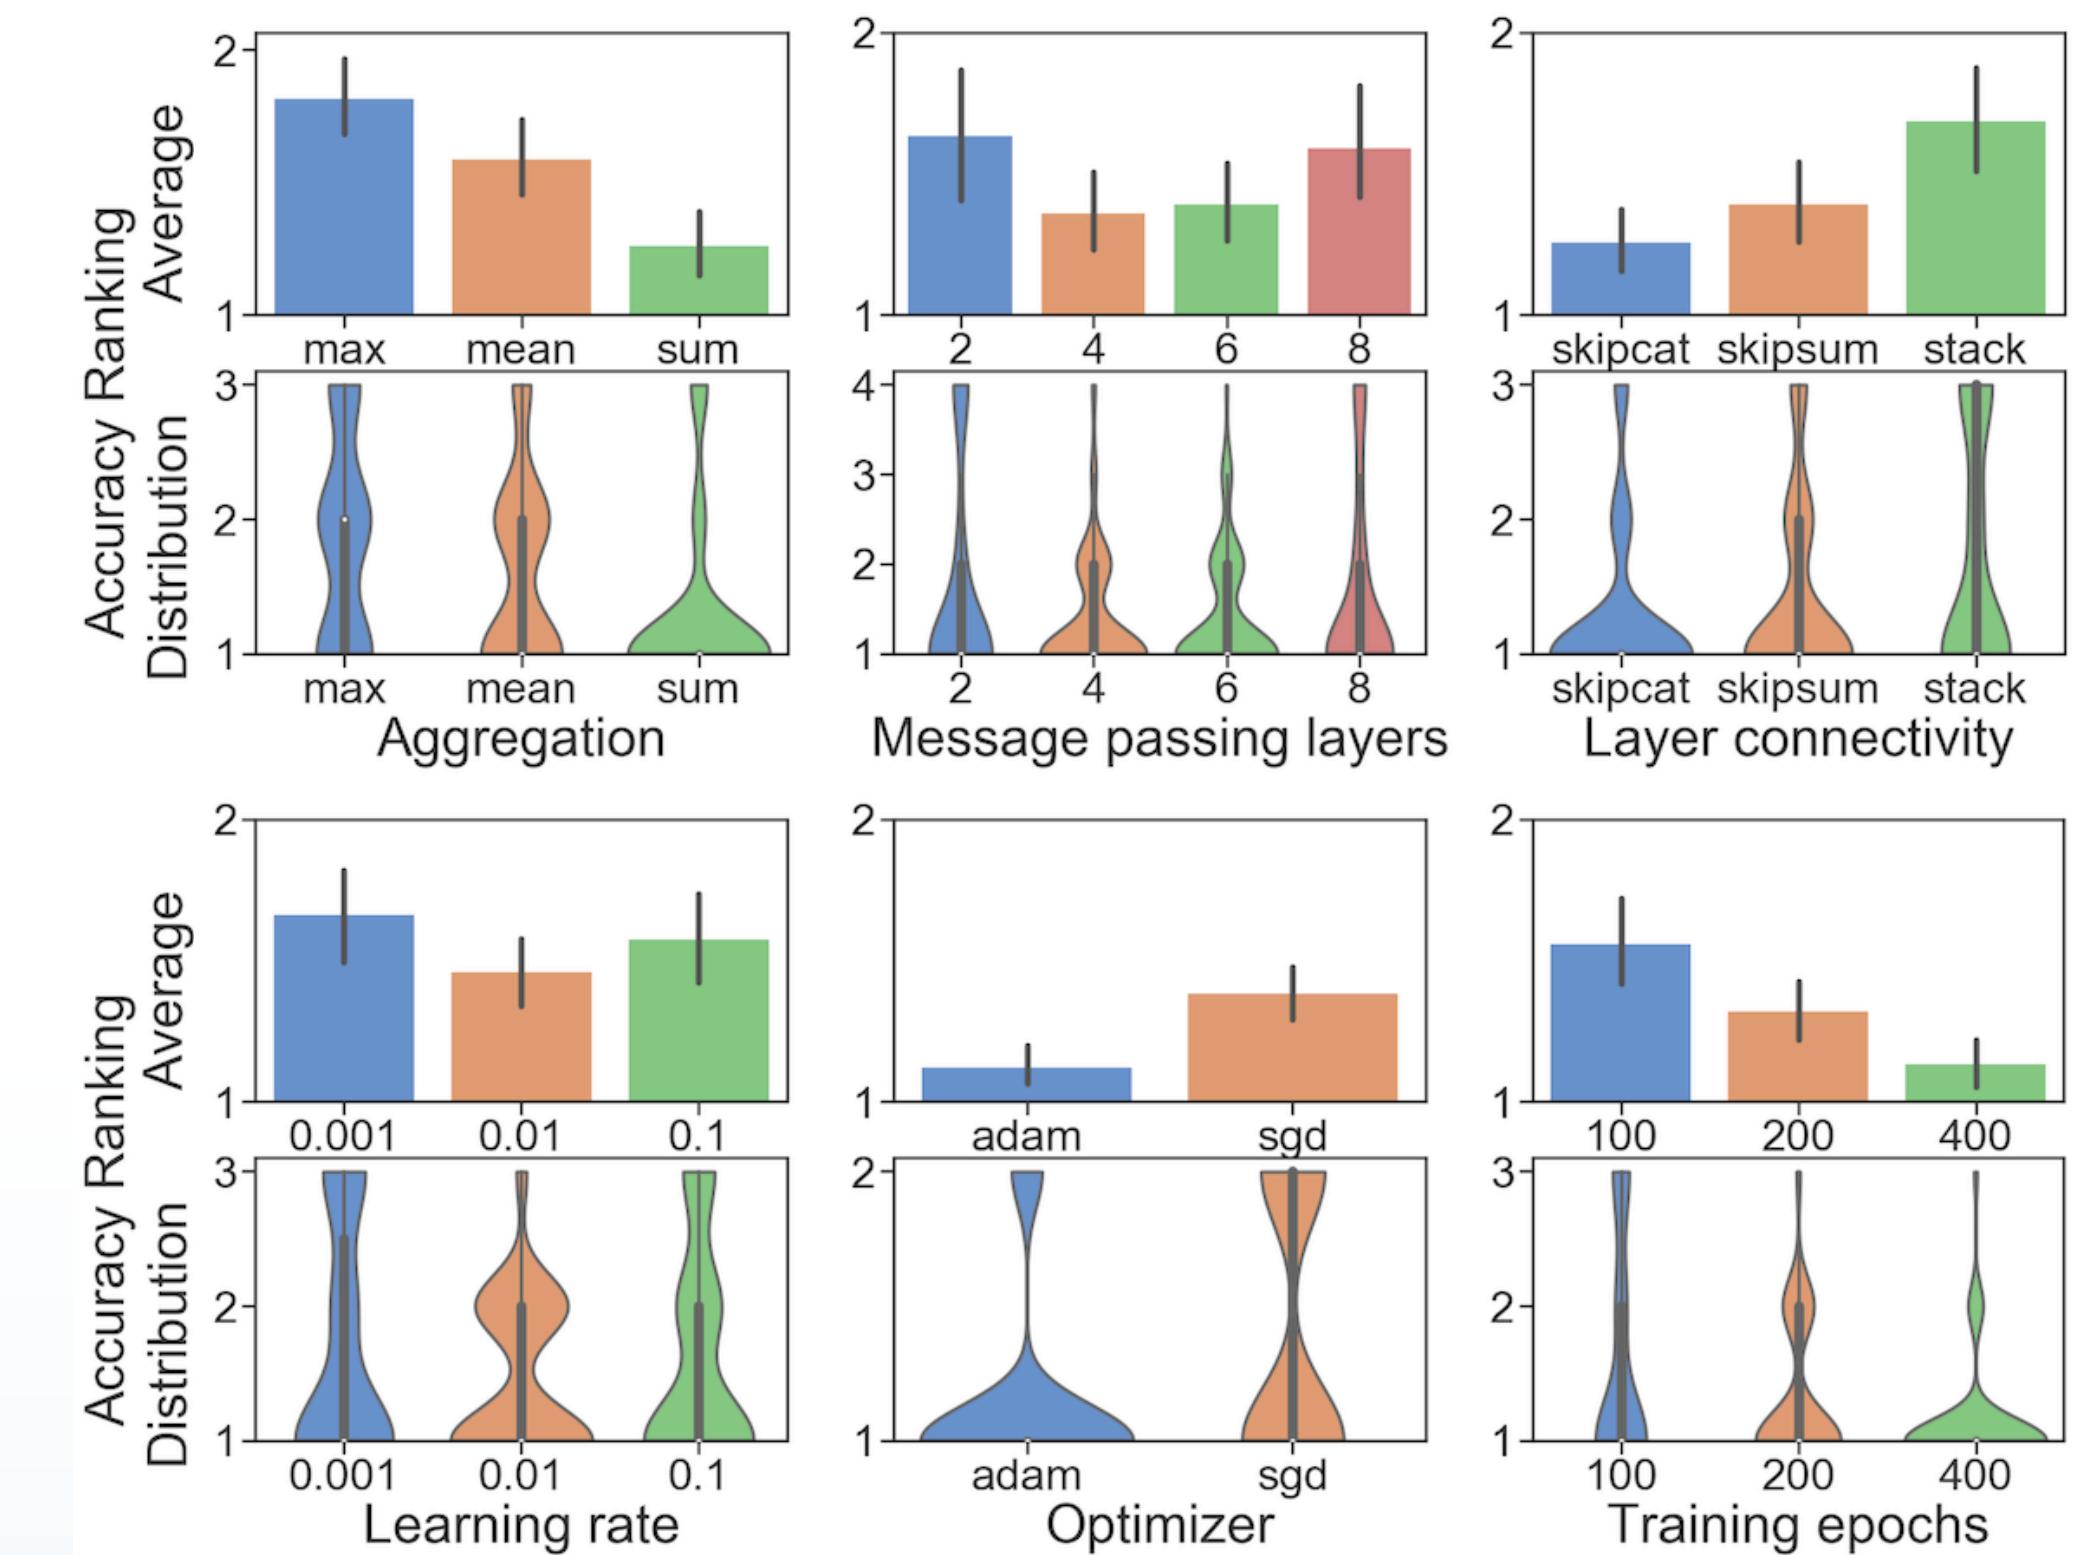
# GraphGym

## Design Space Exploration with GraphGym

*Which GNN is the best for your given task?*

GraphGym aims to design and train GNNs from configurations, using a modularized pipeline:

- Standardized GNN implementation/evaluation
- Design space exploration via simple interfaces to try out thousands of GNN architectures in parallel
- Hyper-parameter search and visualizations

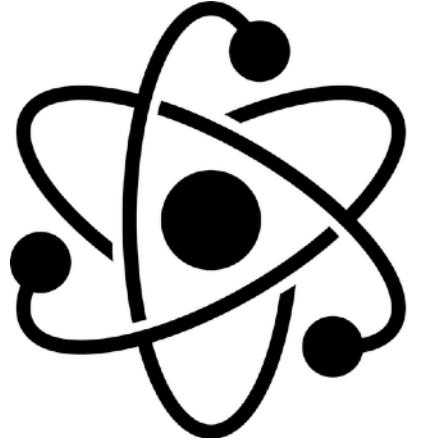




# Future Plans



# Future Plans

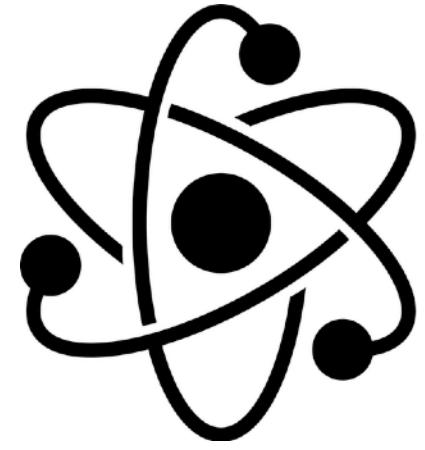


## Temporal Graphs

- (Nearly) all real-world graphs are inherently dynamic
- Support for addition and removal of nodes and edges
- Integration of temporal Graph Neural Networks
- Real-Time In-Stream Inference



# Future Plans



## Temporal Graphs

- (Nearly) all real-world graphs are inherently dynamic
- Support for addition and removal of nodes and edges
- Integration of temporal Graph Neural Networks
- Real-Time In-Stream Inference

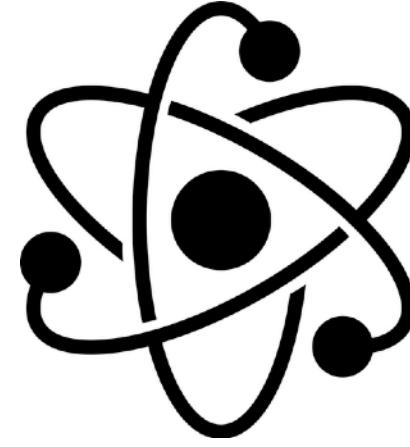


## Distributed Data

- While distributed training is possible, distributing data is currently a user task
- Scaling to billions of nodes via distributing input data
- Partitioning input node and edge features



# Future Plans



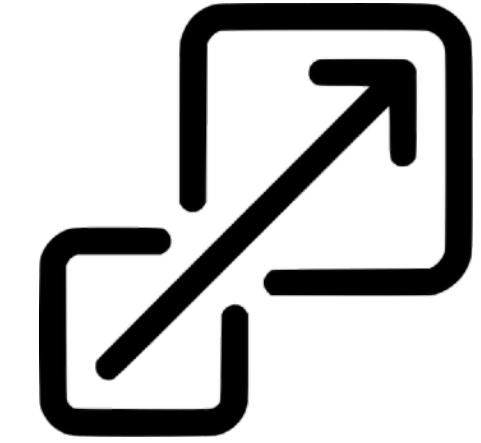
## Temporal Graphs

- (Nearly) all real-world graphs are inherently dynamic
- Support for addition and removal of nodes and edges
- Integration of temporal Graph Neural Networks
- Real-Time In-Stream Inference



## Distributed Data

- While distributed training is possible, distributing data is currently a user task
- Scaling to billions of nodes via distributing input data
- Partitioning input node and edge features



## Auto-Scaling

- PyG should automatically determine the best scalability approach for the given task
- Write GNNs in full-batch mode and let PyG figure out the rest
- GNNAutoScale (ICML 2021)



# Conclusion

 **PyG** bundles the state-of-the-art in Graph Representation Learning

- ✓ 50+ GNN architectures
- ✓ 200+ benchmark datasets
- ✓ Dedicated sparsity-aware CUDA kernels
- ✓ Multi-GPU support
- ✓ Half-Precision support
- ✓ Support for scalability techniques
- ✓ Heterogeneous graph support
- ✓ GNN Design Space Exploration

We are constantly encouraged  
to make PyG even better!



dortmund  
university



**Stanford**  
University

team@pyg.org

<https://pyg.org>

 /pyg-team/pytorch-geometric

license MIT

PRs welcome



`conda install pyg -c pyg`